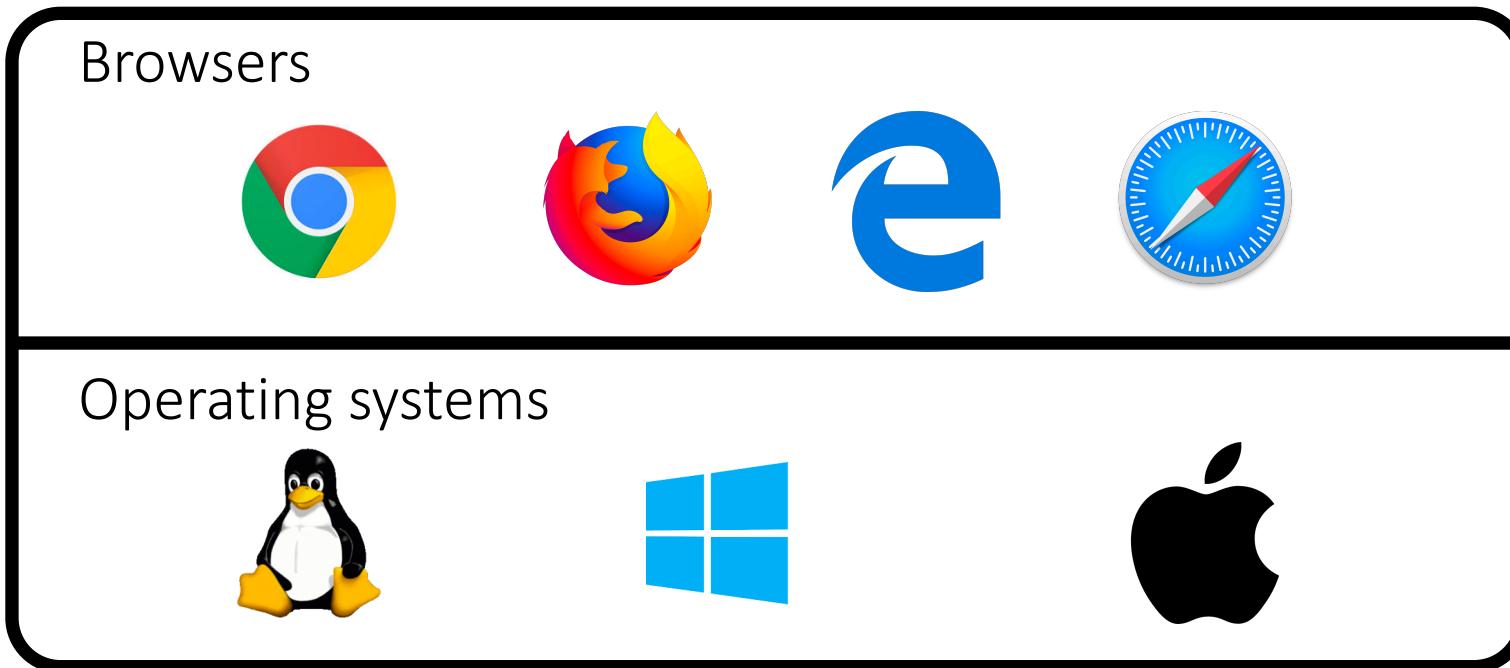


# Concolic Execution Tailored for Hybrid Fuzzing

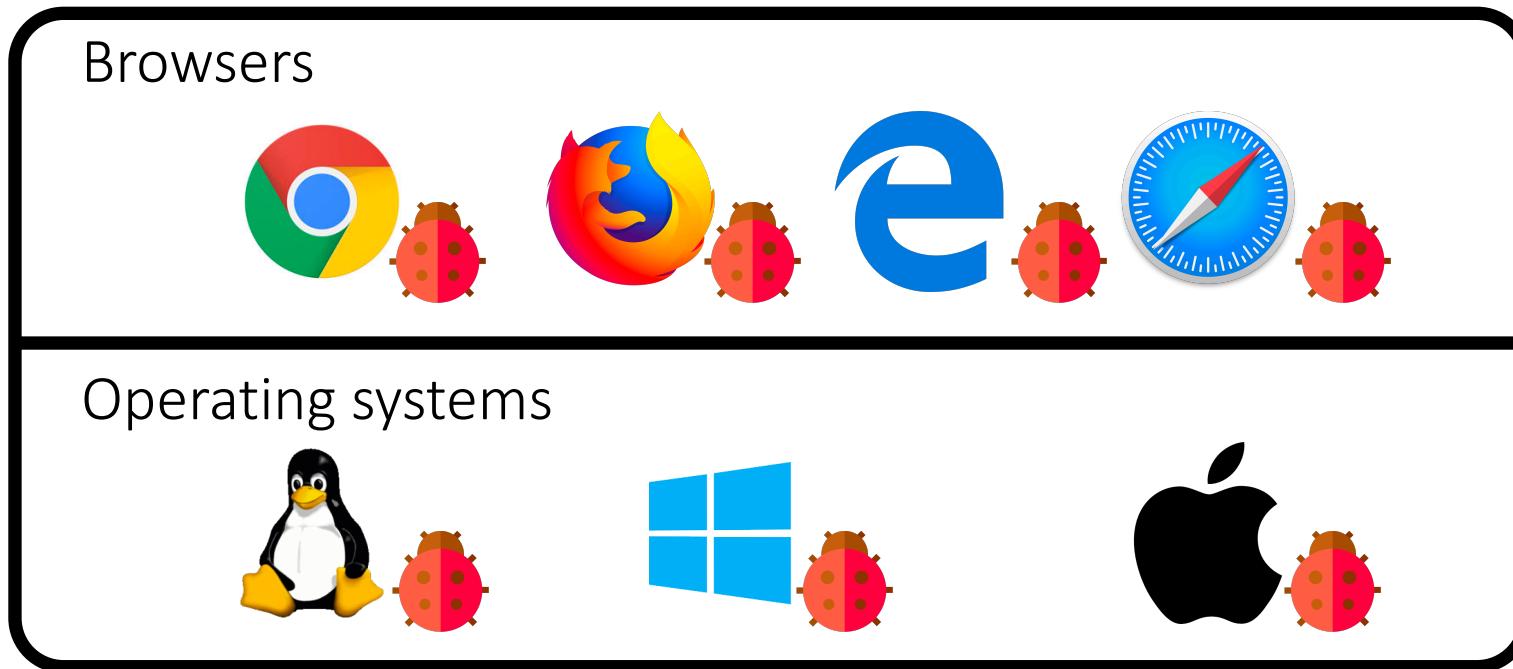
Insu Yun

Georgia Institute of Technology

# We are using many software applications



# We are using many (vulnerable) software applications



# Attack flow (hacking)

Vulnerability

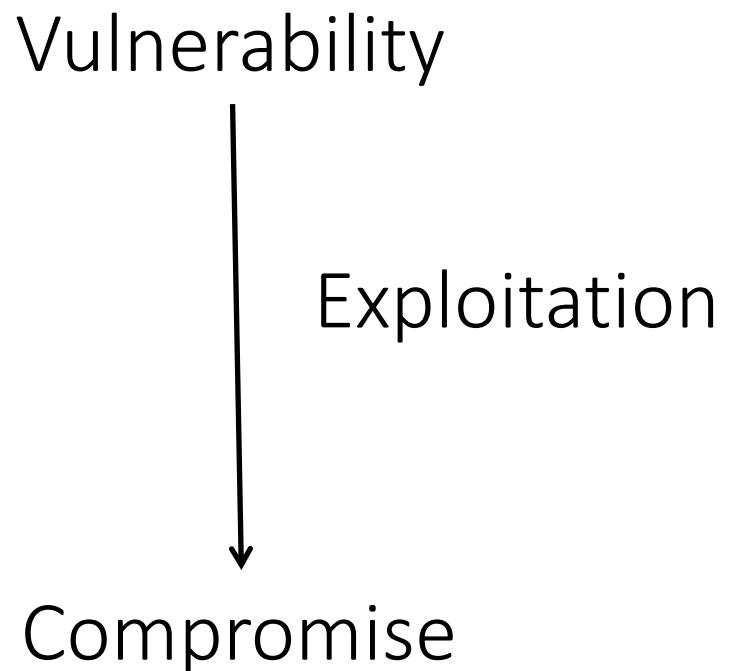
# Attack flow (hacking)

Vulnerability

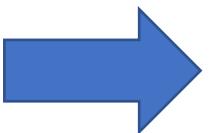
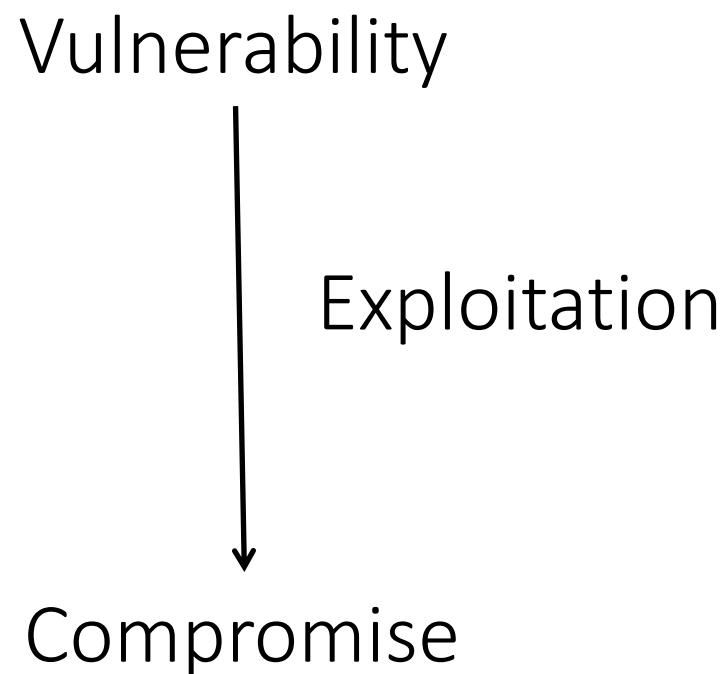


Exploitation

# Attack flow (hacking)



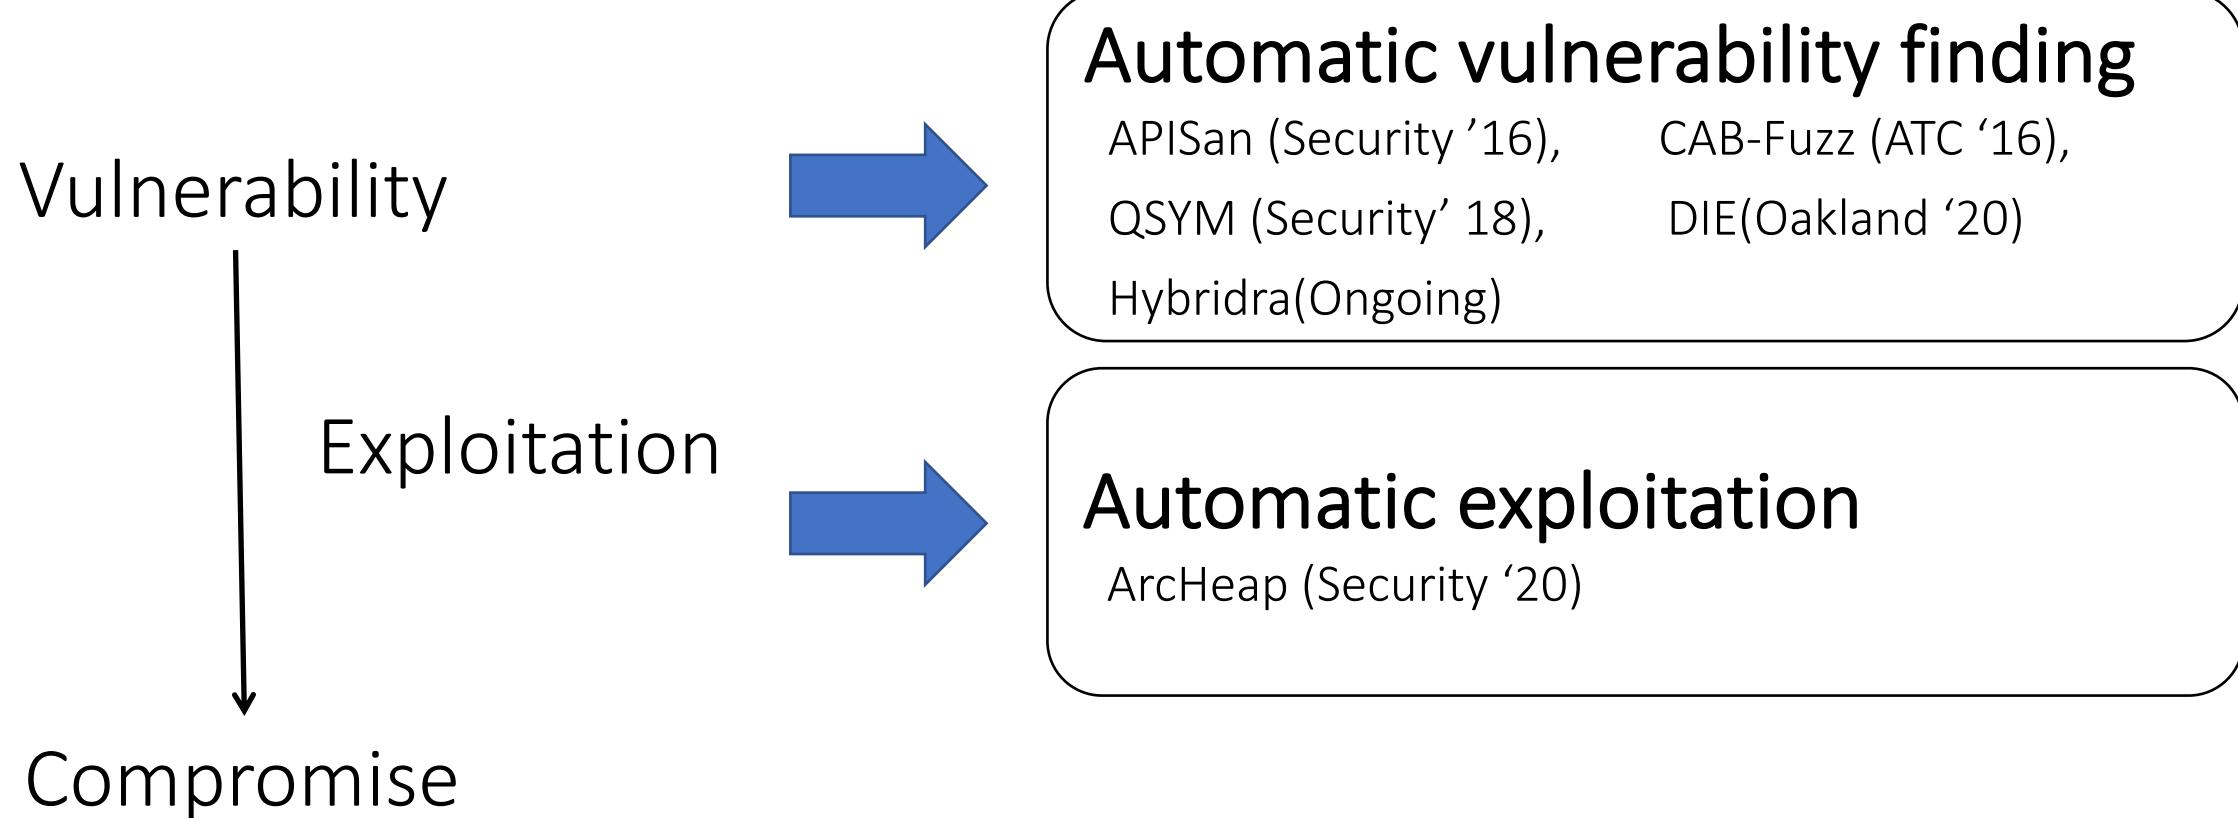
# Attack flow (hacking)



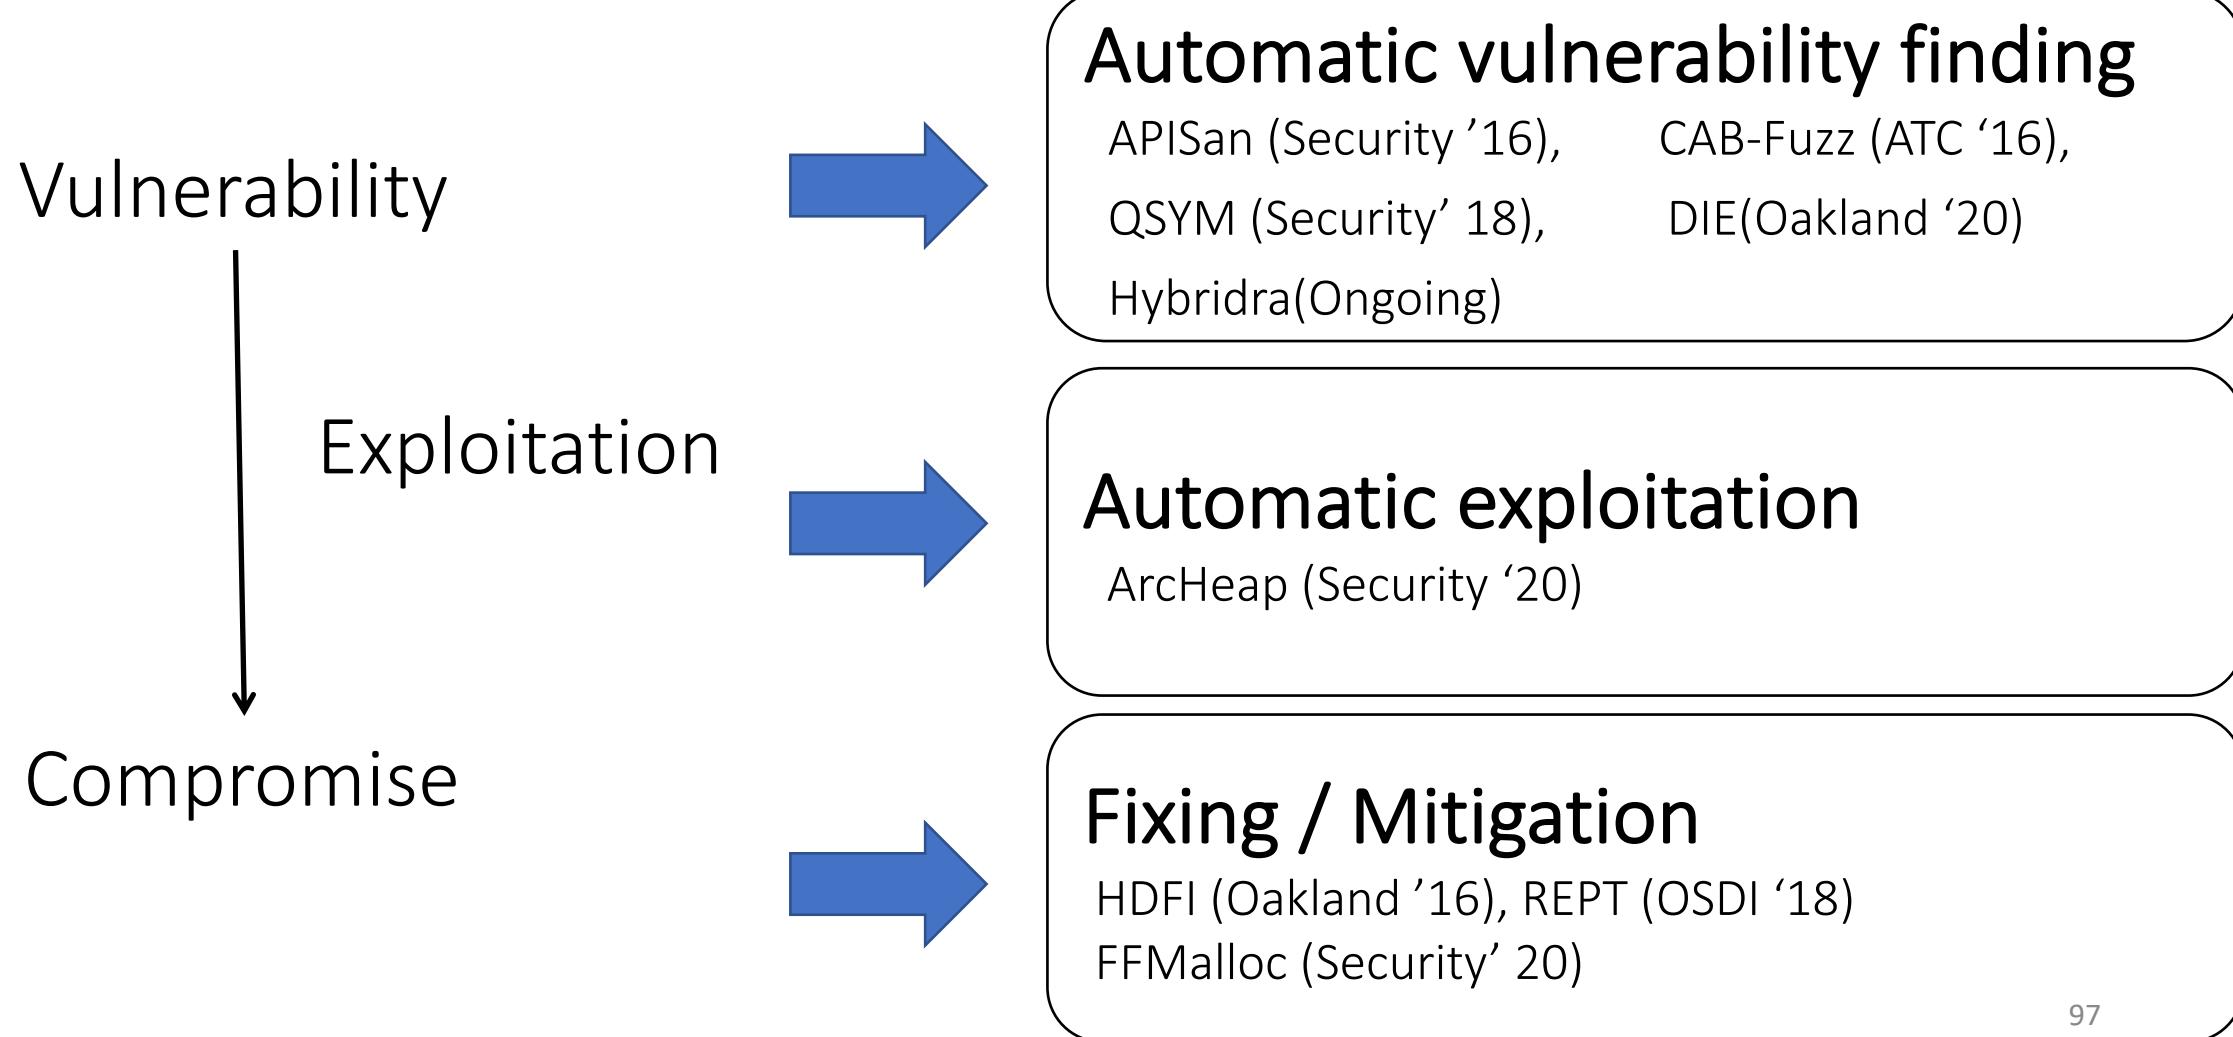
**Automatic vulnerability finding**

APISan (Security '16), CAB-Fuzz (ATC '16),  
QSYM (Security' 18), DIE(Oakland '20)  
Hybridra(Ongoing)

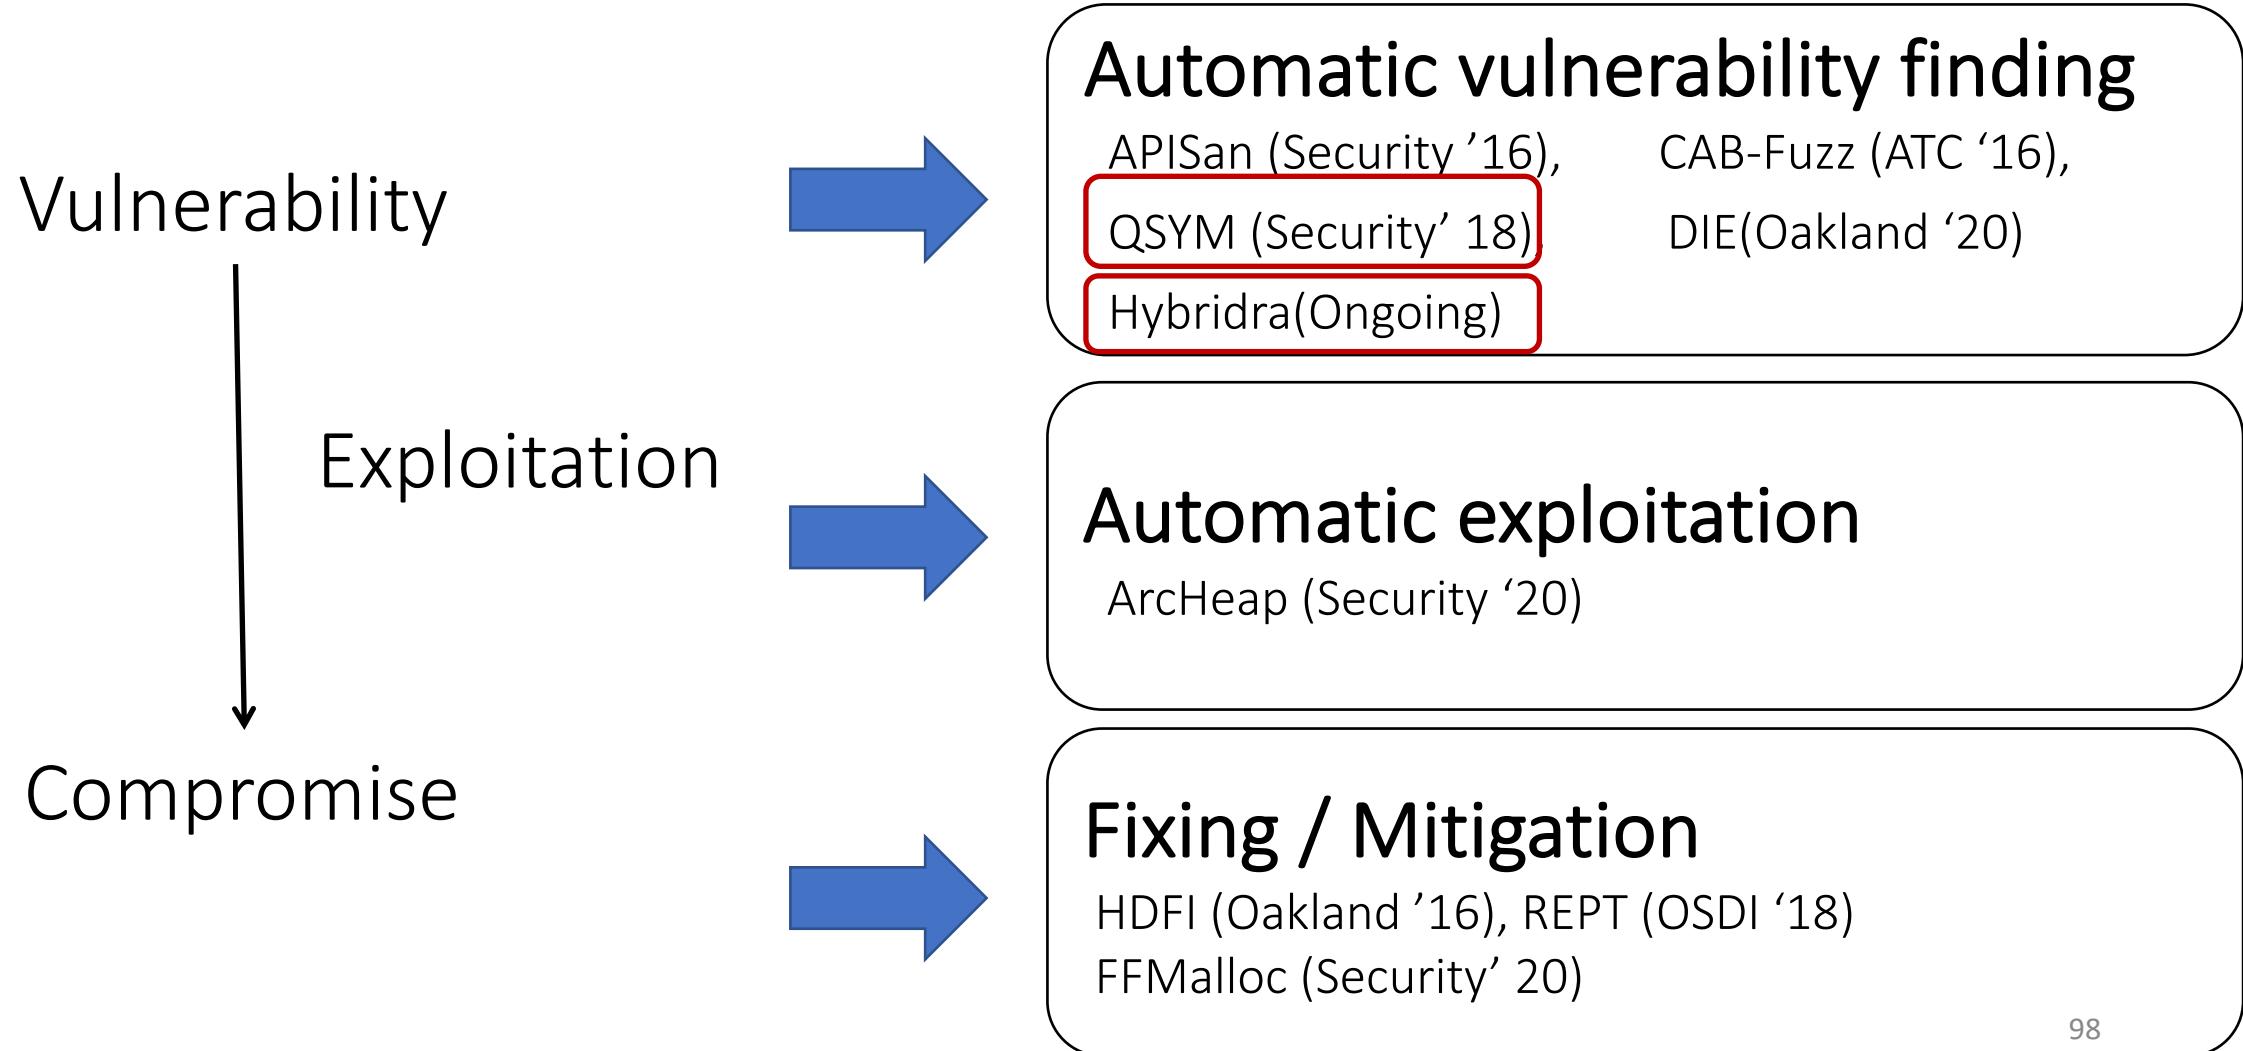
# Attack flow (hacking)



# Attack flow (hacking)



# Attack flow (hacking)



# Today's talk

QSYM: A Binary-level  
Concolic Execution Engine  
for *Hybrid fuzzing*

- Binary
- User applications

Hybridra: A *Hybrid Fuzzer*  
for Kernel File Systems

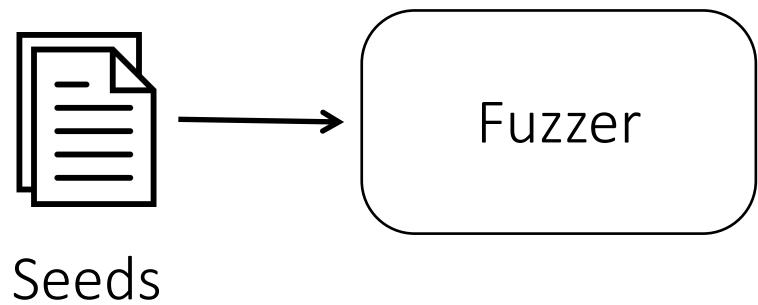
- Source code
- File systems

# Overview of 40-year-old random testing (fuzzing)



Seeds

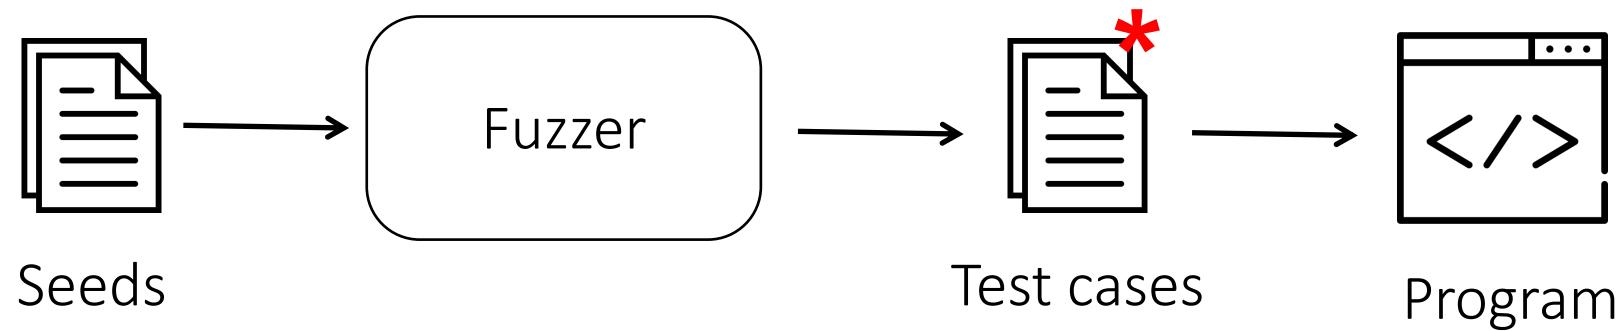
# Overview of 40-year-old random testing (fuzzing)



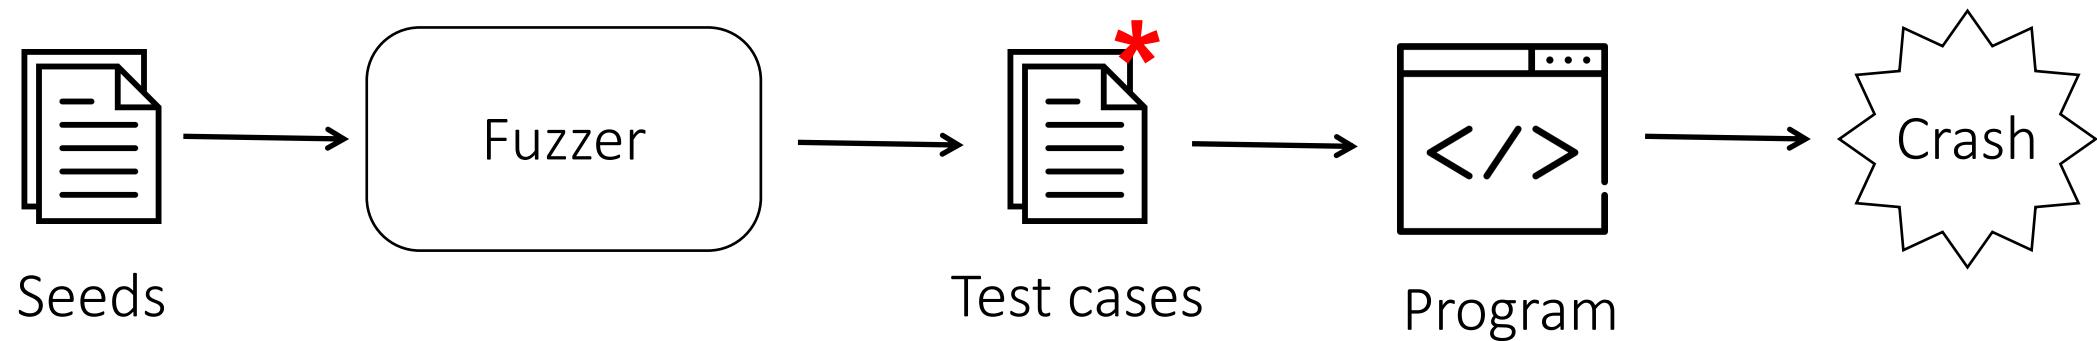
# Overview of 40-year-old random testing (fuzzing)



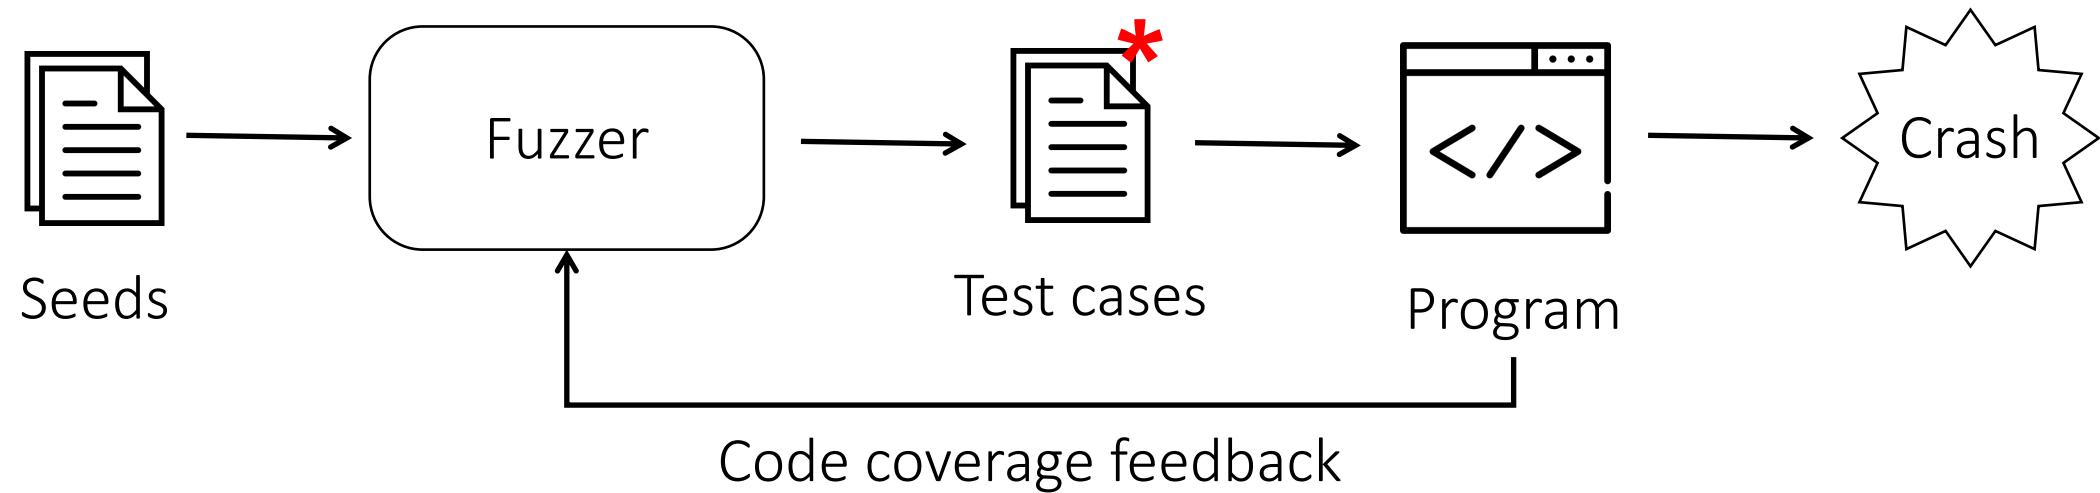
# Overview of 40-year-old random testing (fuzzing)



# Overview of 40-year-old random testing (fuzzing)



# Recent breakthrough: Code coverage feedback



# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

Test cases

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

Test cases

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'PTSY'

Test cases

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'PTSY'

Test cases

x = 'H4SY'

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'PTSY'

Test cases

x = 'H4SY'

x = 'O4SY'

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'PTSY'

Test cases

x = 'H4SY'

x = 'O4SY'

...

# Before code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'PTSY'

Test cases

x = 'H4SY'

x = 'O4SY'

...

$$P(\text{crash}) = 2^{-32}$$

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



```
x = 'E4SY'
```

Seeds

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()  
  
↓
```



x = 'E4SY'

Seeds



x = 'E4SI'

Test cases

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()  
  
↓
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

Test cases

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()  
  
↓
```



x = 'E4SY'

Seeds



x = 'E4SI'

Test cases

x = 'S4SY'

x = 'ETSY'

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()  
  
↓
```



x = 'E4SY'

Seeds



x = 'E4SI'

Test cases

x = 'S4SY'

x = 'ETSY'

x = 'H4SY'

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'H4SY'

Test cases

x = 'S4SY'

x = 'ETSY'

# After code coverage feedback,

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'E4SY'

Seeds



x = 'E4SI'

x = 'S4SY'

x = 'ETSY'

Test cases

x = 'H4SY'

New code  
coverage!

# Generate test cases from a test case that introduces new code coverage

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'EASY'

x = 'H4SY'

Seeds

# Generate test cases from a test case that introduces new code coverage

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'EASY'

x = 'H4SY'

Seeds



x = 'H4SI'

Test cases

# Generate test cases from a test case that introduces new code coverage

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'EASY'

x = 'H4SY'



Test cases

x = 'H4SI'

x = 'HASY '

New code  
coverage!

# Generate test cases from a test case that introduces new code coverage

```
x = input()  
  
if x[0] == 'H':  
    if x[1] == 'A':  
        if x[2] == 'R':  
            if x[3] == 'D':  
                crash()
```



x = 'EASY'

x = 'H4SY'



x = 'H4SI'

x = 'HASY '

Seeds

Test cases

New code  
coverage!

$$P(\text{crash}) = 2^{-32}$$

# Generate test cases from a test case that introduces new code coverage

```
x = input()
if x[0] == 'H':
    if x[1] == 'A':
        if x[2] == 'R':
            if x[3] == 'D':
                crash()
```



`x = 'EASY'`

`x = 'H4SY'`

Seeds



`x = 'H4SI'`

`x = 'HASY '`

Test cases

New code  
coverage!

$$P(\text{crash}) = \cancel{2^{-32}} = 2^{-8} \times 2^{-2} = 2^{-10}$$

Per-byte      4 bytes

# Coverage-guided fuzzing is effective

```
american fuzzy lop 0.47b (readpng)
process timing
  run time : 0 days, 0 hrs, 4 min, 43 sec
  last new path : 0 days, 0 hrs, 0 min, 26 sec
  last uniq crash : none seen yet
  last uniq hang : 0 days, 0 hrs, 1 min, 51 sec
cycle progress
  now processing : 38 (19.4%)
  paths timed out : 0 (0.0%)
  stages tested : 32/8
  now trying interest: 32/8
  stage execs : 0/9990 (0.00%)
  total execs : 654k
  exec speed : 2306/sec
fuzzing strategy details
  dict types: 88/14.4k, 6/14.4k, 6/14.4k
  byte flips: 0/1804, 0/1786, 1/1750
  arithmetics: 31/126k, 3/45.6k, 1/17.8k
  known ints: 1/1254k, 6/65.8k, 6/78.2k
  havoc: 1/1254k, 0/0
  trim: 2876 B/931 (61.45% gain)
overall results
  cycles done : 0
  total paths : 195
  uniq crashes : 0
  uniq hangs : 1
map coverage
  map density : 1217 (7.43%)
  count coverage : 1.55 bits/tuple
  max tuple depth : 1
  favored paths : 128 (65.64%)
  new edges on : 85 (43.59%)
  total crashes : 0 (0 unique)
  total hangs : 1 (1 unique)
pitch geometry
  level 0 : 3
  pending : 178
  pend fav : 114
  imported : 0
  variants : 0
  latent : 0
```

AFL

- Fuzzer developed by Google
- Re-discover coverage-guided fuzzing
- Found hundreds of bugs in many programs e.g.,) Safari, Firefox, OpenSSL, ...

# Coverage-guided fuzzing is effective

```
american fuzzy top 0.47b (readpng)
process timing
  run time : 0 days, 0 hrs, 4 min, 43 sec
  last new path : 0 days, 0 hrs, 0 min, 26 sec
  last seen crash : none seen yet
  last uniq hang : 0 days, 0 hrs, 1 min, 51 sec
cycle progress
  now processing : 38 (19.4%)
  paths timed out : 0 (0.0%)
  stages tested : 32/8
  now trying interest: 32/8
  stage execs : 0/9990 (0.00%)
  total execs : 654k
  exec speed : 2306/sec
fuzzing strategy details
  dict types: 88/14.4k, 6/14.4k, 6/14.4k
  byte flips: 0/1804, 0/1786, 1/1750
  arithmetic: 31/126k, 3/45.6k, 1/17.8k
  known ints: 1/1254k, 6/65.8k, 6/78.2k
  havoc: 1/1254k, 0/0
  trim: 2876 B/931 (61.45% gain)
overall results
  cycles done : 0
  total paths : 195
  uniq crashes : 0
  uniq hangs : 1
map coverage
  map density : 1217 (7.43%)
  count coverage : 1.55 bits/tuple
  max edge depth : 128
  favored paths : 128 (65.64%)
  new edges on : 85 (43.59%)
  total crashes : 0 (0 unique)
  total hangs : 1 (1 unique)
pitch geometry
  level 0 : 3
  pending : 178
  pend fav : 114
  imported : 0
  variants : 0
  latent : 0
```

AFL



[LLVM Home](#) | [Documentation](#) »

**libFuzzer – a library for coverage-guided fuzz testing.**

libFuzzer

- Fuzzer developed by Google
- Re-discover coverage-guided fuzzing
- Found hundreds of bugs in many programs e.g.,) Safari, Firefox, OpenSSL, ...

- LLVM community developed
- A library to include random testing as a part of projects  
e.g.,) LLVM, Chromium, Tensorflow, ...

# Coverage-guided fuzzing is effective

```
american fuzzy top 0.47b (readpng)
process timing
  run time : 0 days, 0 hrs, 4 min, 43 sec
  last new path : 0 days, 0 hrs, 26 sec
  last unique crash : none seen yet
  last uniq hang : 0 hrs, 1 min, 51 sec
overall results
  cycles done : 0
  total paths : 195
  uniq crashes : 0
  uniq hangs : 1
cycle progress
  now processing : 38 (19.49%)
  paths timed out : 0 (0.00%)
  stages tested : 32/8
  now trying interest: 32/8
  stage execs : 0/9990 (0.00%)
  total execs : 654k
  exec speed : 2306/sec
fuzzing strategy details
  dict types: 88/14.4k, 6/14.4k, 6/14.4k
  byte flips: 0/1804, 0/1786, 1/1750
  arithmetics: 31/126k, 3/45.6k, 1/17.8k
  known ints: 1/125k, 1/65.8k, 6/78.2k
  havoc: 1/254k, 0/0
  trim: 2876 B/931 (61.45% gain)
map coverage
  map density : 1217 (7.43K)
  count coverage : 12.55 bits/tuple
  max depth : 128 (65.64%)
  favored paths : 85 (43.59%)
  new edges on : 0 (0 unique)
  total crashes : 0 (0 unique)
  total hangs : 1 (1 unique)
pitch geometry
  level 3 pending : 178
  pend fav : 114 imported : 0
  var pending : 0 latent : 0
```

AFL



[LLVM Home](#) | [Documentation](#) »

[libFuzzer – a library for coverage-guided fuzz testing.](#)

libFuzzer



OSS-Fuzz

- Fuzzer developed by Google
- Re-discover coverage-guided fuzzing
- Found hundreds of bugs in many programs e.g.,) Safari, Firefox, OpenSSL, ...
- LLVM community developed
- A library to include random testing as a part of projects  
e.g.,) LLVM, Chromium, Tensorflow, ...
- Use Google's cloud resources to fuzz open-source software
- 4 trillion test cases a week

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
  
// 459684 == 6782  
if x * x == 459684 :  
    crash()
```

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```



Seeds

```
x = 0
```

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



```
x = 0
```

Seeds



```
x = 3
```

Test cases

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

Test cases

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

x = 942

Test cases

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

x = 942

Test cases

x = 512

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

x = 942

Test cases

x = 512

x = 28

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

x = 942

Test cases

x = 512

x = 28

...

# Limitations of fuzzing: Randomly hard-to-find

```
x = int(input())  
// 459684 == 6782  
if x*x == 459684 :  
    crash()
```



x = 0

Seeds



x = 3

x = 452

x = 942

Test cases

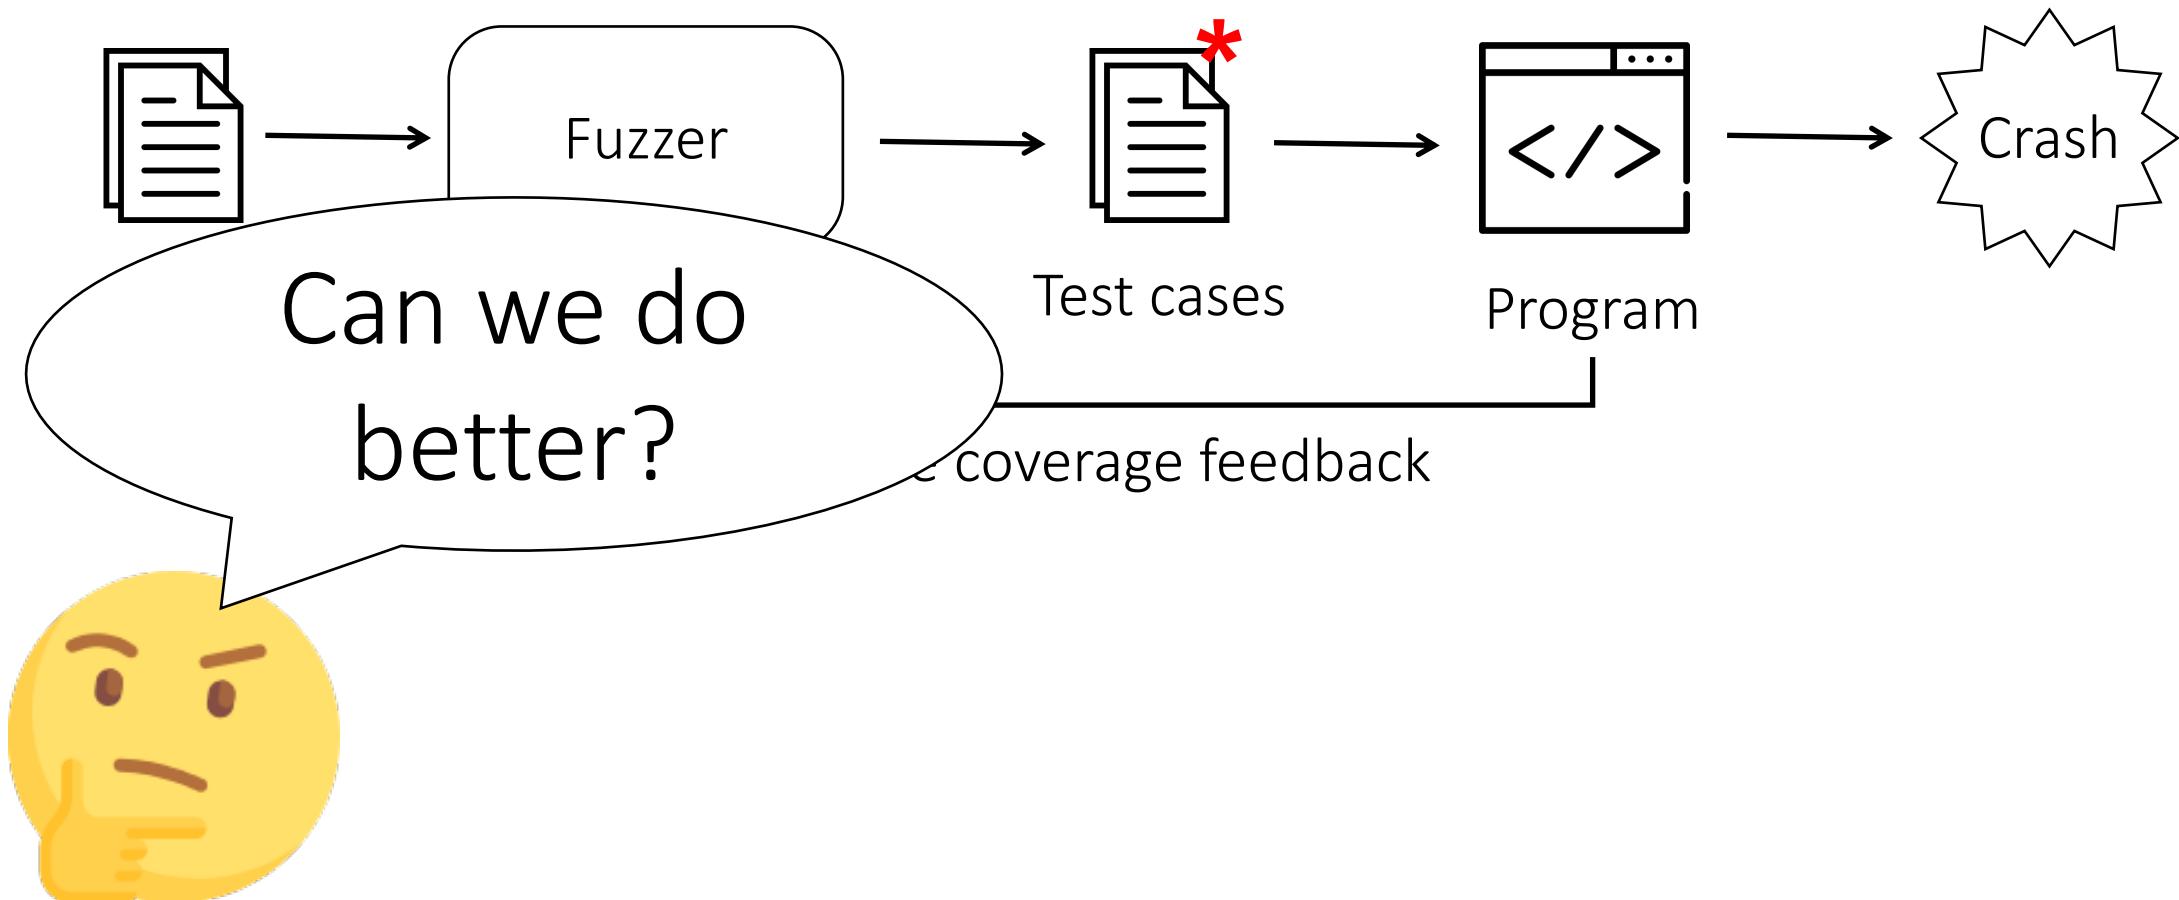
x = 512

x = 28

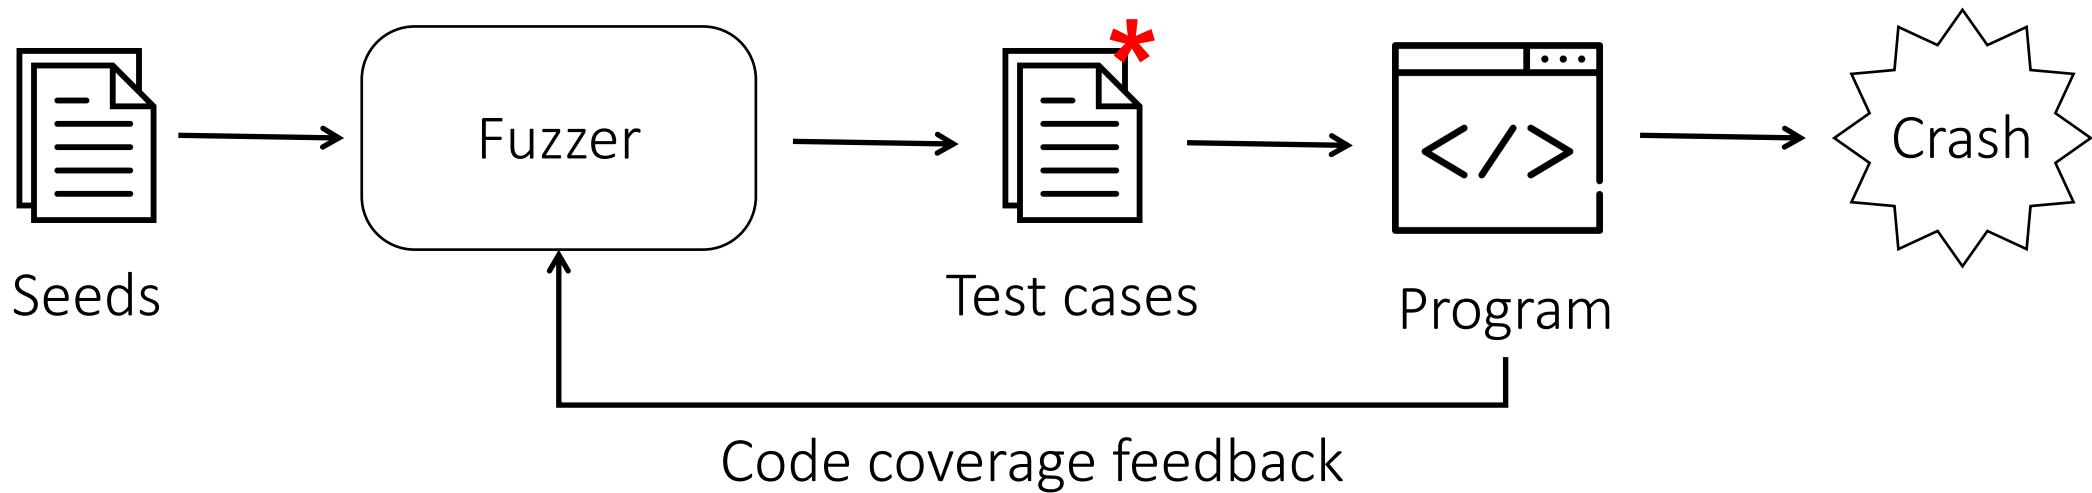
...

$$P(\text{crash}) = 2^{-32}$$

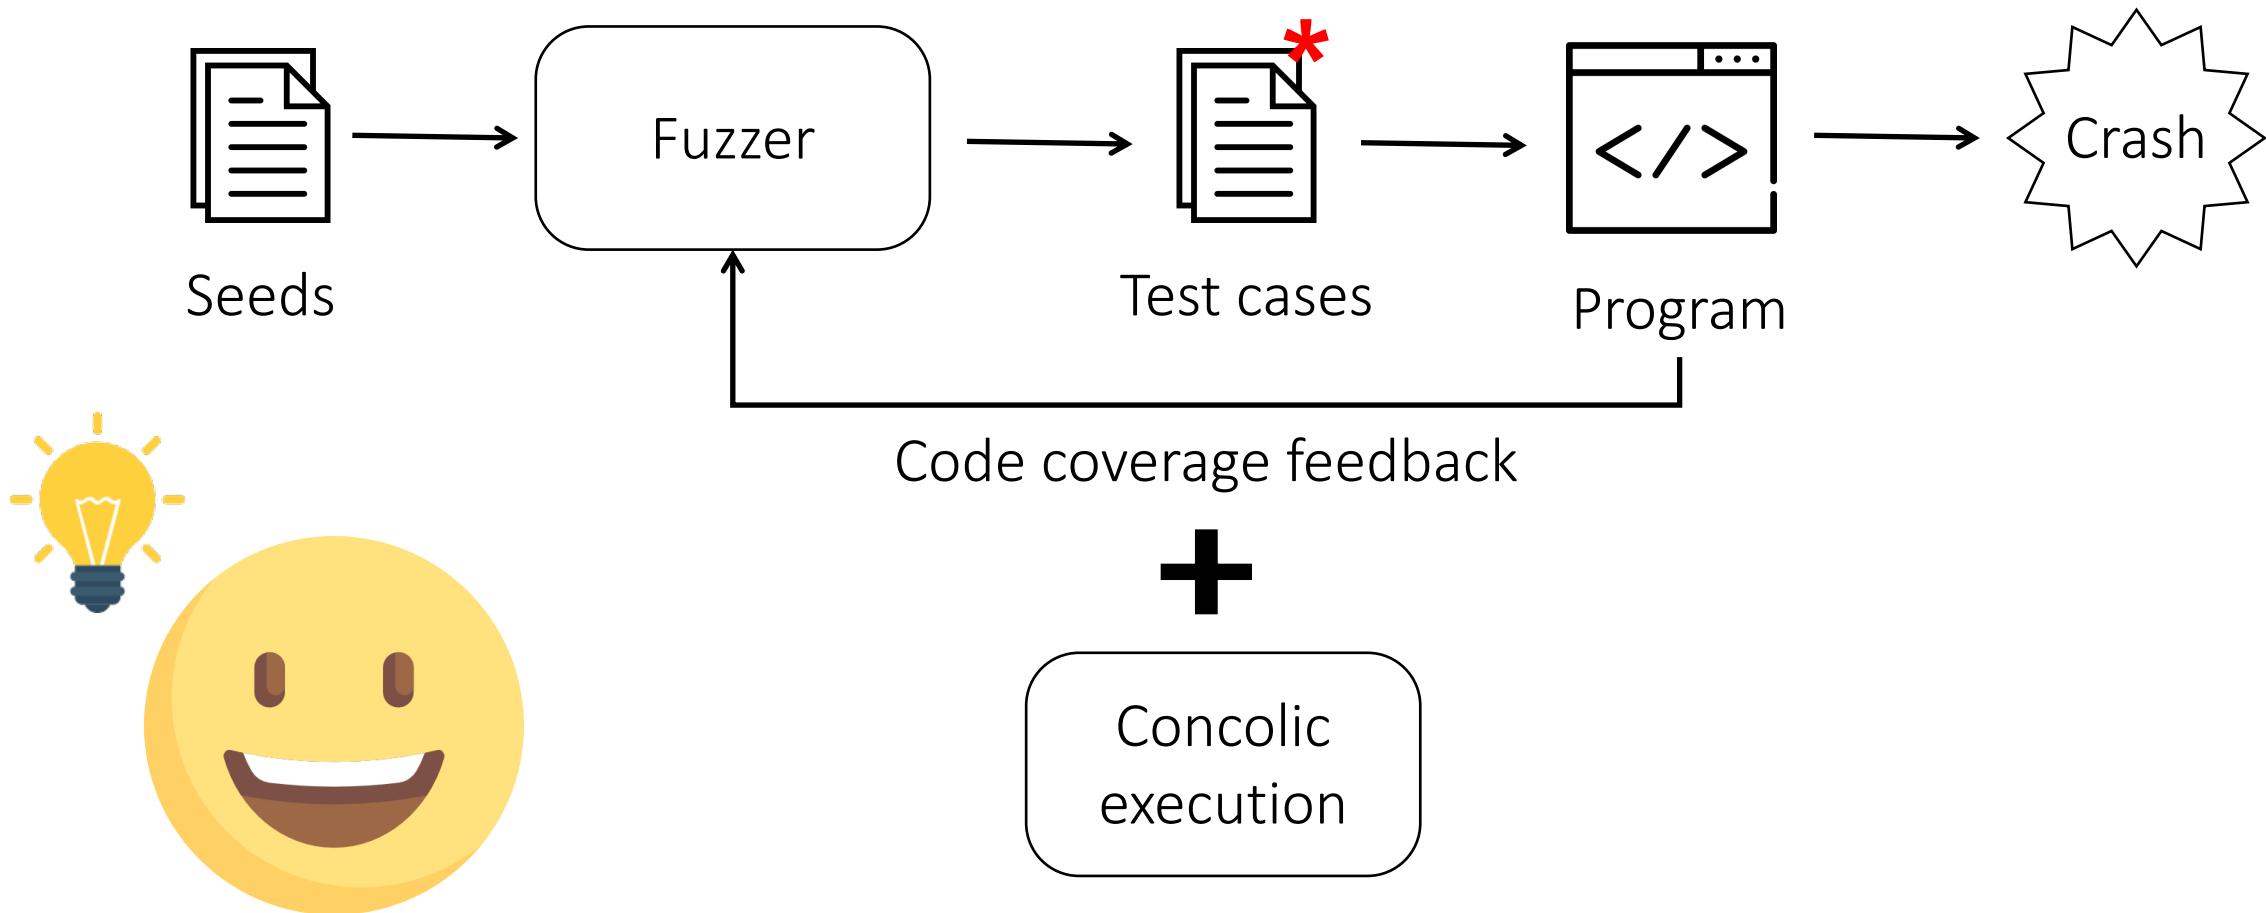
# Limitations of coverage-guided fuzzing



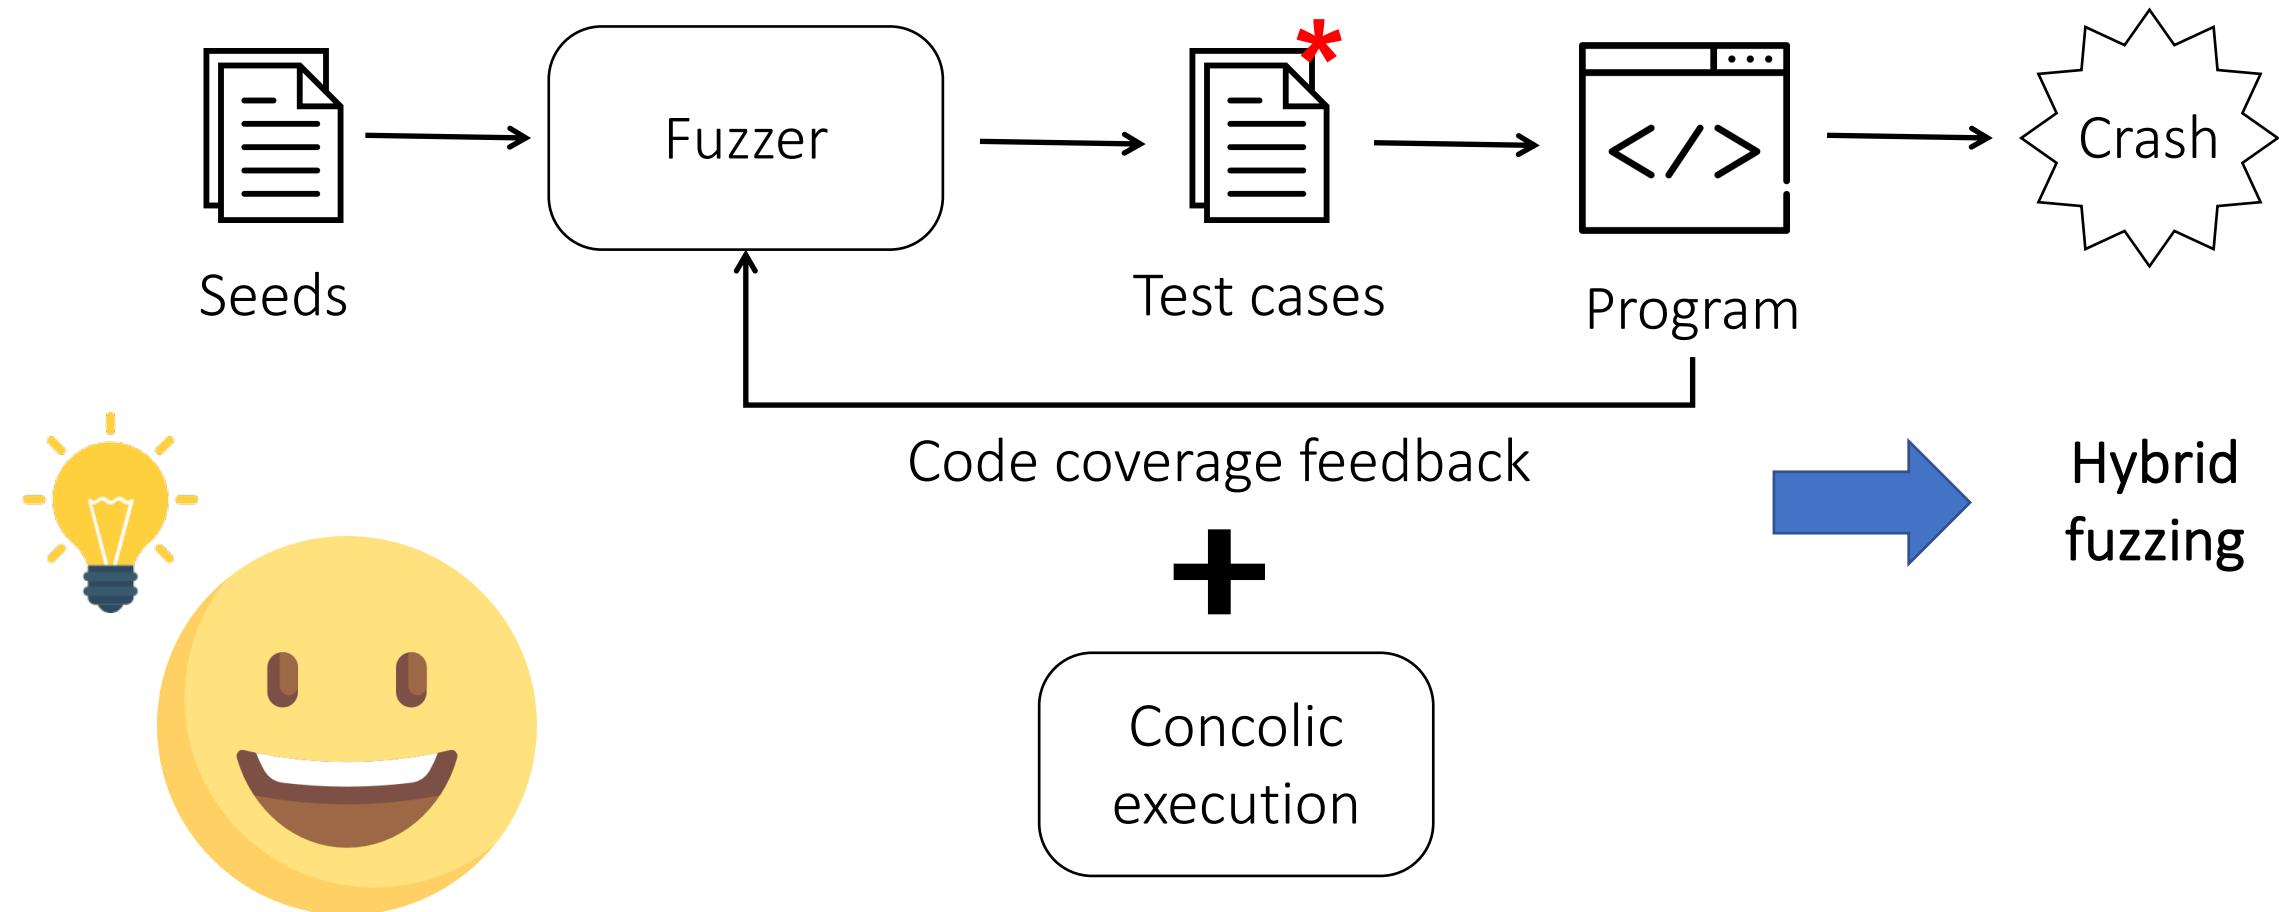
# Limitations of coverage-guided fuzzing



# Limitations of coverage-guided fuzzing



# Limitations of coverage-guided fuzzing



Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
```

```
if x * x == 459684 :  
    crash()
```

# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```

```
x = input()
```

# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```

```
x = input()
```

```
if x * x == 459684 :
```

# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```

```
x = input()
```

```
if x * x == 459684 :
```

```
x = 0
```

# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```

```
x = input()
```

```
if x * x == 459684 :
```

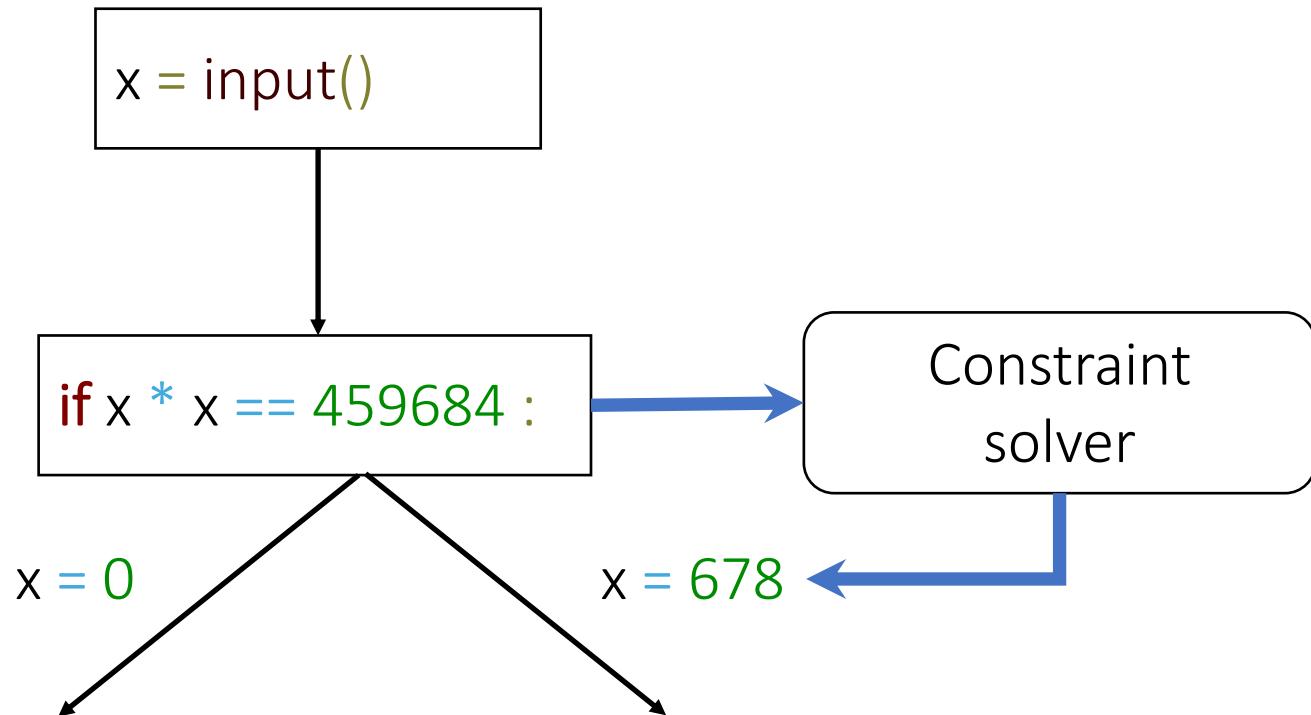
```
Constraint  
solver
```

```
x = 0
```

# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

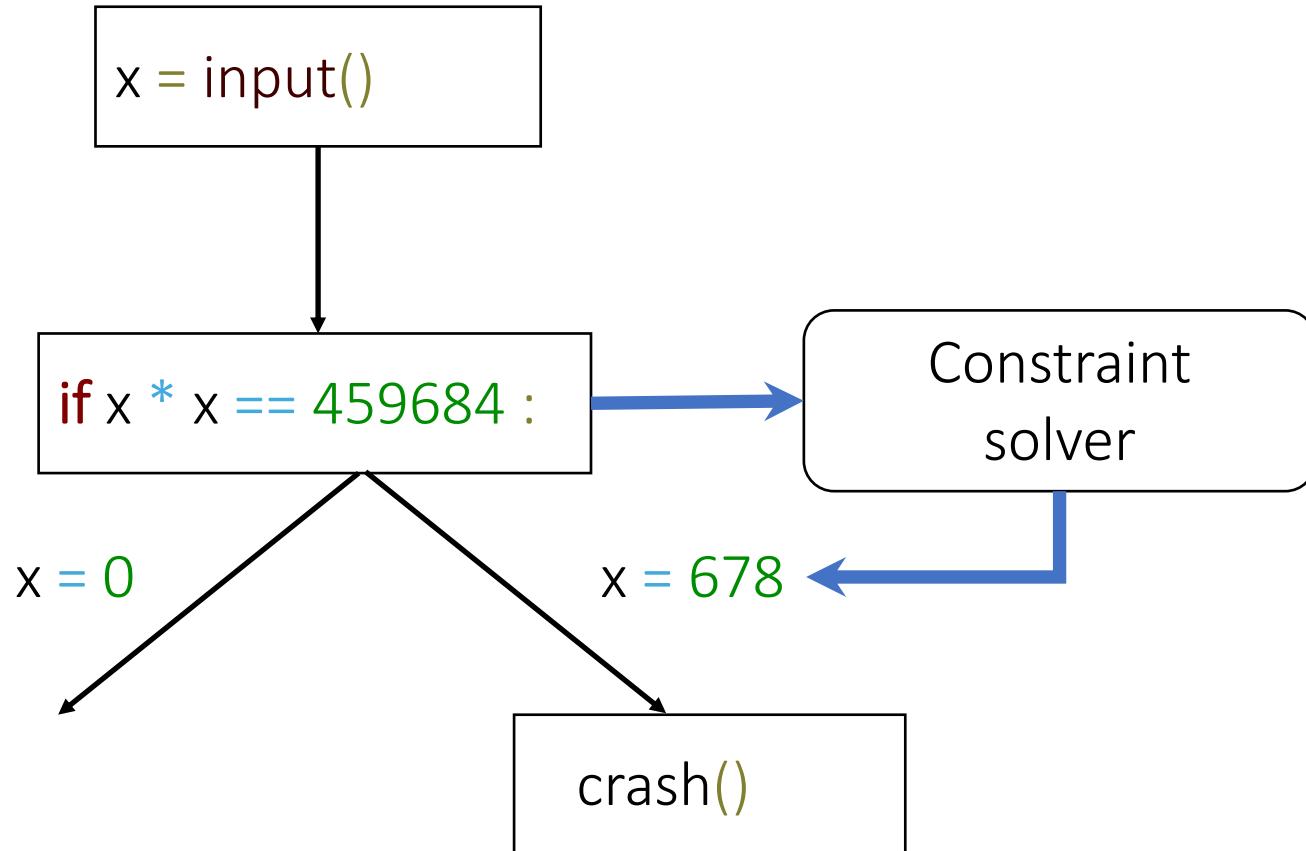
```
// 459684 == 6782
if x * x == 459684 :
    crash()
```



# Concolic execution can help fuzzing by finding hard-to-find test cases

```
x = int(input())
```

```
// 459684 == 6782
if x * x == 459684 :
    crash()
```



Hybrid fuzzing has achieved great success in small-scale study (DARPA Cyber Grand Challenge)



# Hybrid fuzzing has achieved great success in small-scale study (DARPA Cyber Grand Challenge)

- Organized by DARPA in 2017
- Build a system to find bugs, exploit and patch on binaries
- Over 100 teams → 7 teams were qualified (include our team)
- Almost \$4 million for prize money



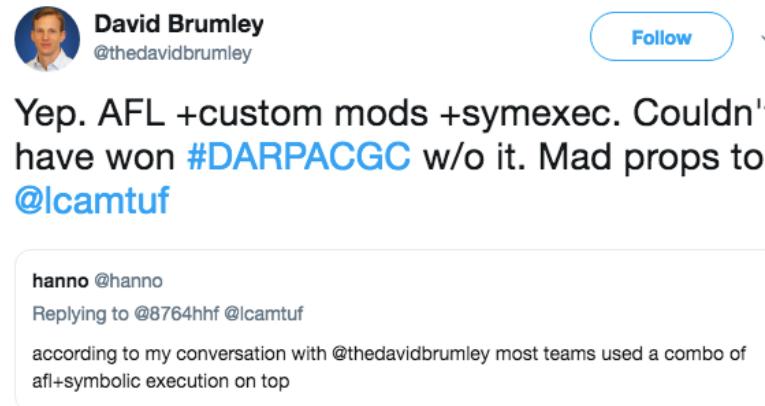
# Hybrid fuzzing has achieved great success in small-scale study (DARPA Cyber Grand Challenge)

- Organized by DARPA in 2017
- Build a system to find bugs, exploit and patch on binaries
- Over 100 teams → 7 teams were qualified (include our team)
- Almost \$4 million for prize money
- *Small binaries: a few KB*



# Hybrid fuzzing has achieved great success in small-scale study (DARPA Cyber Grand Challenge)

- The winner from CMU used hybrid fuzzing



**David Brumley**  
@thedavidbrumley

Follow

Yep. AFL +custom mods +symexec. Couldn't have won **#DARPACGC** w/o it. Mad props to **@lcamtuf**

**hanno** @hanno  
Replies to @8764hhf @lcamtuf  
according to my conversation with @thedavidbrumley most teams used a combo of afl+symbolic execution on top

- Shellphish open-sourced their tool, Driller
  - Won 3<sup>rd</sup> place in CGC competition
  - Found 6 new crashes: cannot be found by fuzzing or concolic execution

But, hybrid fuzzing fails to scale real-world applications

Cannot find ANY bug in  
real-world software  
using Driller!



Current concolic executors suffer several problems to be used in hybrid fuzzing

Generating  
constraints is too slow

Current concolic executors suffer several problems to be used in hybrid fuzzing

Generating  
constraints is too slow

Not effective in  
generating test cases

Current concolic executors suffer several problems to be used in hybrid fuzzing

Generating  
constraints is too slow

Not effective in  
generating test cases

Symbolic emulation is well-known to be much slower than concrete execution

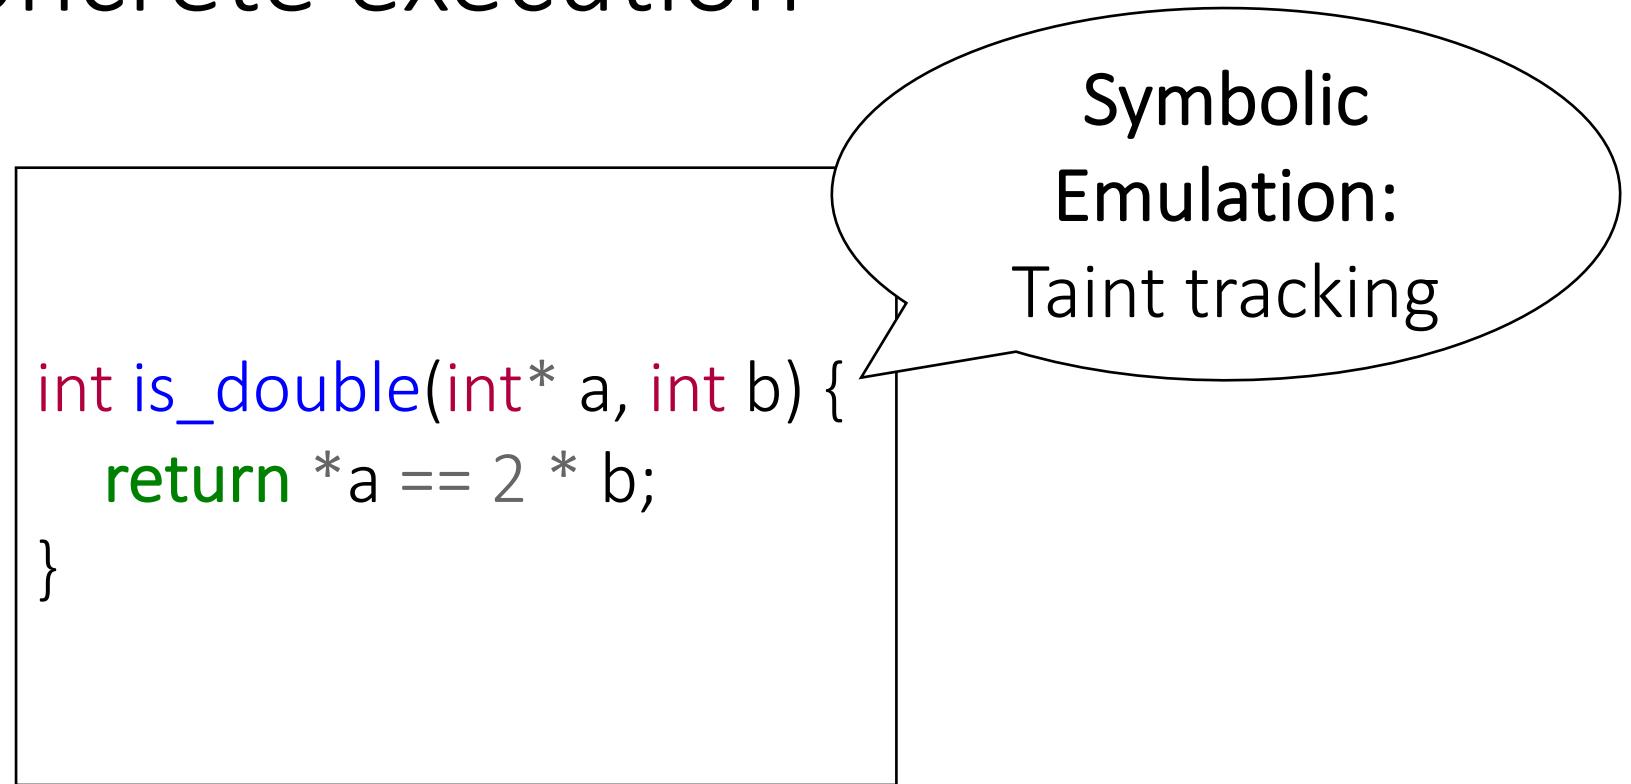
```
int is_double(int* a, int b) {  
    return *a == 2 * b;  
}
```

Symbolic emulation is well-known to be much slower than concrete execution

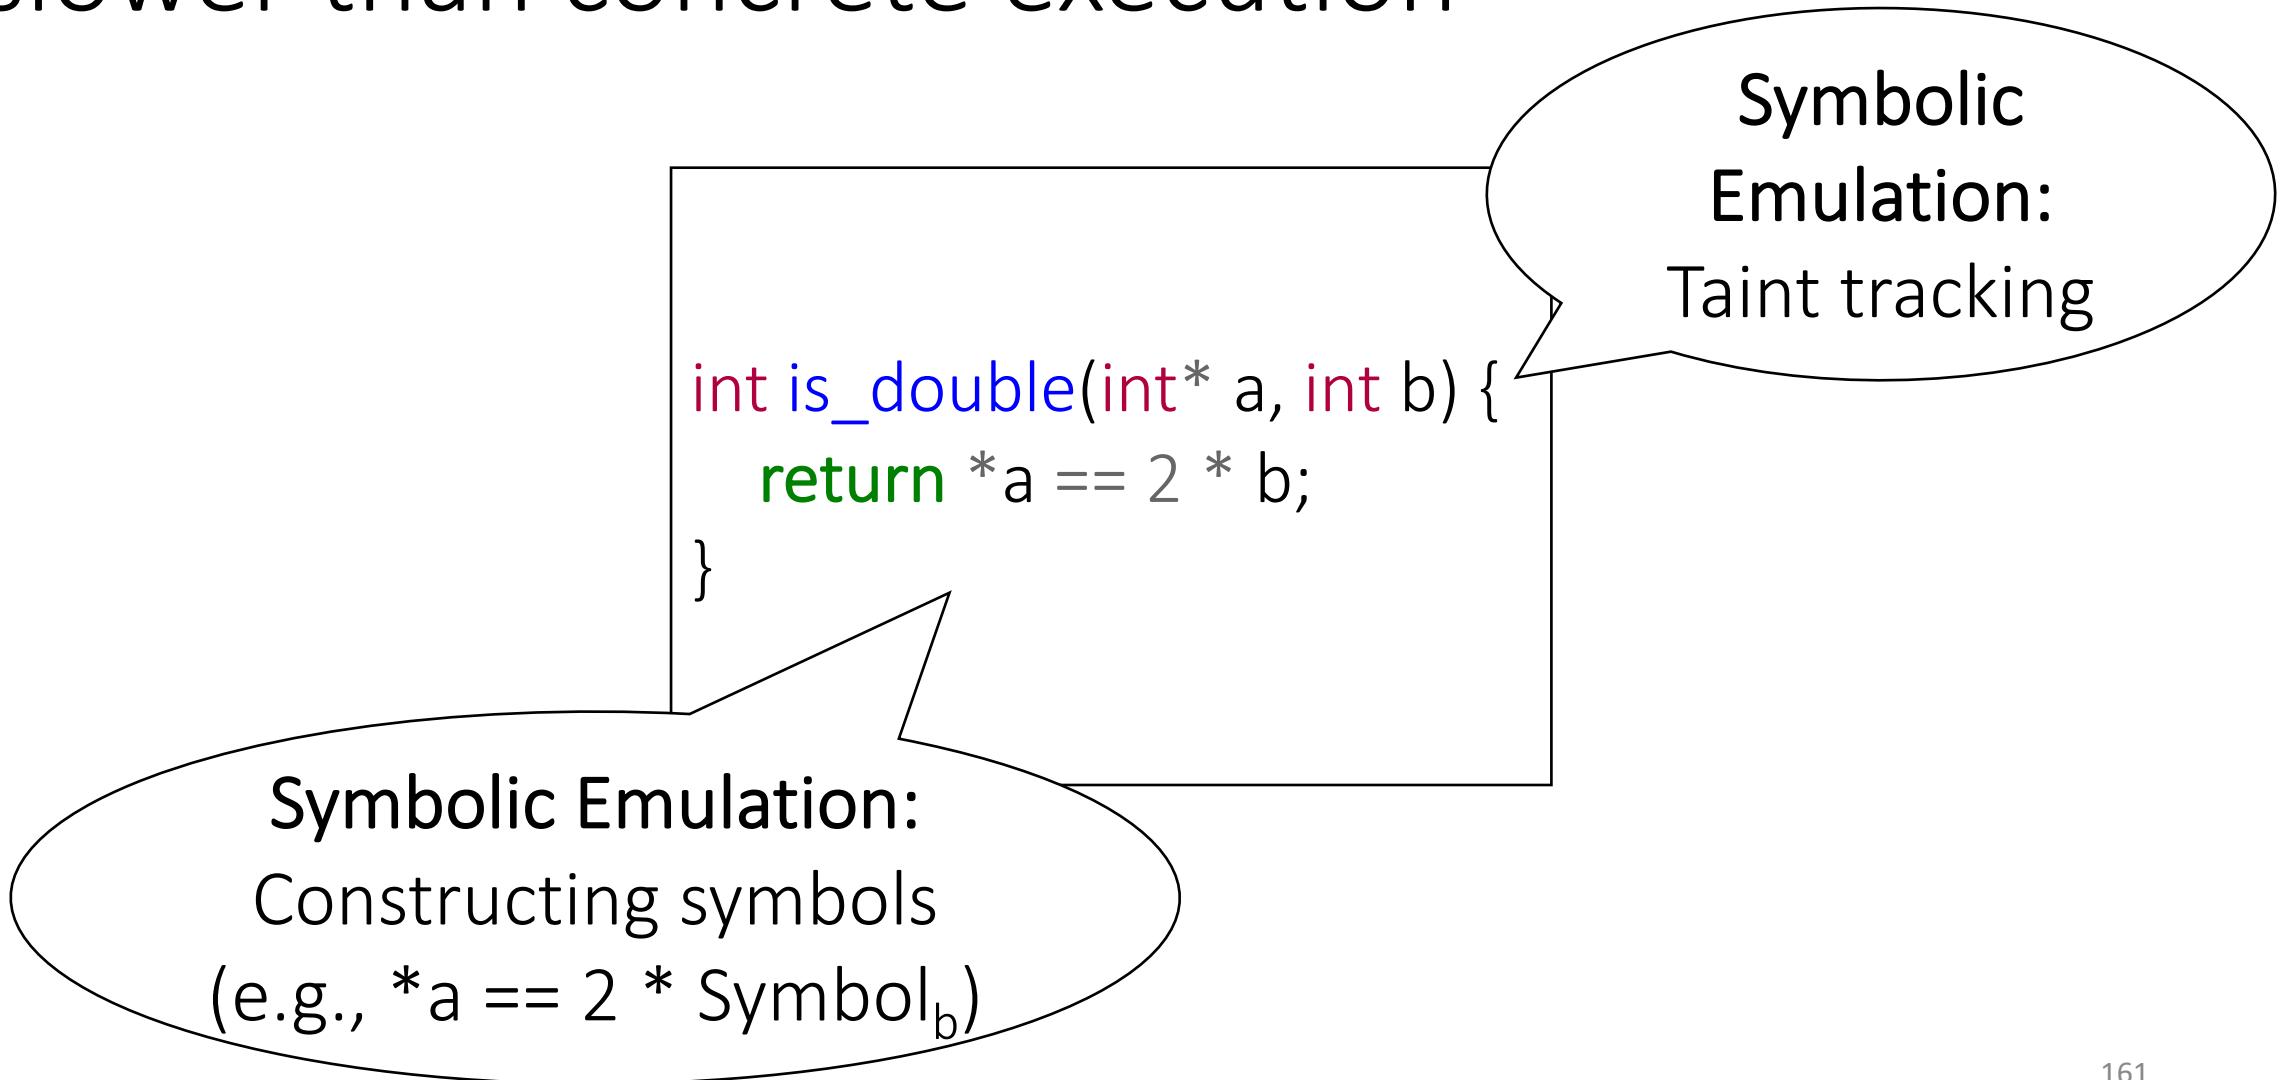
Concrete  
Execution:  
Just Execute!

```
int is_double(int* a, int b) {  
    return *a == 2 * b;  
}
```

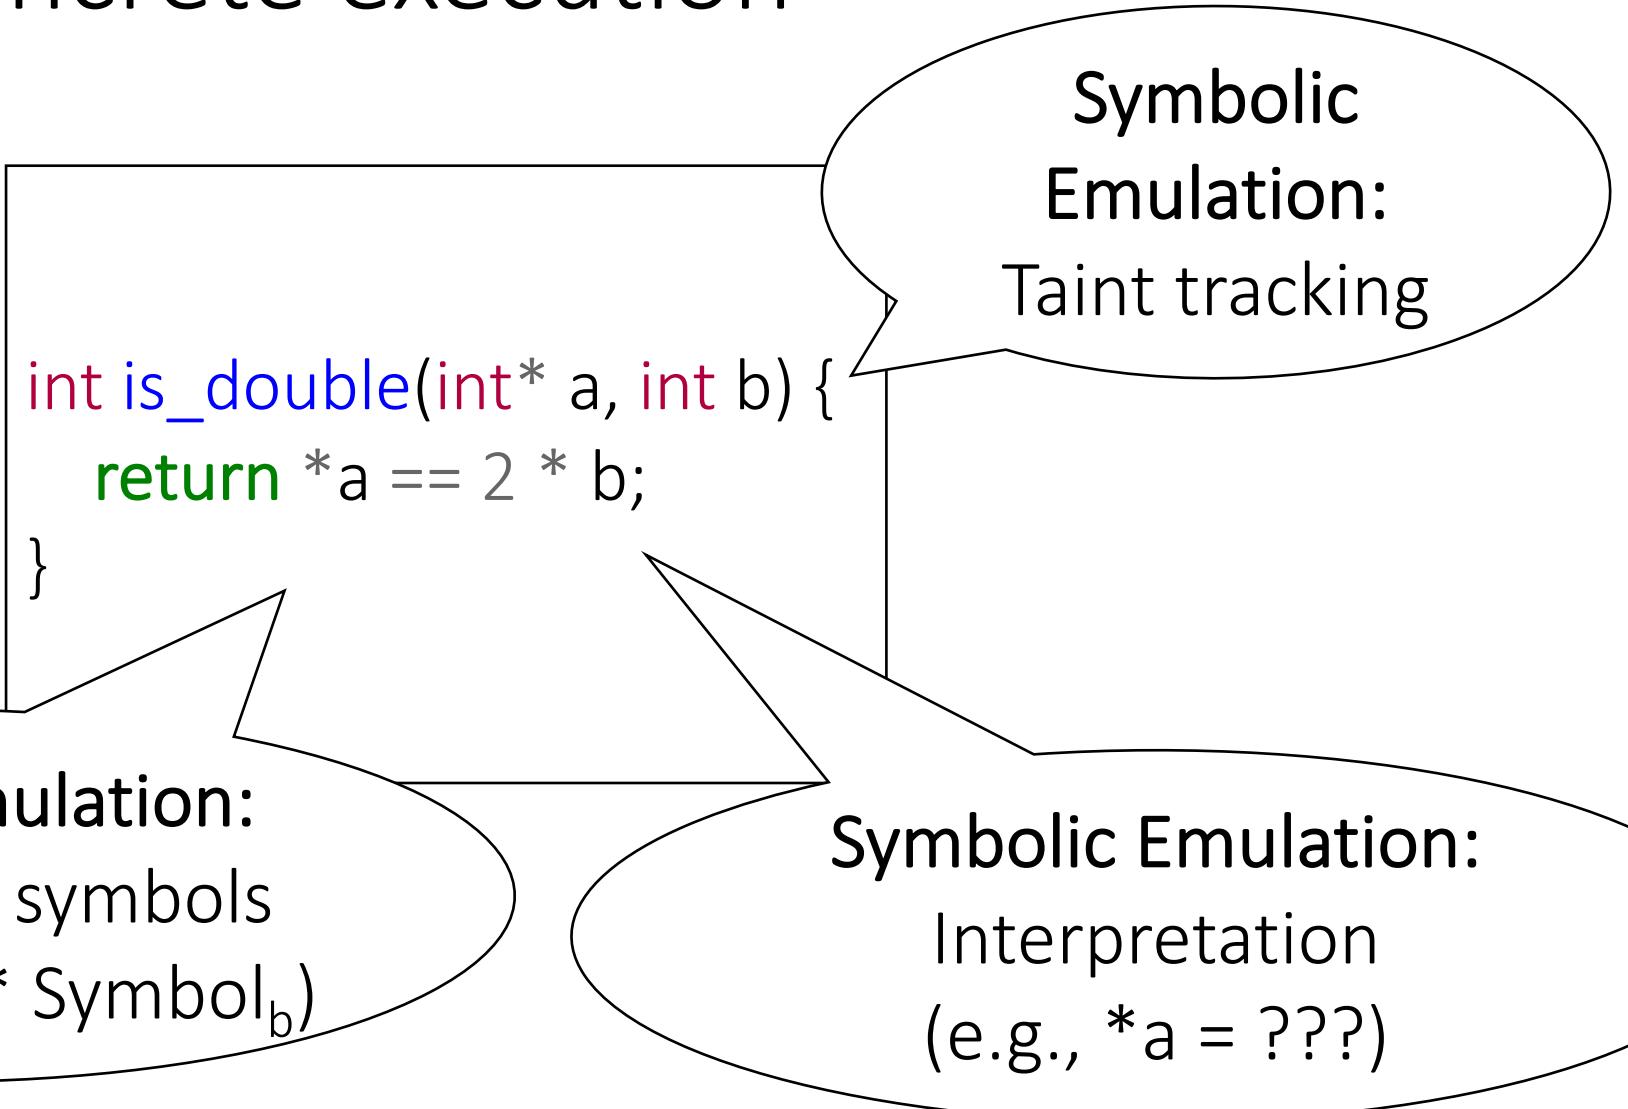
Symbolic emulation is well-known to be much slower than concrete execution



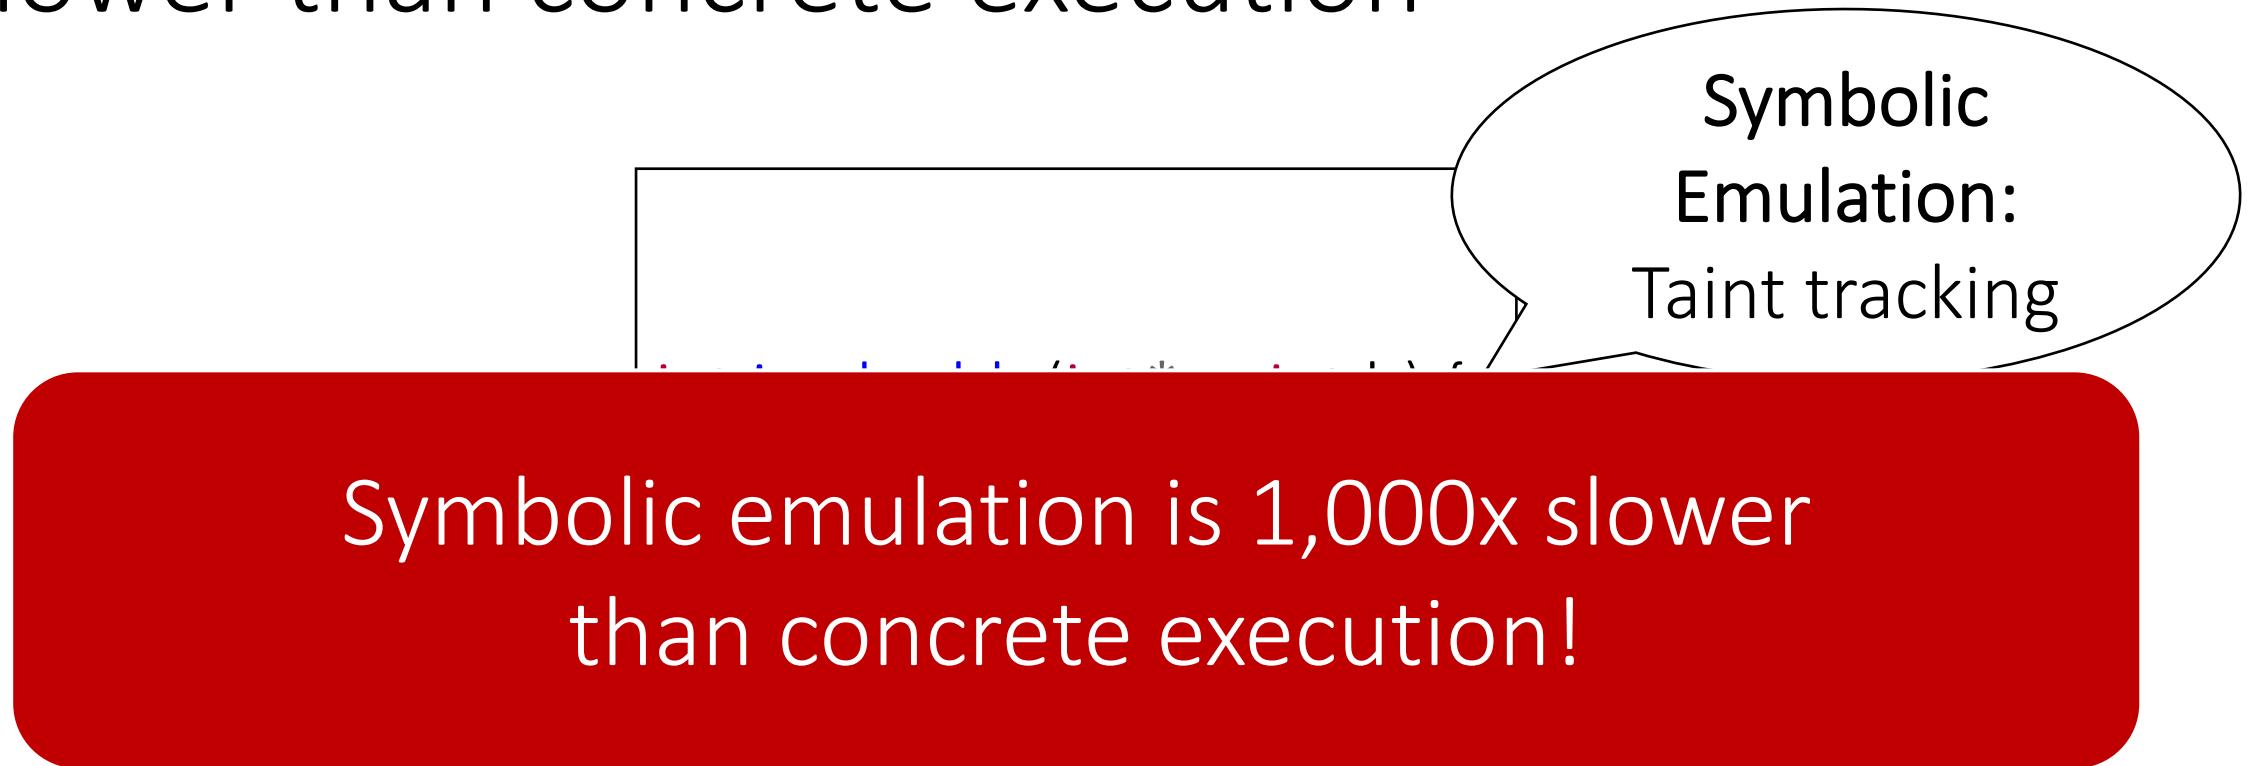
Symbolic emulation is well-known to be much slower than concrete execution



Symbolic emulation is well-known to be much slower than concrete execution



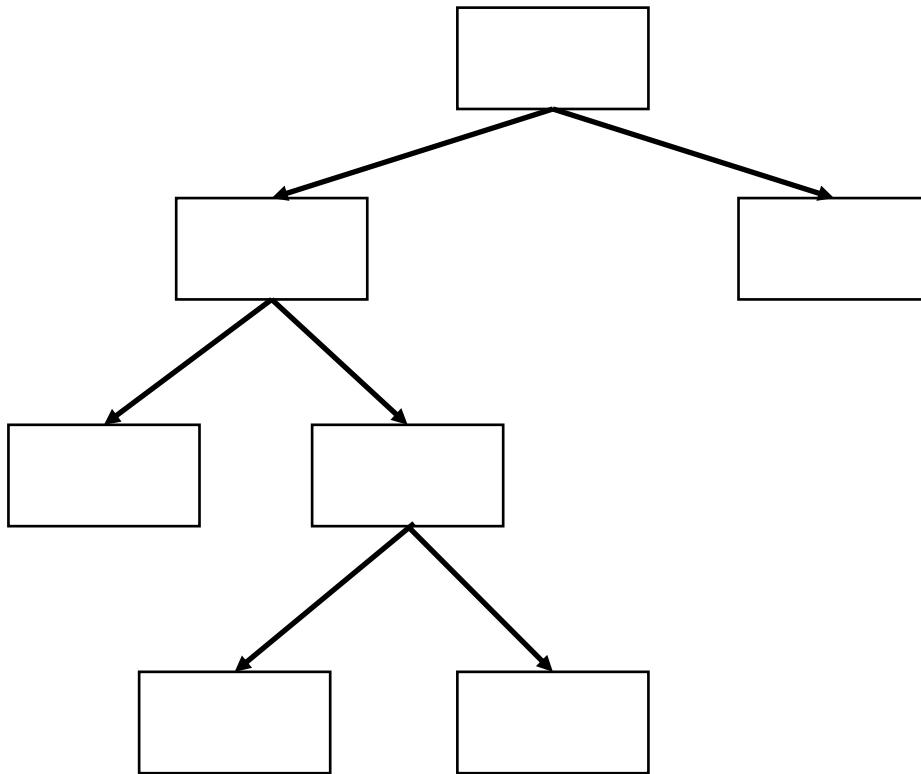
Symbolic emulation is well-known to be much slower than concrete execution



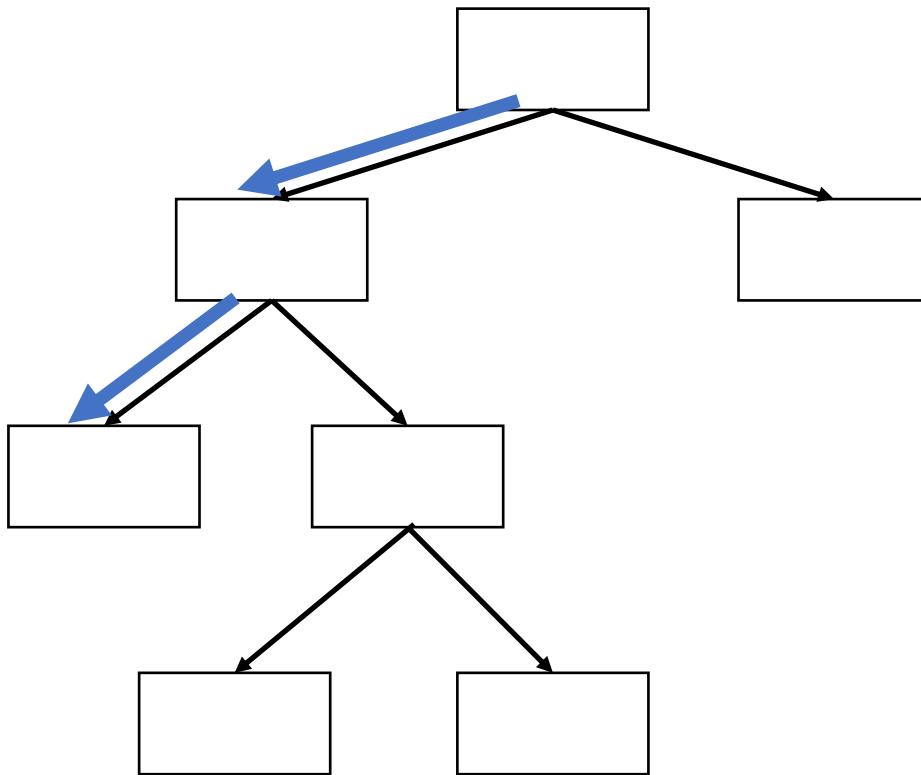
Symbolic Emulation:  
Constructing symbols  
(e.g.,  $*a == 2 * \text{Symbol}_b$ )

Symbolic Emulation:  
Interpretation  
(e.g.,  $*a = ???$ )

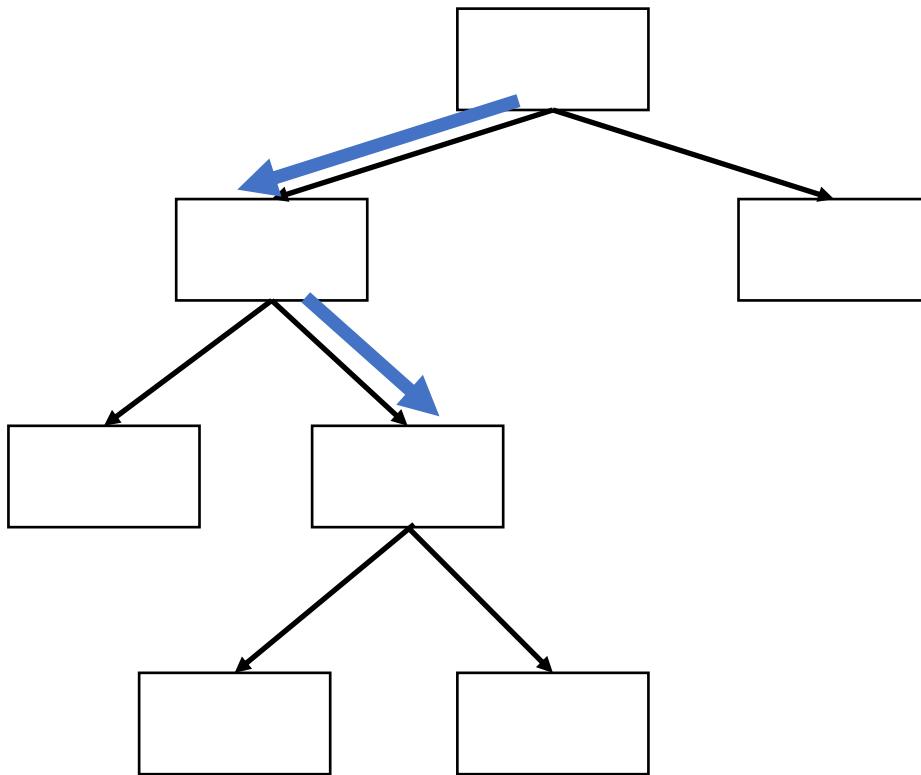
State forking can be used to solve this problem



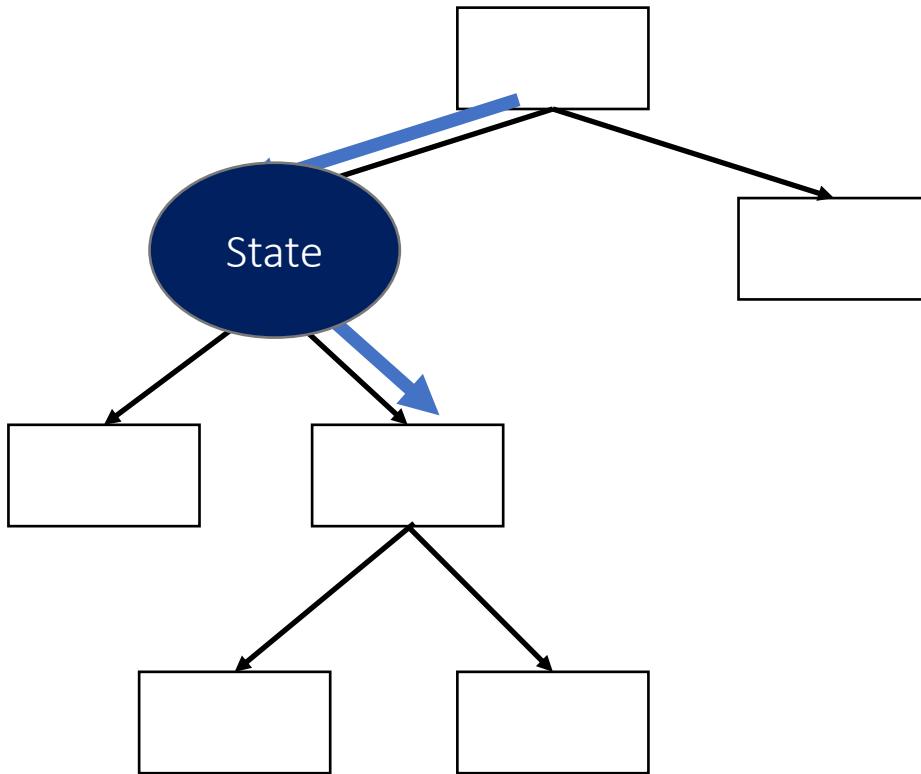
State forking can be used to solve this problem



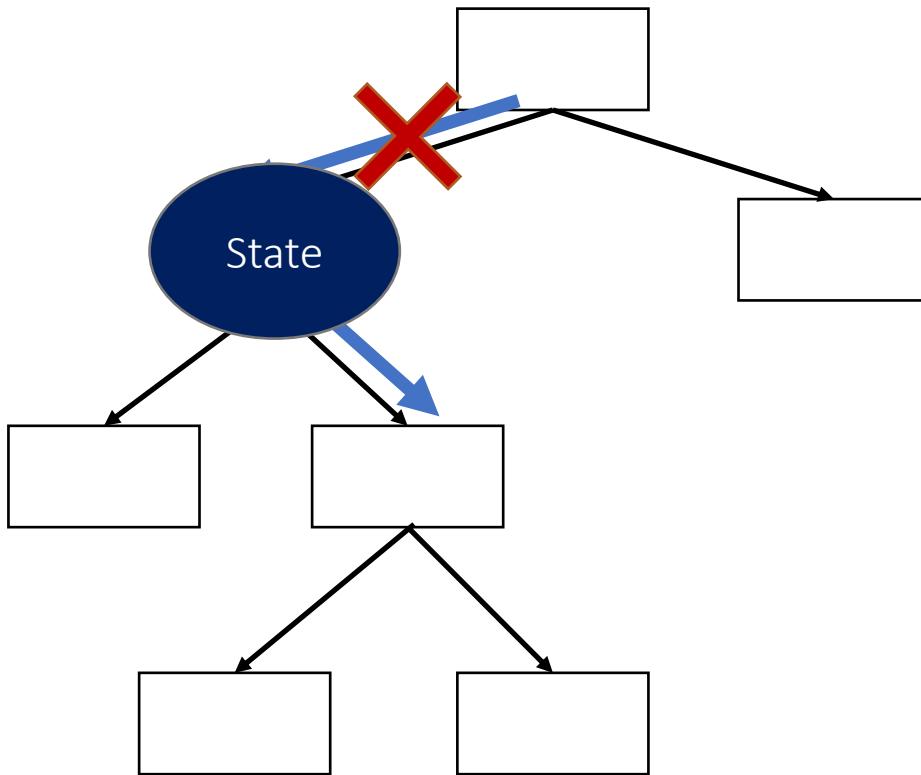
State forking can be used to solve this problem



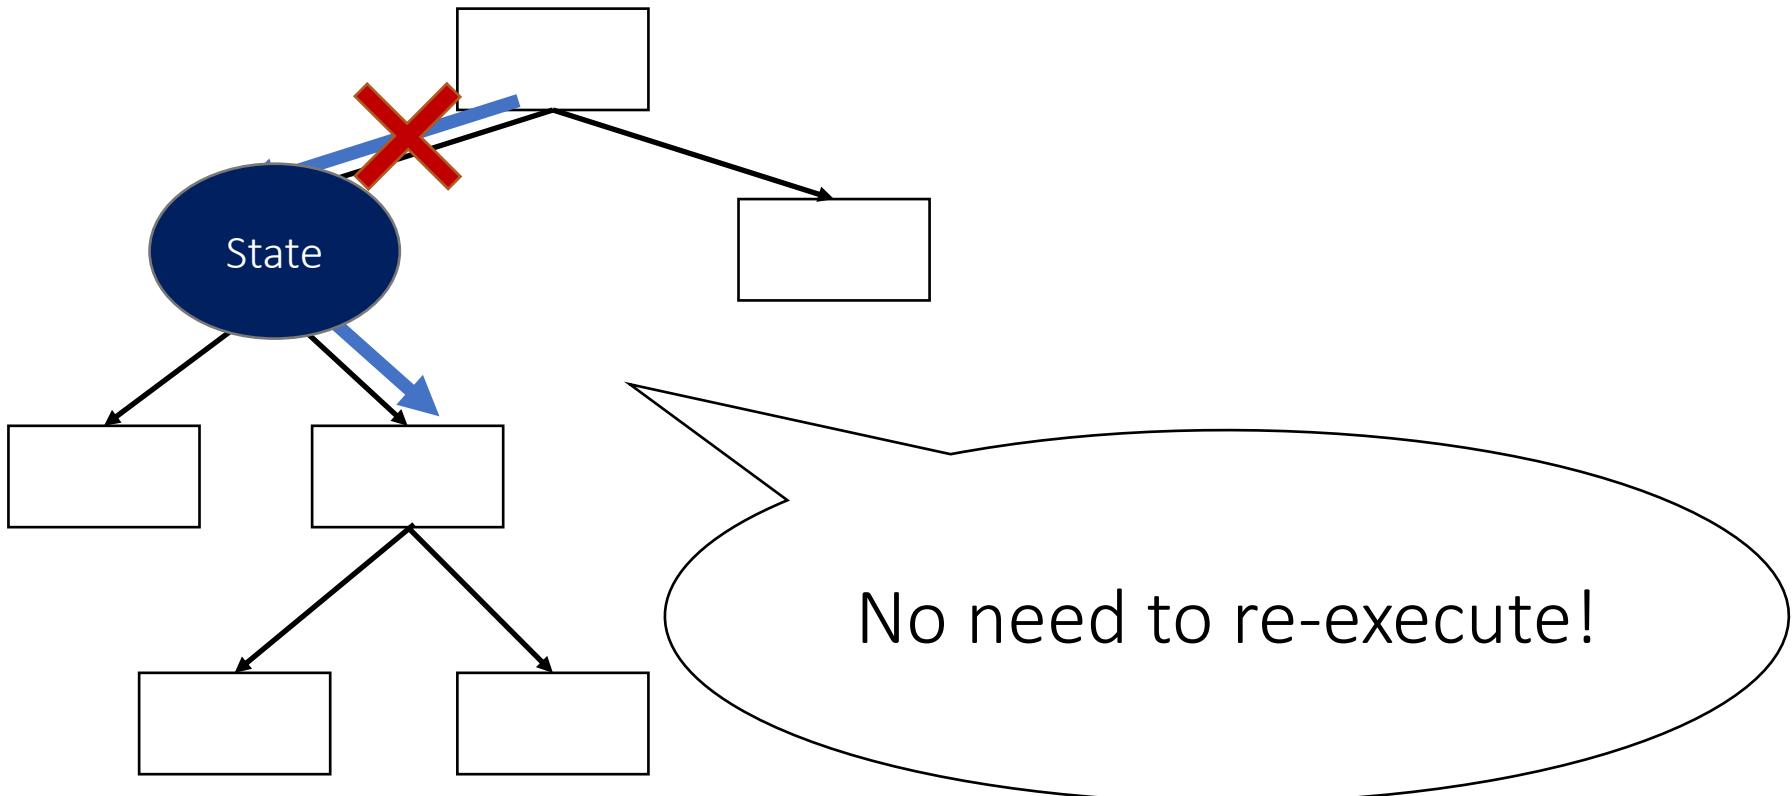
# State forking can be used to solve this problem



# State forking can be used to solve this problem



# State forking can be used to solve this problem



State forking is limited in hybrid fuzzing

# State forking is limited in hybrid fuzzing

- # of states is enormous in a complex real-world program  
=> Large performance overhead
- Hybrid fuzzing explores paths randomly following fuzzing  
=> Low reusability

# State forking is limited in hybrid fuzzing

- # of states is enormous in a complex real-world program  
=> Large performance overhead
- Hybrid fuzzing explores paths randomly following fuzzing  
=> Low reusability

State forking cannot help slow symbolic emulation  
in hybrid fuzzing!

Current concolic executors have several problems to be used in hybrid fuzzing

Generating  
constraints is too slow

Not effective in  
generating test cases

# Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```

# Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```



Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```



Analyze every routine for completeness!

# Soundness of concolic execution incurs unnecessary re-execution

```
1 // 'x' is symbolic and 'x' == 0 in a given input
2 int soundess(int x) {
3     if (x == 0)
4         do_something();
5
6     if (x * x == 1234 * 1234)
7         crash();
8 }
```

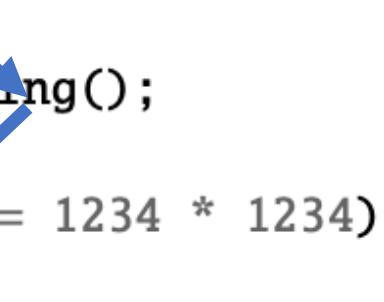
# Soundness of concolic execution incurs unnecessary re-execution

```
1 // 'x' is symbolic and 'x' == 0 in a given input
2 int soundess(int x) {
3     if (x == 0)
4         do_something();
5
6     if (x * x == 1234 * 1234)
7         crash();
8 }
```



# Soundness of concolic execution incurs unnecessary re-execution

```
1 // 'x' is symbolic and 'x' == 0 in a given input
2 int soundess(int x) {
3     if (x == 0)
4         do_something();
5
6     if (x * x == 1234 * 1234)
7         crash();
8 }
```



Unsatisfiable!

# Soundness of concolic execution incurs unnecessary re-execution

```
1 // 'x' is symbolic and 'x' == 0 in a given input
2 int soundess(int x) {
3     if (x == 0)
4         do_something();
5
6     if (x * x == 1234 * 1234)
7         crash();
8 }
```



Unsatisfiable!

Don't make possibly incorrect test cases for soundness!

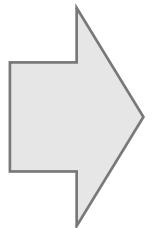
# Our approach

Generating  
constraints is too slow

Not effective in  
generating test cases

# Our approach

Generating  
constraints is too slow

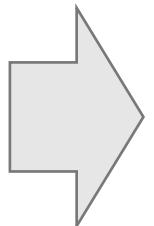


Systematic approach  
for fast symbolic  
emulation

Not effective in  
generating test cases

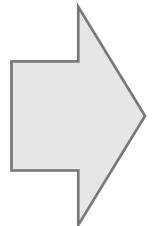
# Our approach

Generating  
constraints is too slow



Systematic approach  
for fast symbolic  
emulation

Not effective in  
generating test cases



New heuristics for  
hybrid fuzzing

# Our approach: QSYM

QSYM

Systematic approach  
for fast symbolic  
emulation

Instruction-level  
concolic execution  
(For binary)

New heuristics for  
hybrid fuzzing

Optimistic solving and  
basic block pruning

# Our approach: Hybridra

Hybridra

Systematic approach  
for fast symbolic  
emulation

Compilation-based  
concolic execution  
(For source code)

New heuristics for  
hybrid fuzzing

Staged reduction  
+ Heuristics from QSYM

# Related work: Whitebox fuzzing

- Goal
  - Hybrid fuzzing: Make a test case for fuzzing
  - Whitebox fuzzing: Explore a program state solely
- Exploration
  - Hybrid fuzzing: Random
  - Whitebox fuzzing: Systematic
- Strategy: Hybrid fuzzing's can be more aggressive thanks to coverage-guided fuzzing (e.g., optimistic solving)

# Today's talk

QSYM: A Binary-level  
Concolic Execution Engine  
for Hybrid fuzzing

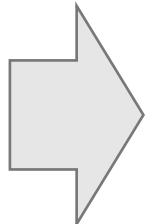
- Binary
- User applications

Hybridra: A Hybrid Fuzzer  
for Kernel File Systems

- Source code
- File systems

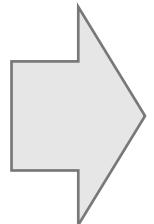
Our system, QSYM, addresses these issues by introducing several key ideas

Generating constraints is too slow



Instruction-level  
concolic execution  
(For binary)

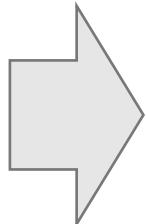
Not effective in  
generating test cases



Optimistic solving and  
basic block pruning

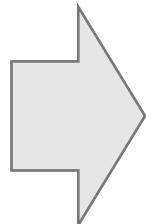
Our system, QSYM, addresses these issues by introducing several key ideas

Generating constraints is too slow



Instruction-level  
concolic execution  
(For binary)

Not effective in  
generating test cases



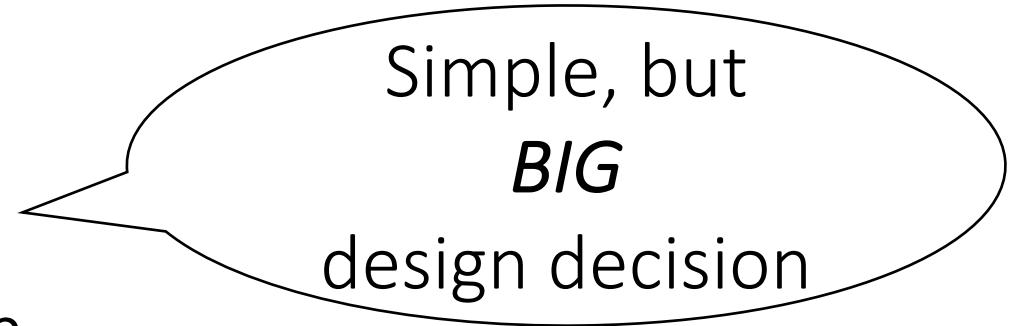
Optimistic solving and  
basic block pruning

QSYM has made several design decisions for improving performance

- Discarding intermediate layer
- Instruction-level symbolic execution

# QSYM has made several design decisions for improving performance

- Discarding intermediate layer
- Instruction-level symbolic execution



Simple, but  
*BIG*  
design decision

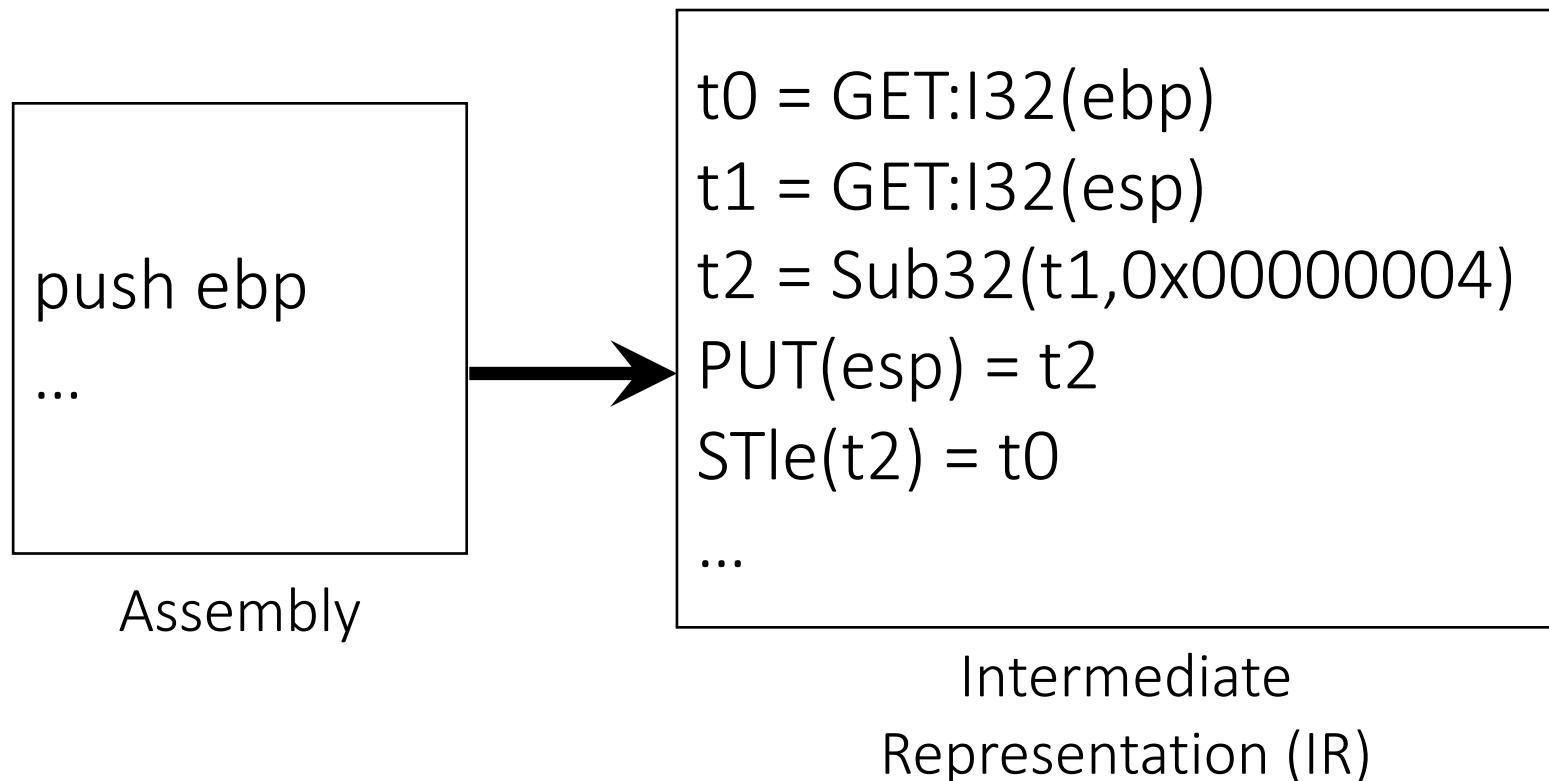
# Hybrid fuzzing in a closer look

```
push ebp
```

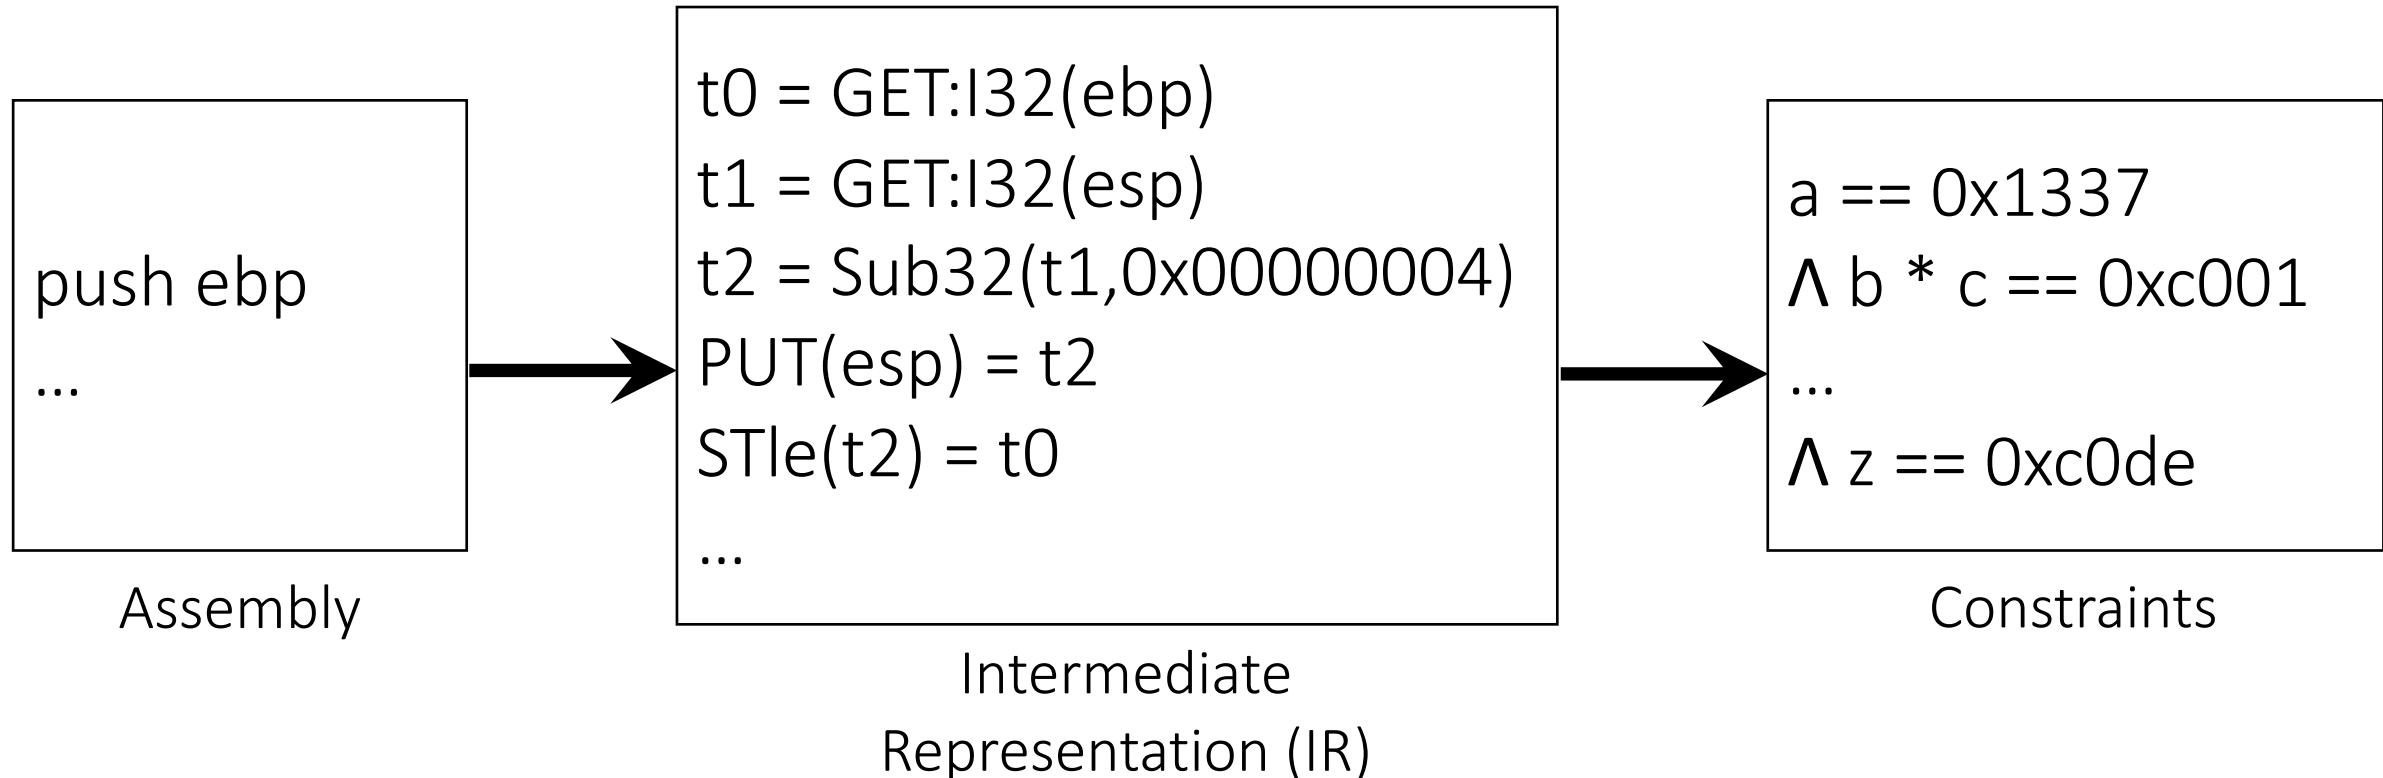
```
...
```

Assembly

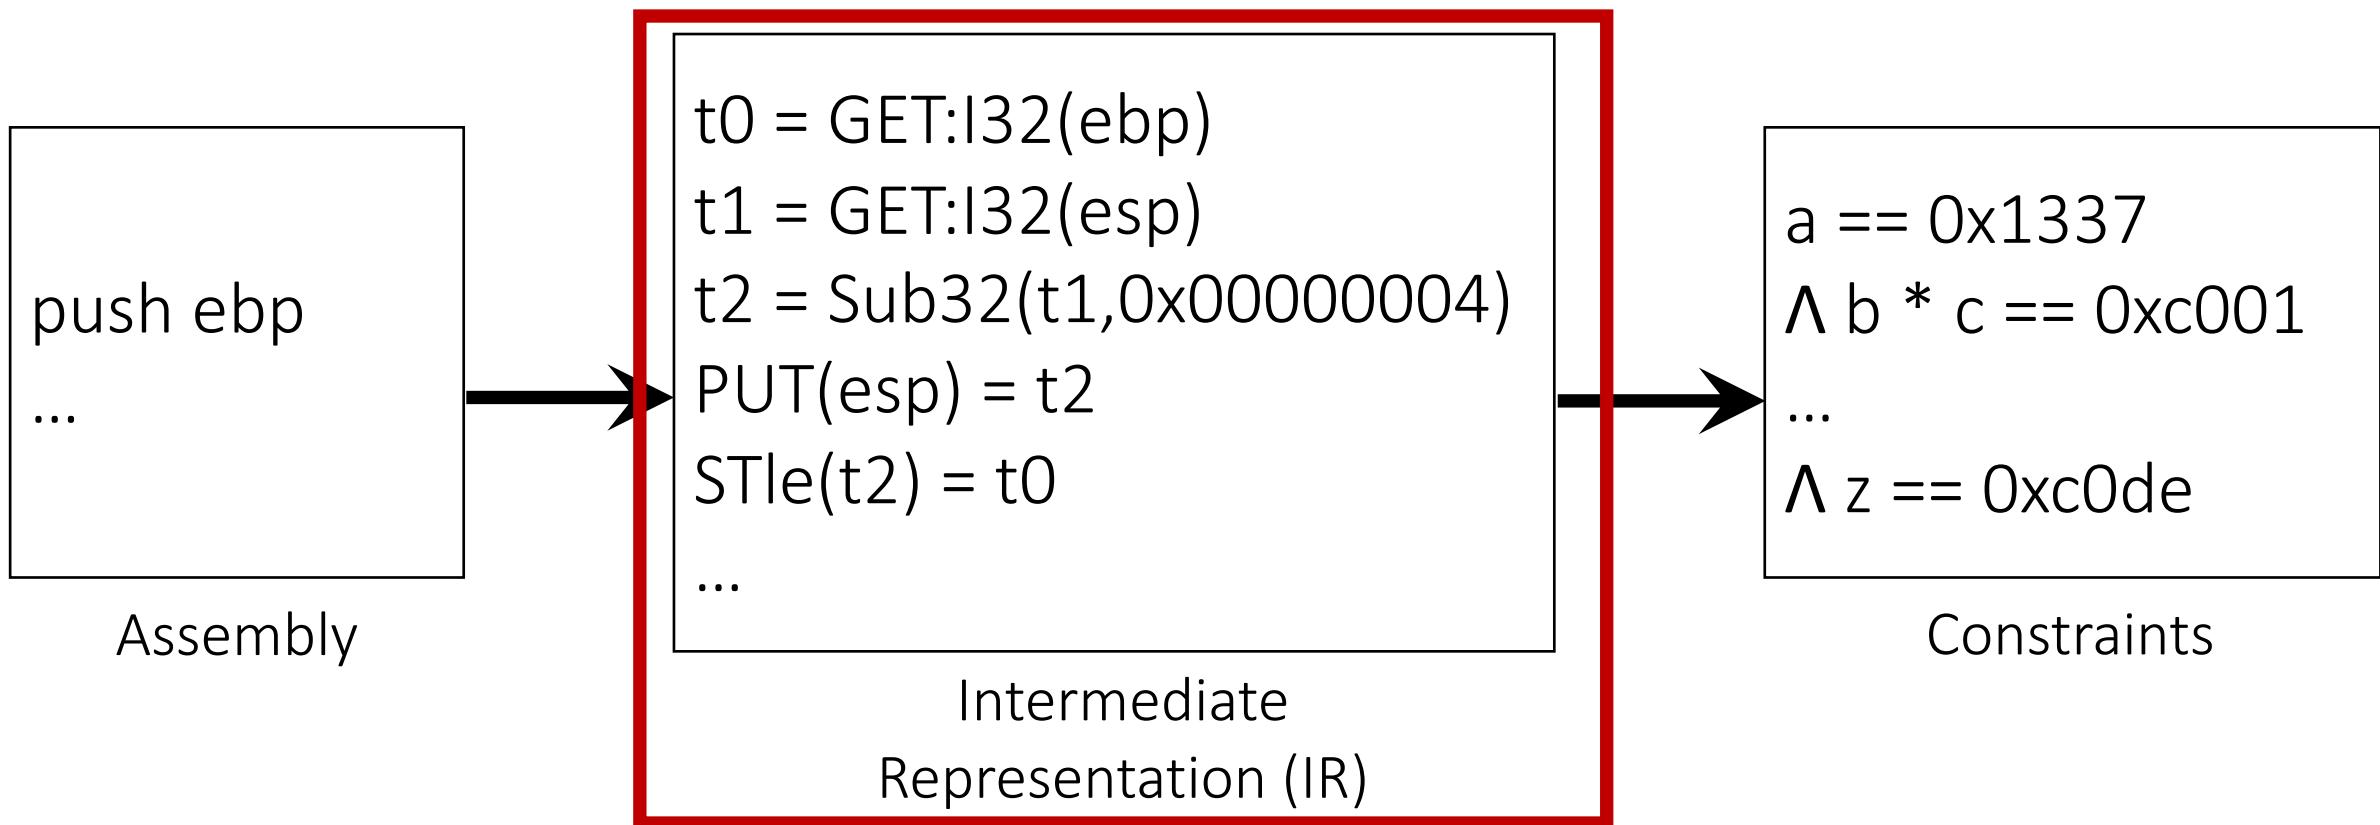
# Hybrid fuzzing in a closer look



# Hybrid fuzzing in a closer look

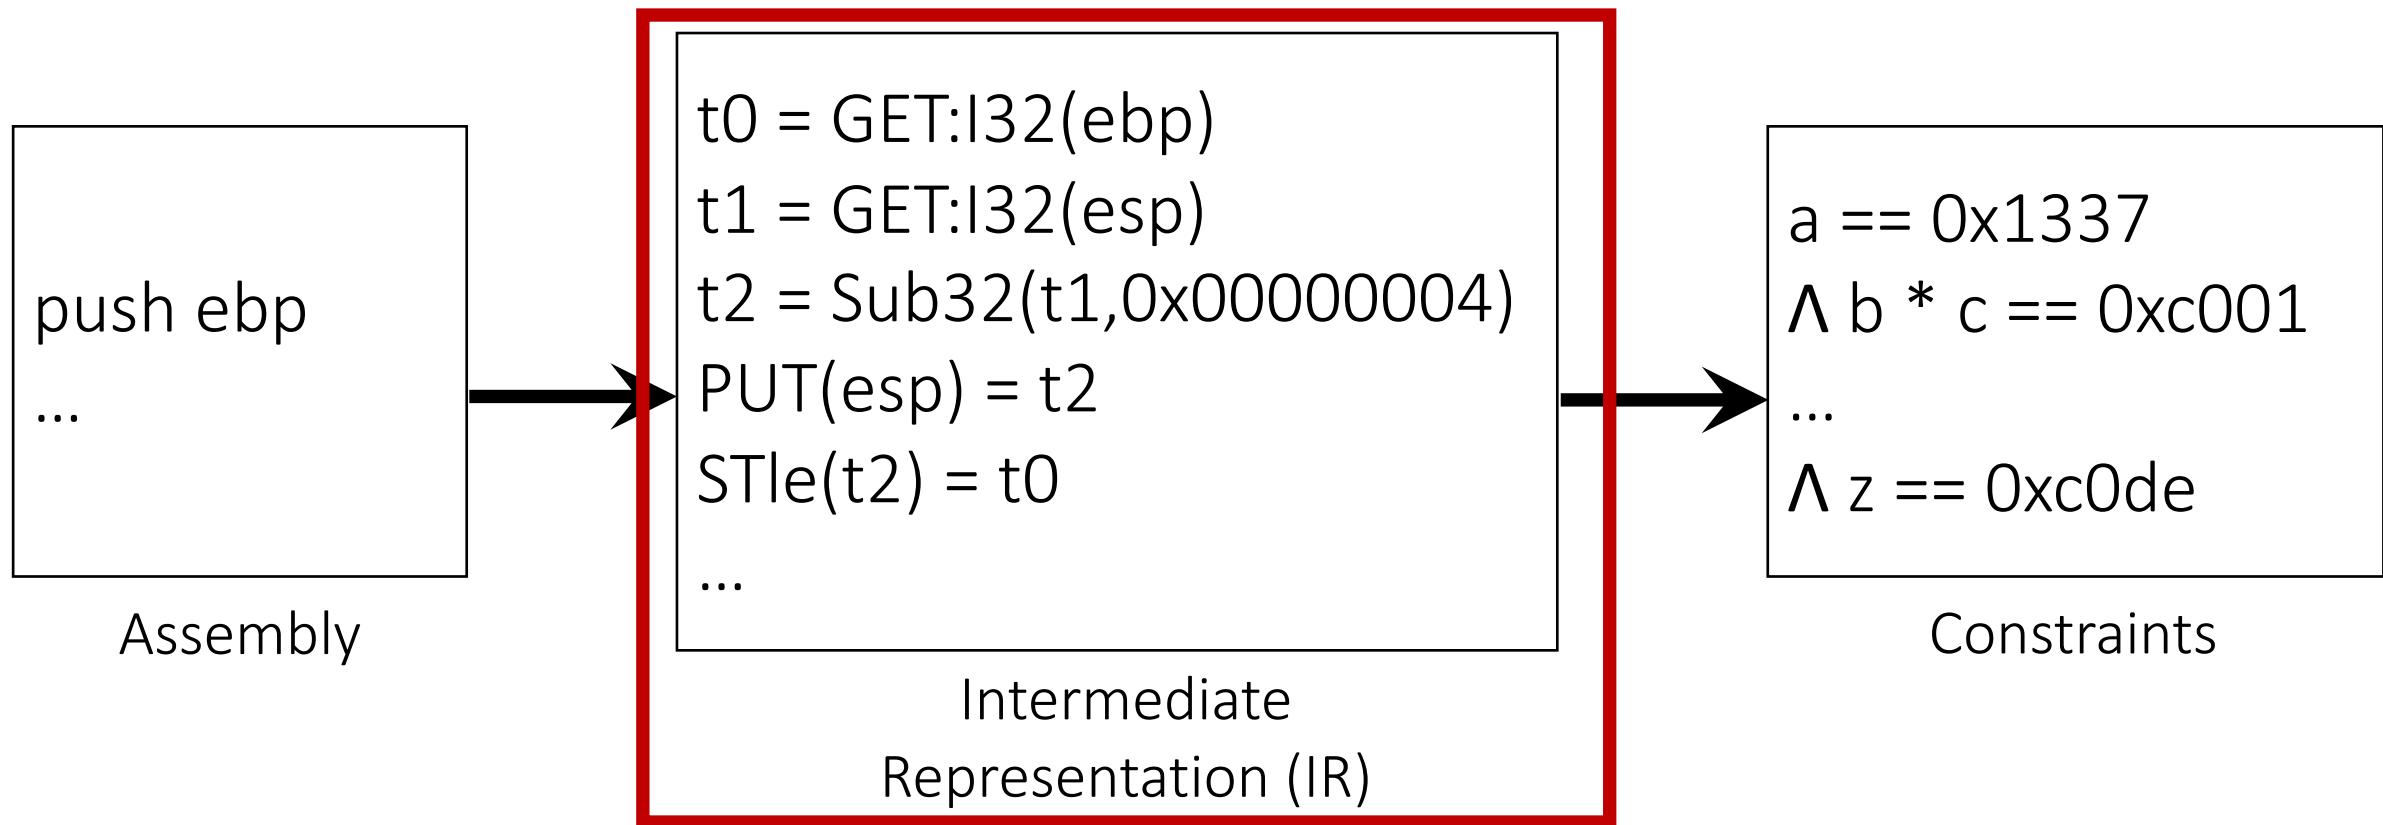


# Hybrid fuzzing in a closer look



Good: Simplifying implementations  
e.g., 981 in x86 vs 115 in VEX

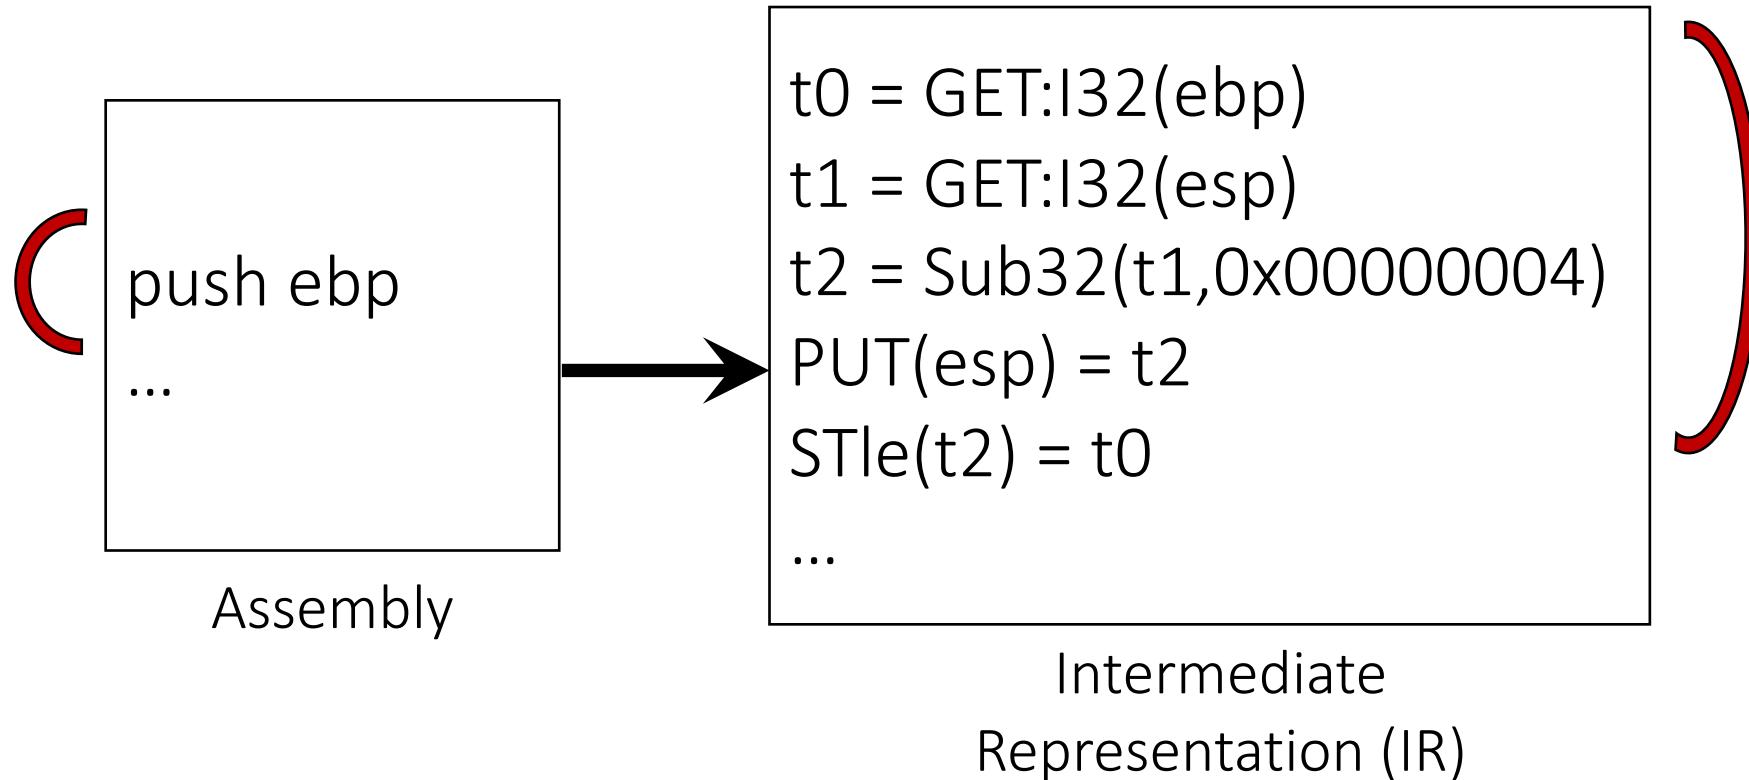
# Hybrid fuzzing in a closer look



Good: Simplifying implementations  
e.g., 981 in x86 vs 115 in VEX

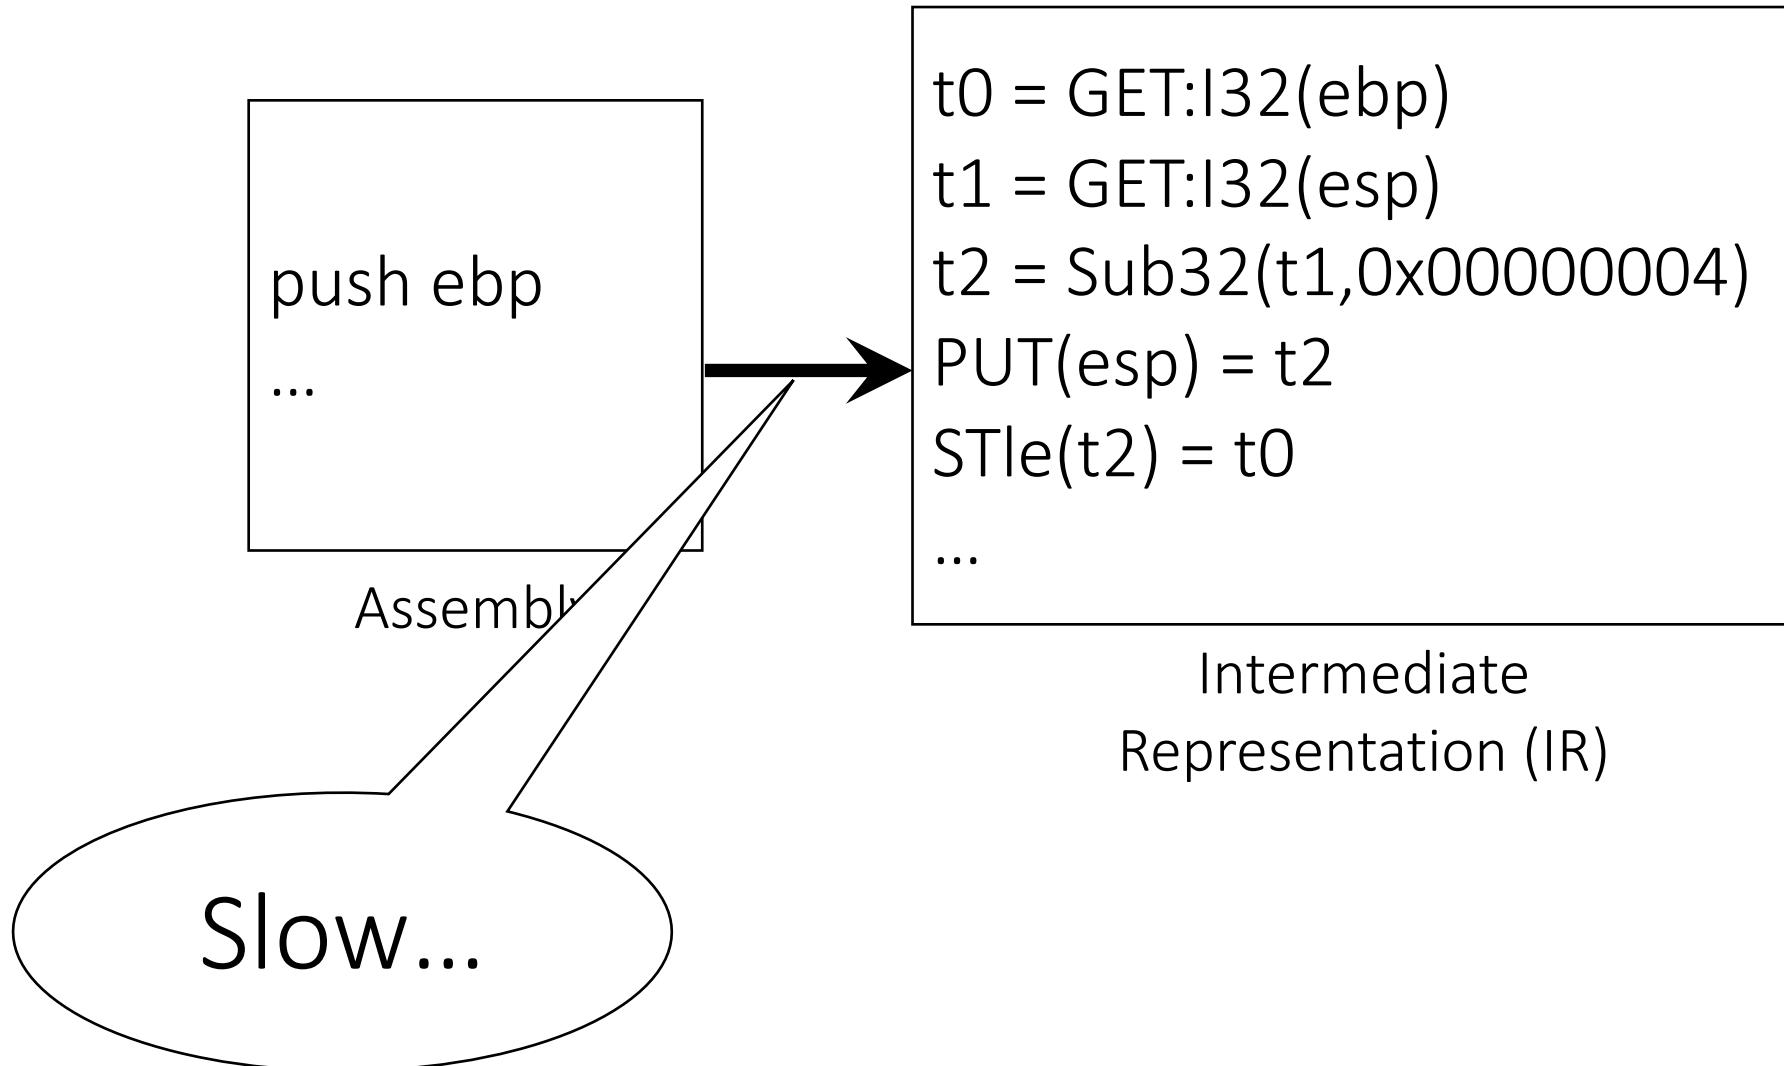
Bad: Performance bottleneck

# Problems of IR: The number of instructions increase

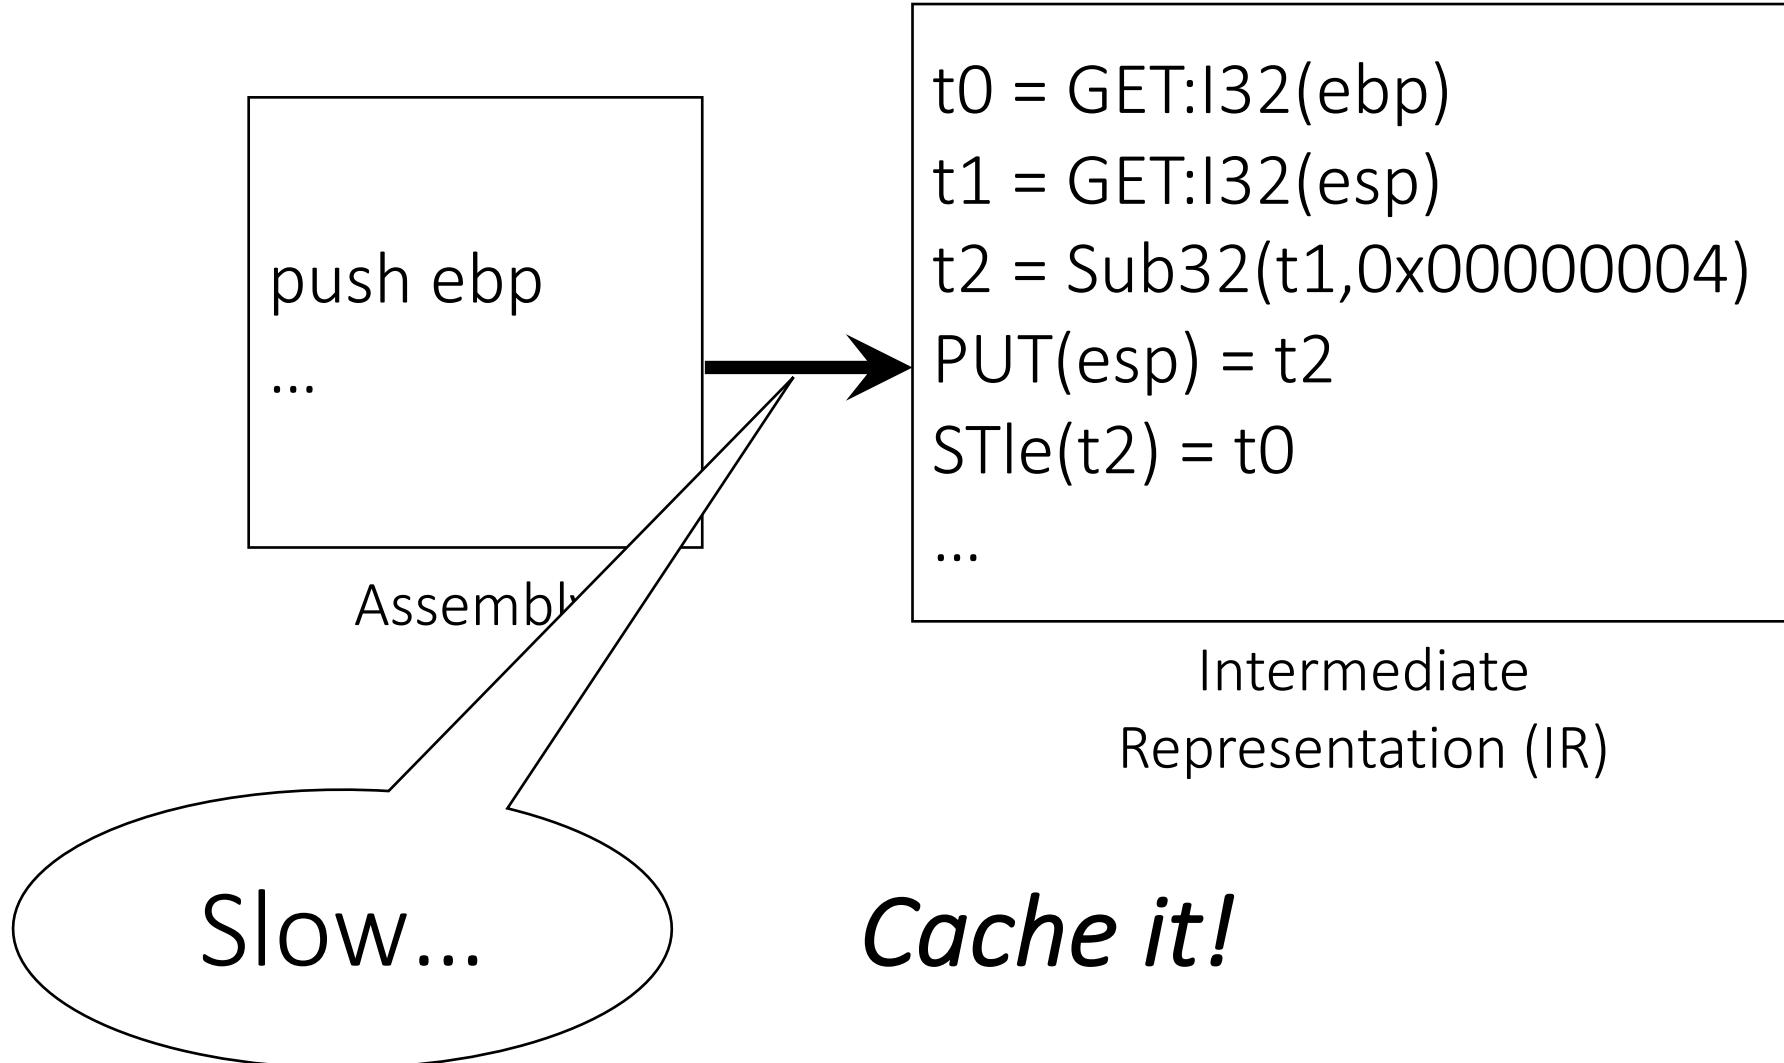


**4.96x increase  
on average!**

# Problems of IR: Slow transformation speed



# Problems of IR: Slow transformation speed



# Side effects of caching: Basic-block granularity

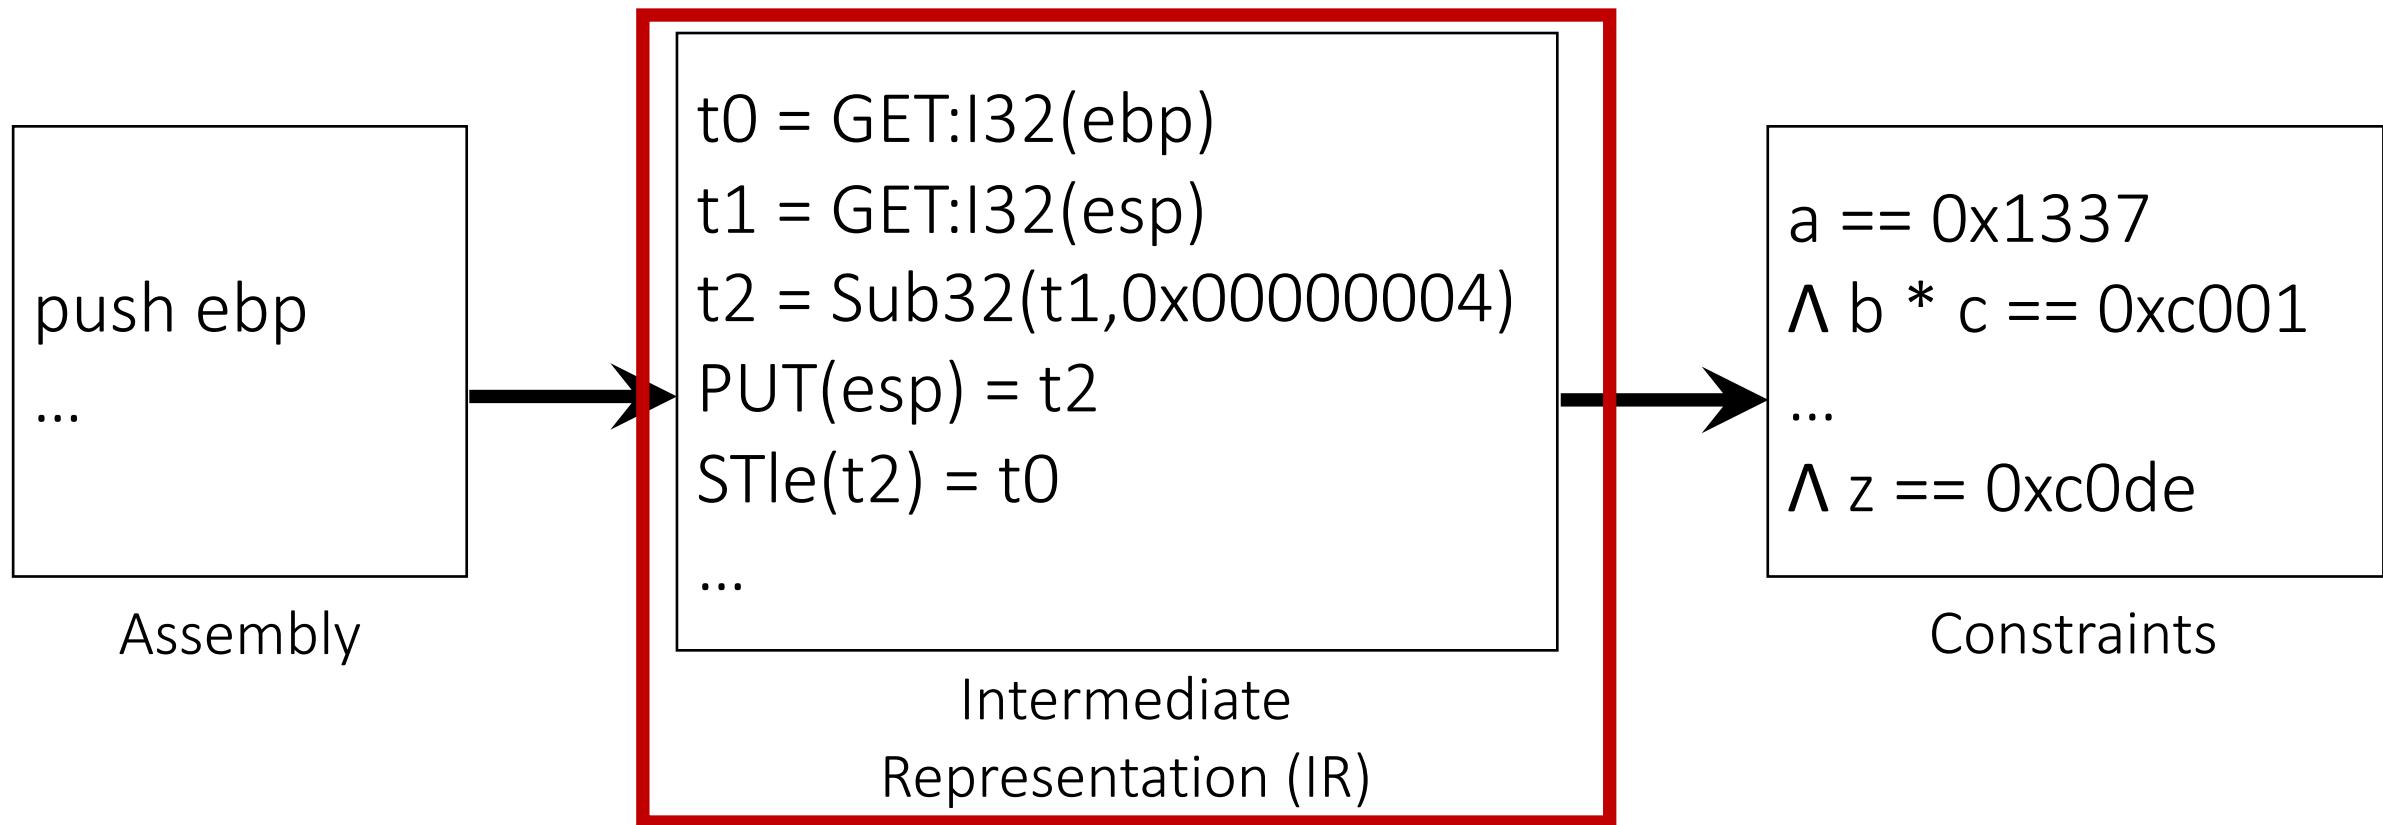
# Side effects of caching: Basic-block granularity

- Cache lookup is also slow
- Use basic-block granularity for caching
  - i.e., transform a basic block into IR and cache
- Unfortunately, **30%** of instructions in a basic block are symbolic  
→ 70% of instructions are executed without need

# Side effects of caching: Basic-block granularity

- Cache lookup is also slow
- Use basic-block granularity for caching
  - i.e., transform a basic block into IR and cache
- Unfortunately, **30%** of instructions in a basic block are symbolic  
→ 70% of instructions are executed without need

# How to solve this challenge?

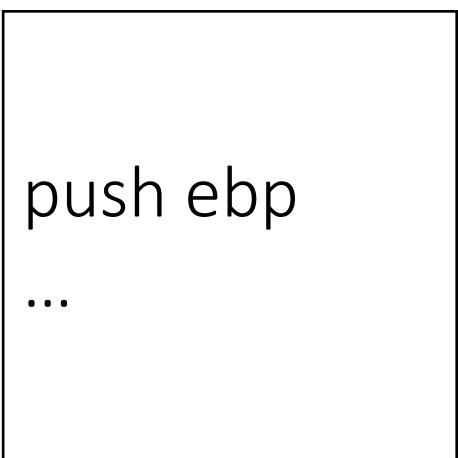


Good: Simplifying implementations  
e.g., 981 in x86 vs 115 in VEX

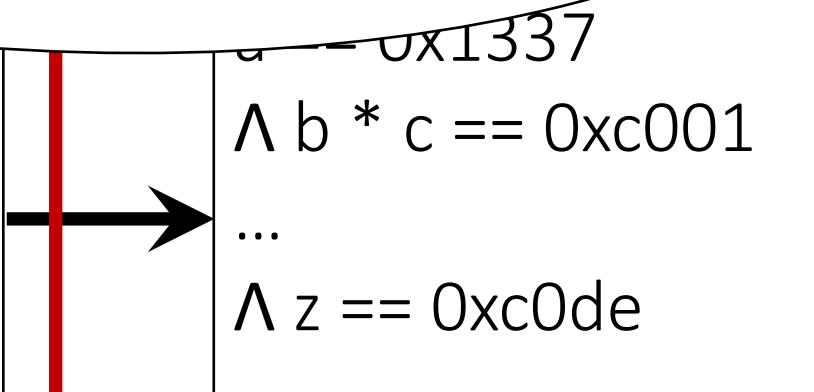
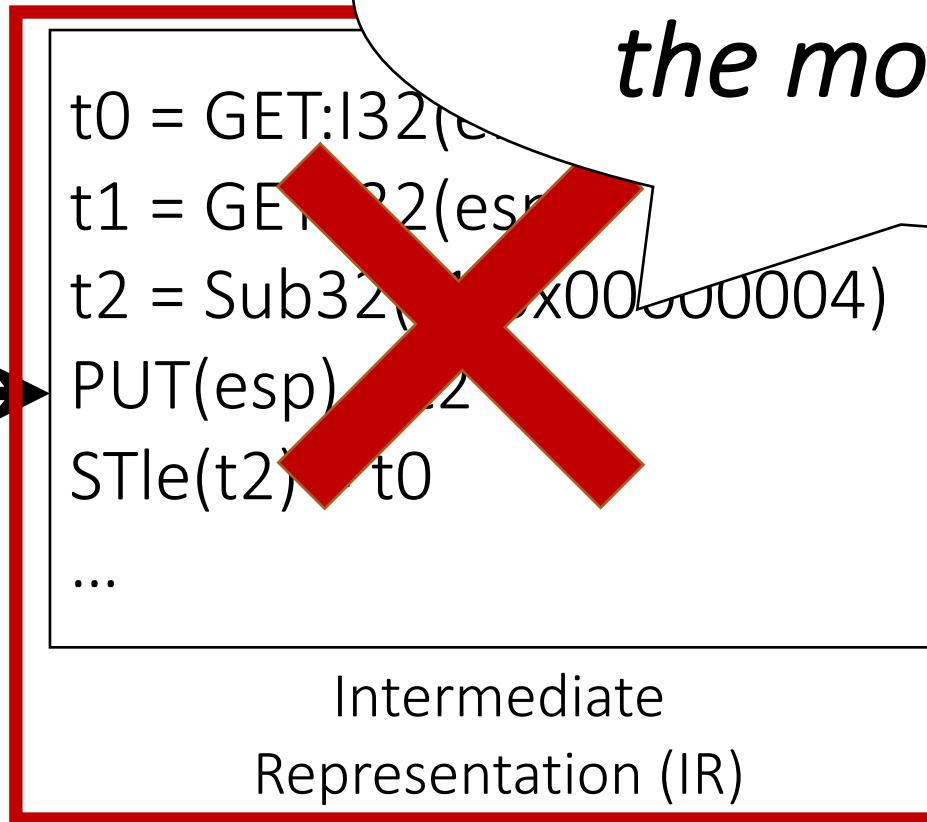
Bad: Performance bottleneck

# How to solve this challenge?

Performance is  
*the most important!*



Assembly



Constraints

Good: Simplifying implementations  
e.g., 981 in x86 vs 115 in VEX

Bad: Performance bottleneck

This is a non-trivial job (LoC)

4.7K

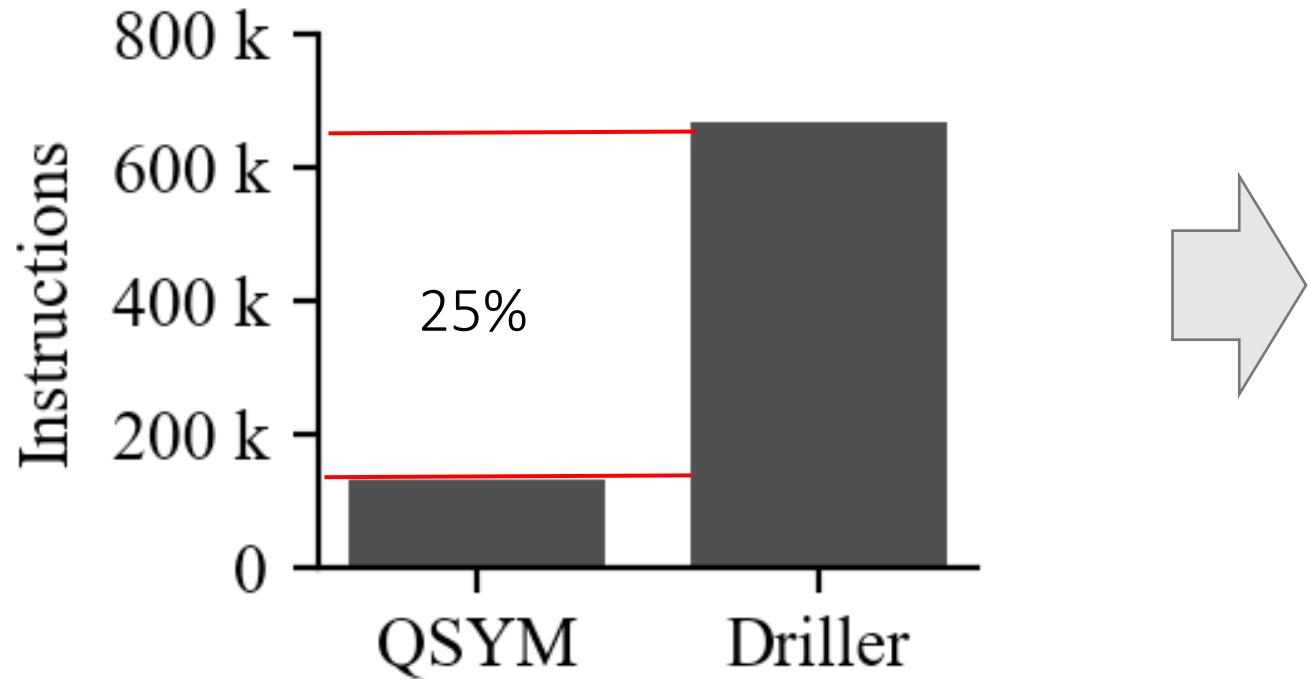
Driller

13K

QSYM

# QSYM reduces the number of symbolically executed instructions

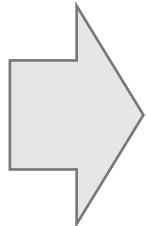
- 126 CGC binaries



*2.5x end-to-end  
performance  
Improvement*

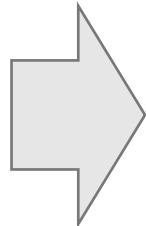
Our system, QSYM, addresses these issues by introducing several key ideas

Generating constraints is too slow



Instruction-level  
concolic execution  
(For binary)

Not effective in  
generating test cases



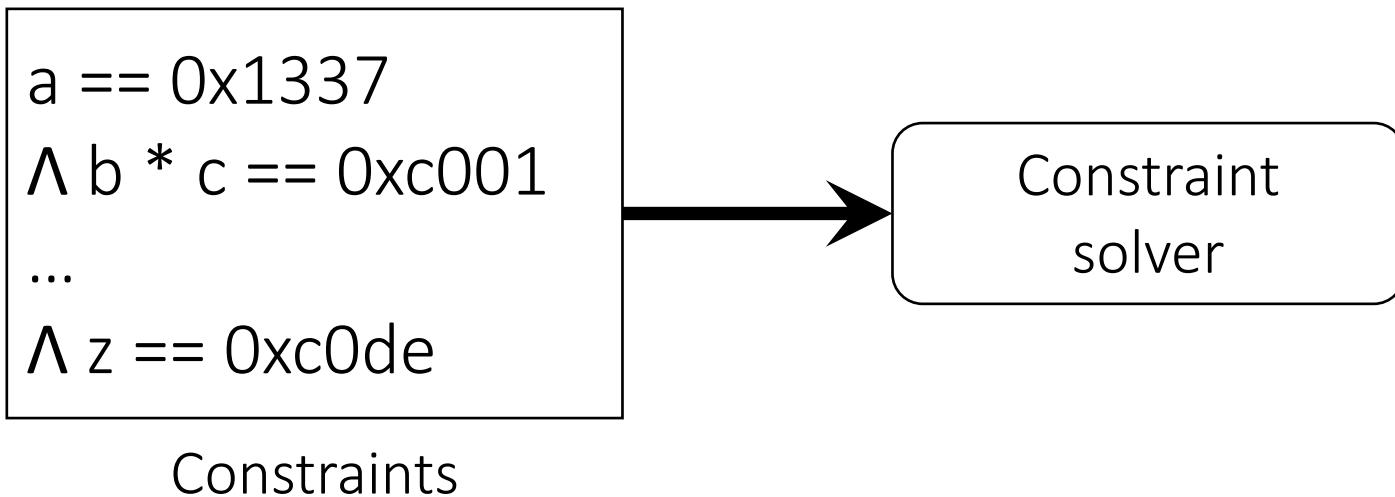
Optimistic solving and  
basic block pruning

Constraint solving can generate a test case  
that meets given constraints

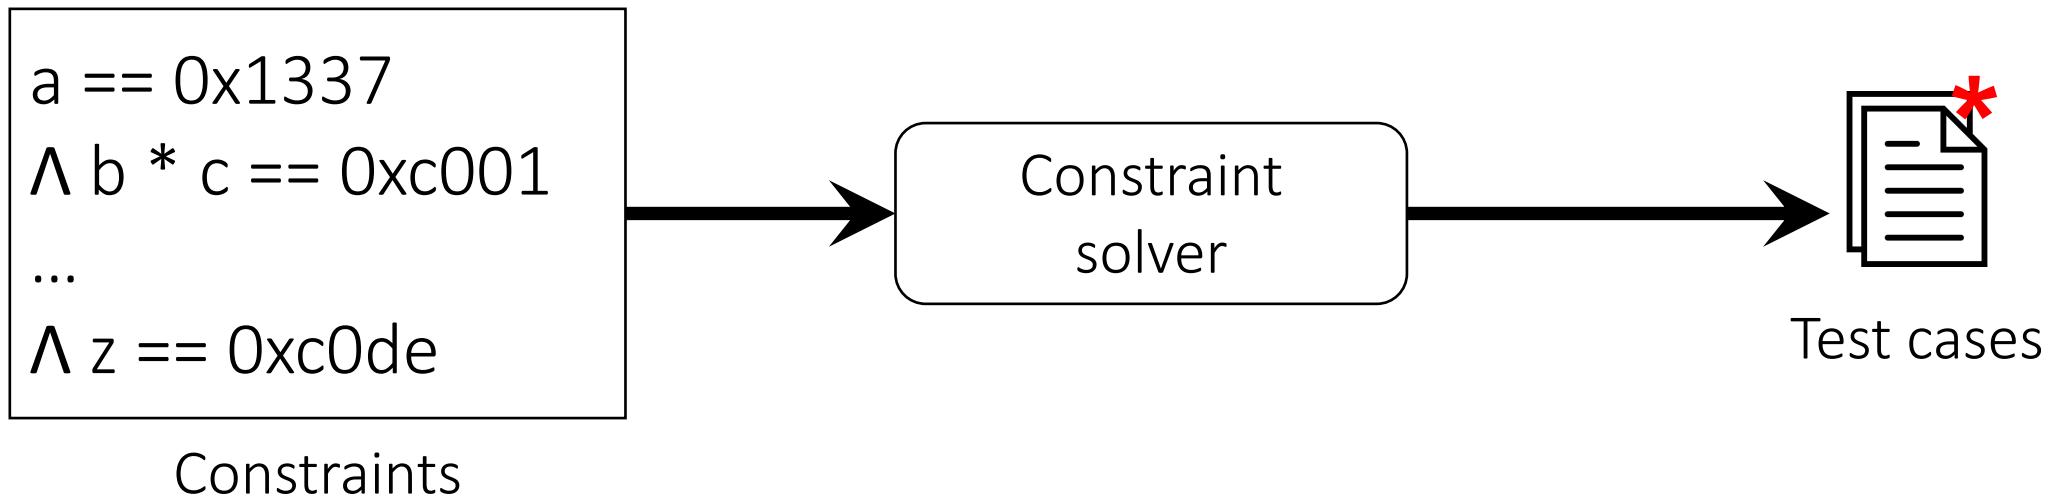
```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

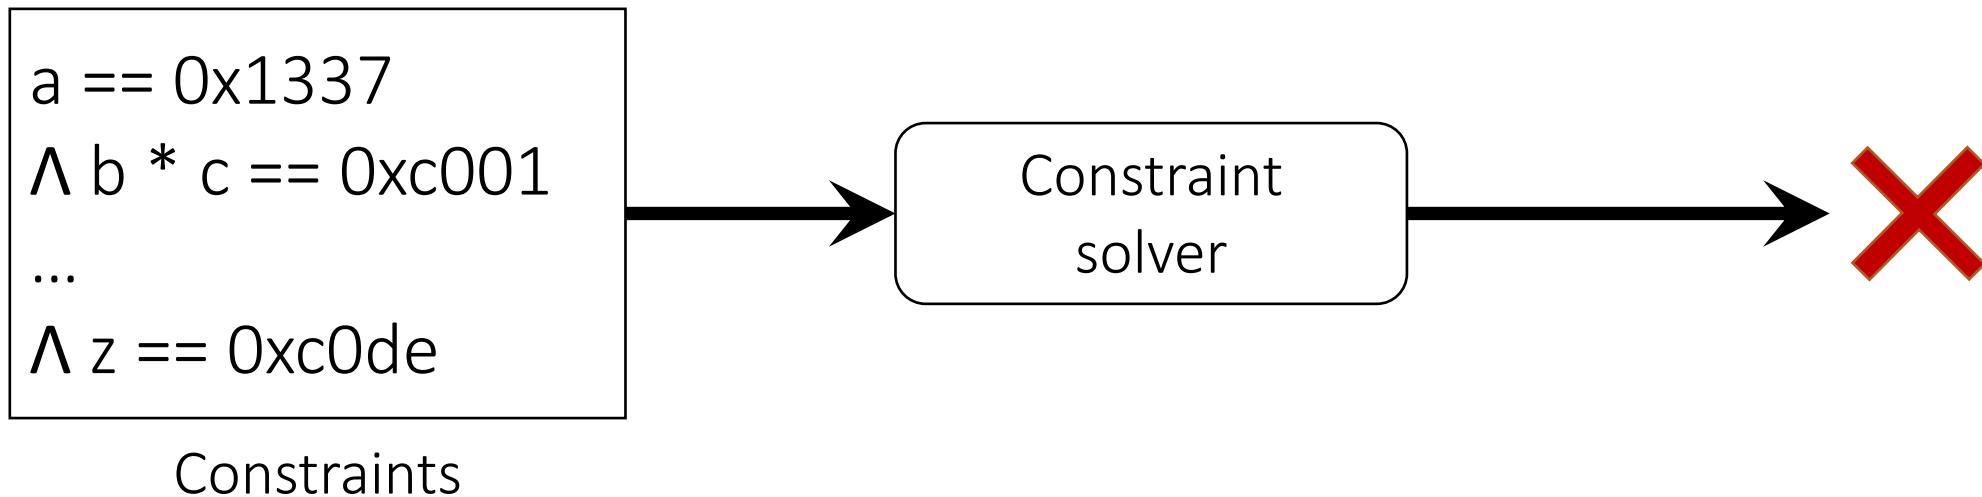
Constraint solving can generate a test case that meets given constraints



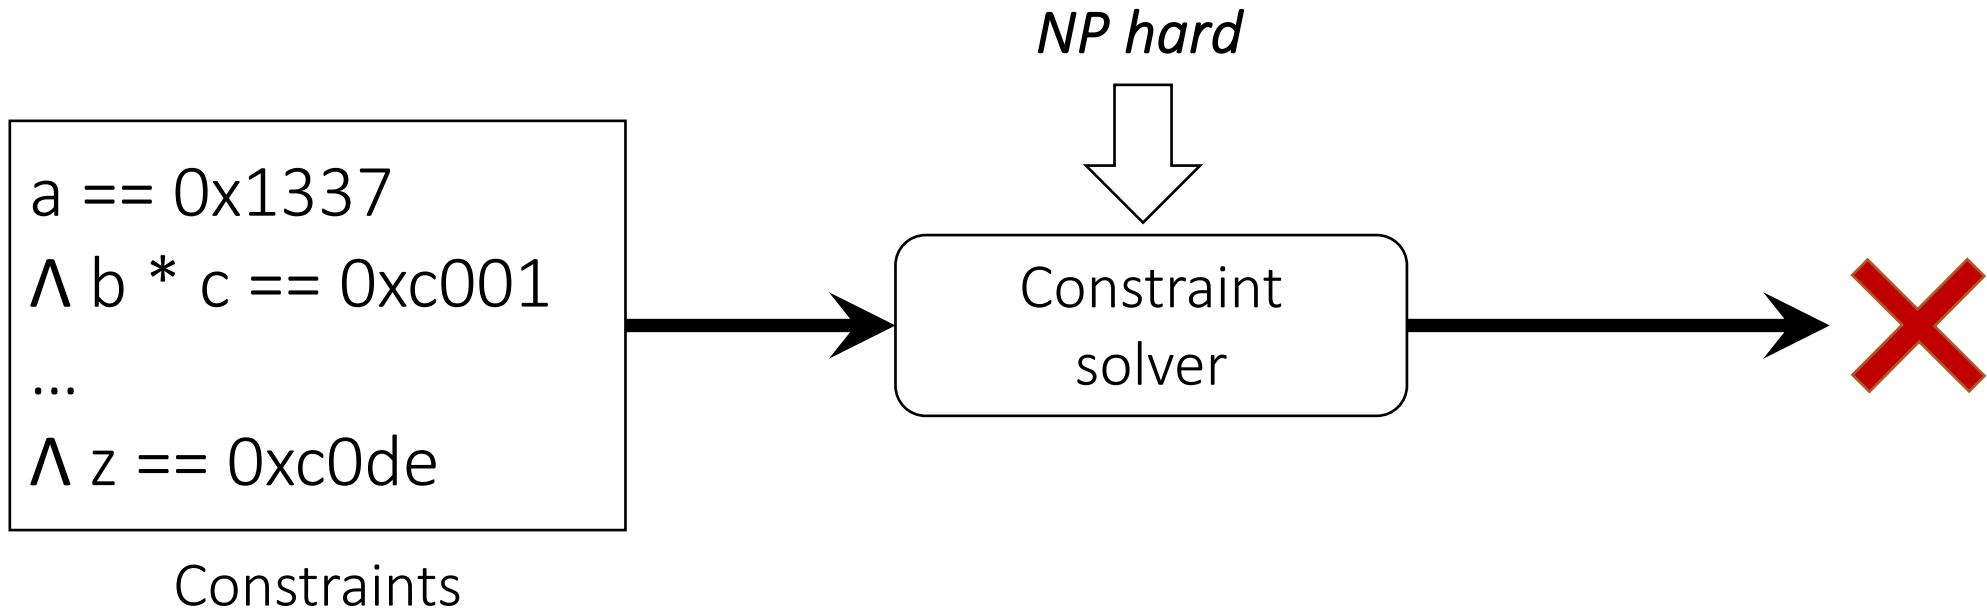
Constraint solving can generate a test case that meets given constraints



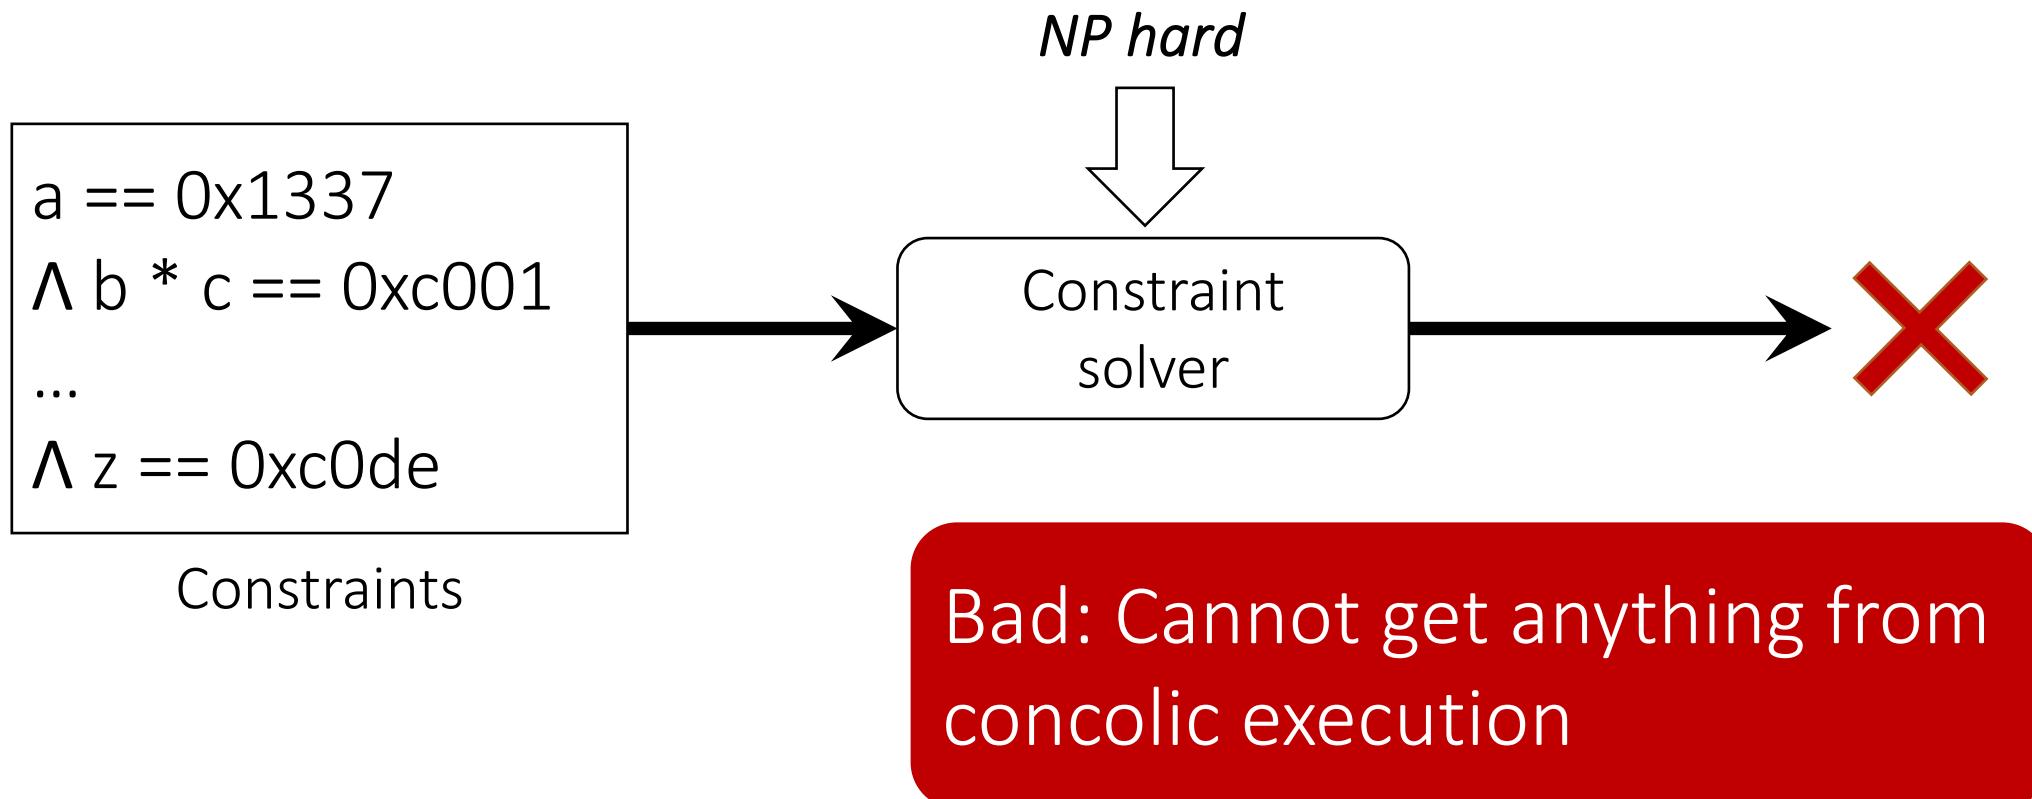
Constraint solving **CANNOT** generate a test case that meets given constraints



Constraint solving **CANNOT** generate a test case that meets given constraints



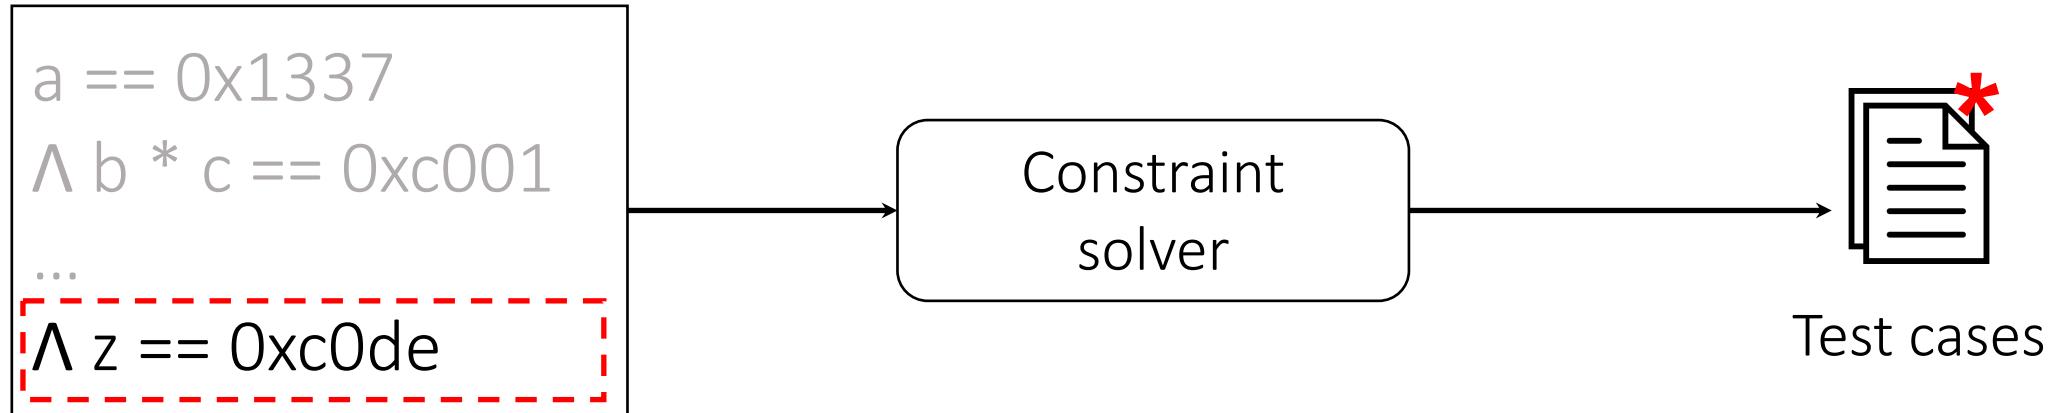
# Constraint solving CANNOT generate a test case that meets given constraints



QSYM solves partial constraints to find some test cases

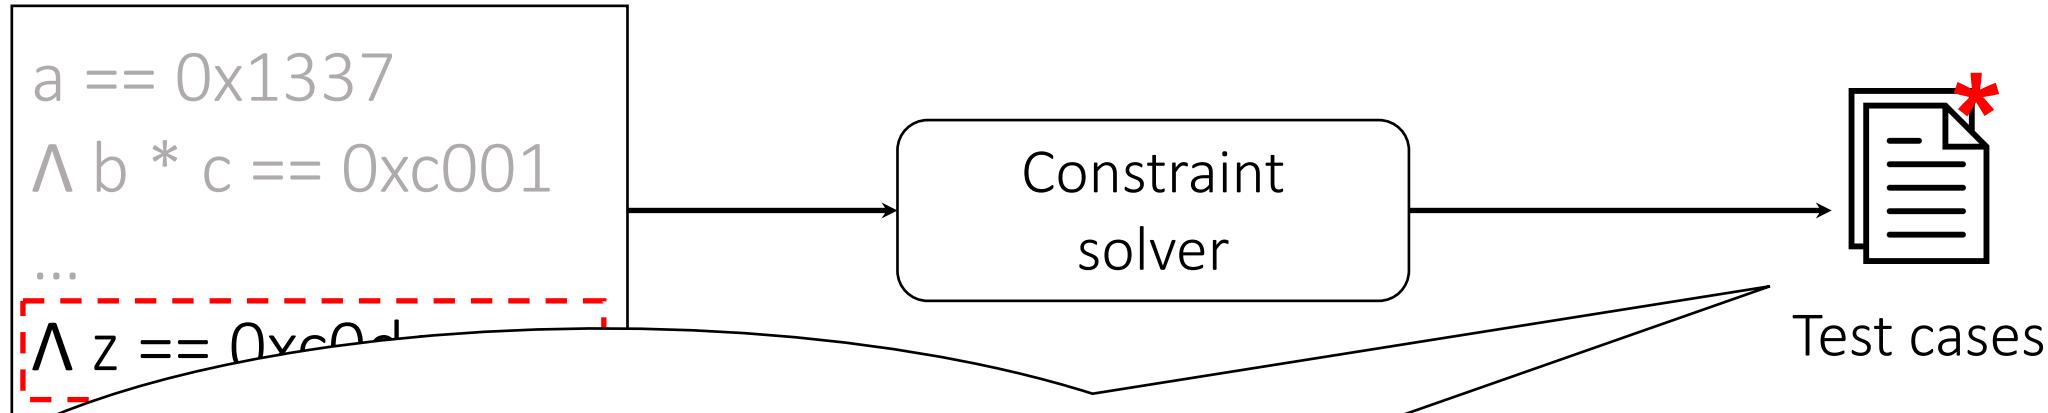
```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

# QSYM solves partial constraints to find some test cases



Good: Can get test cases from  
concolic execution

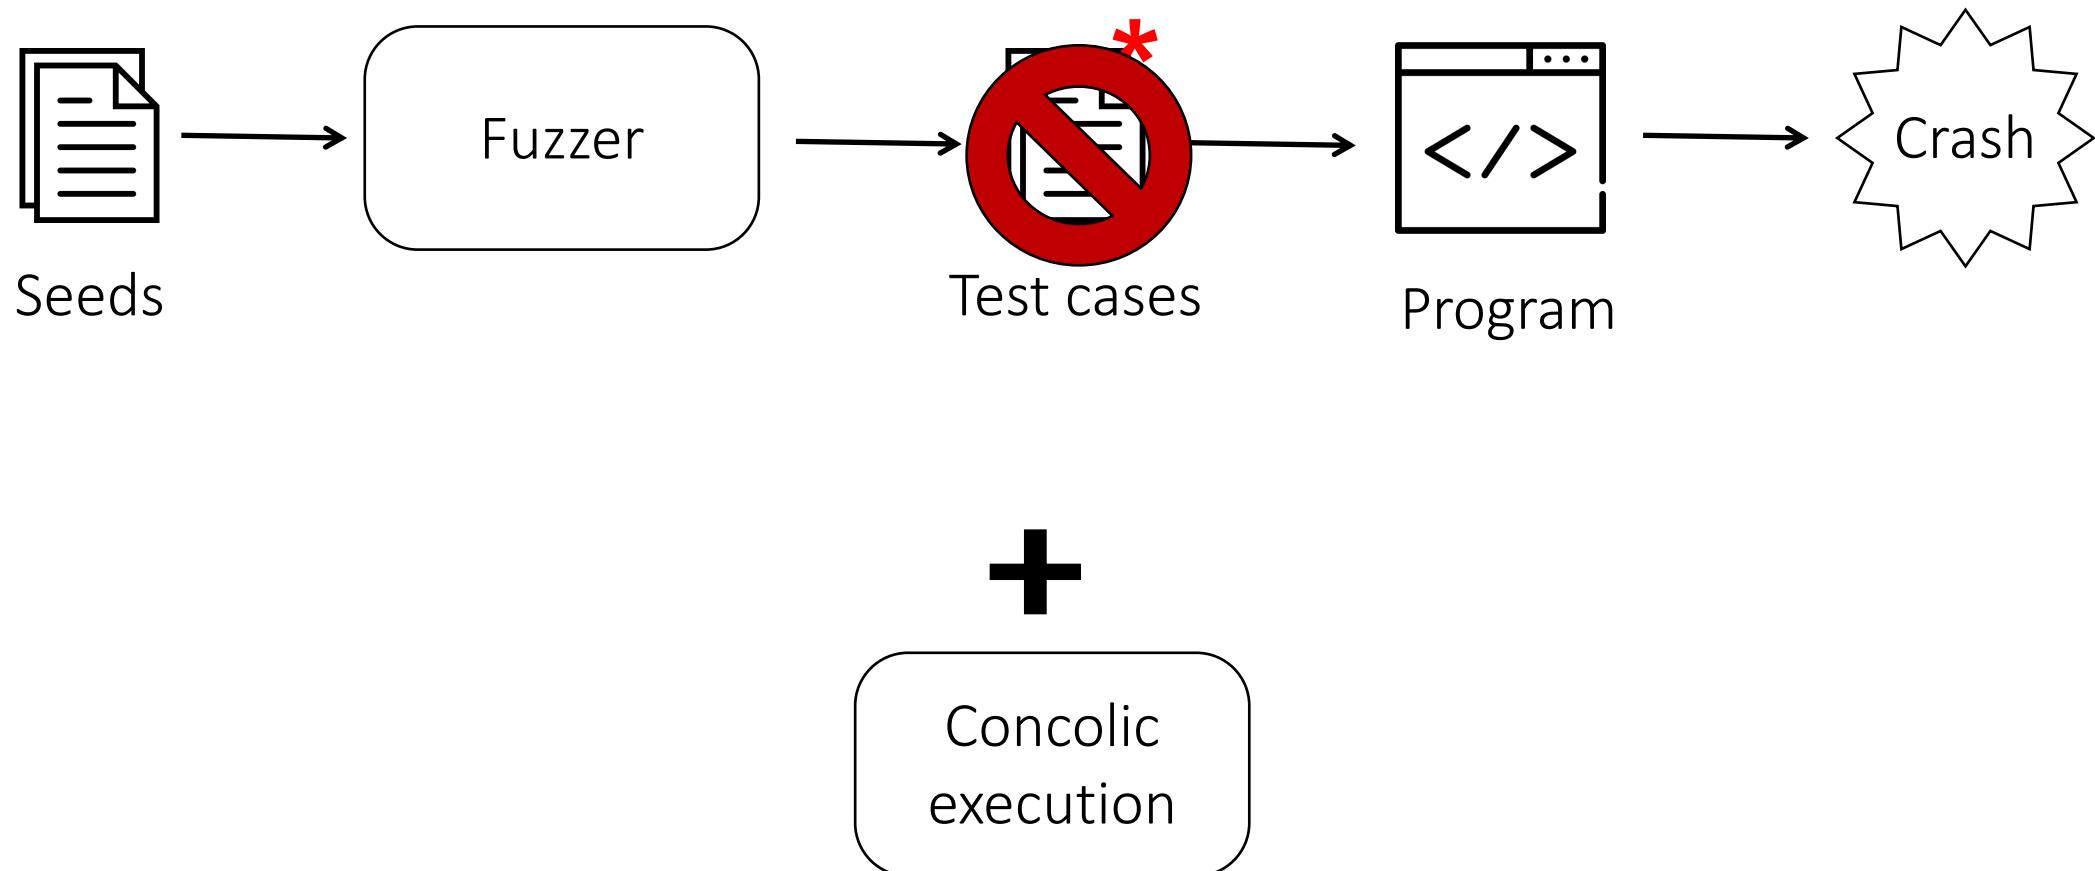
# QSYM solves partial constraints to find some test cases



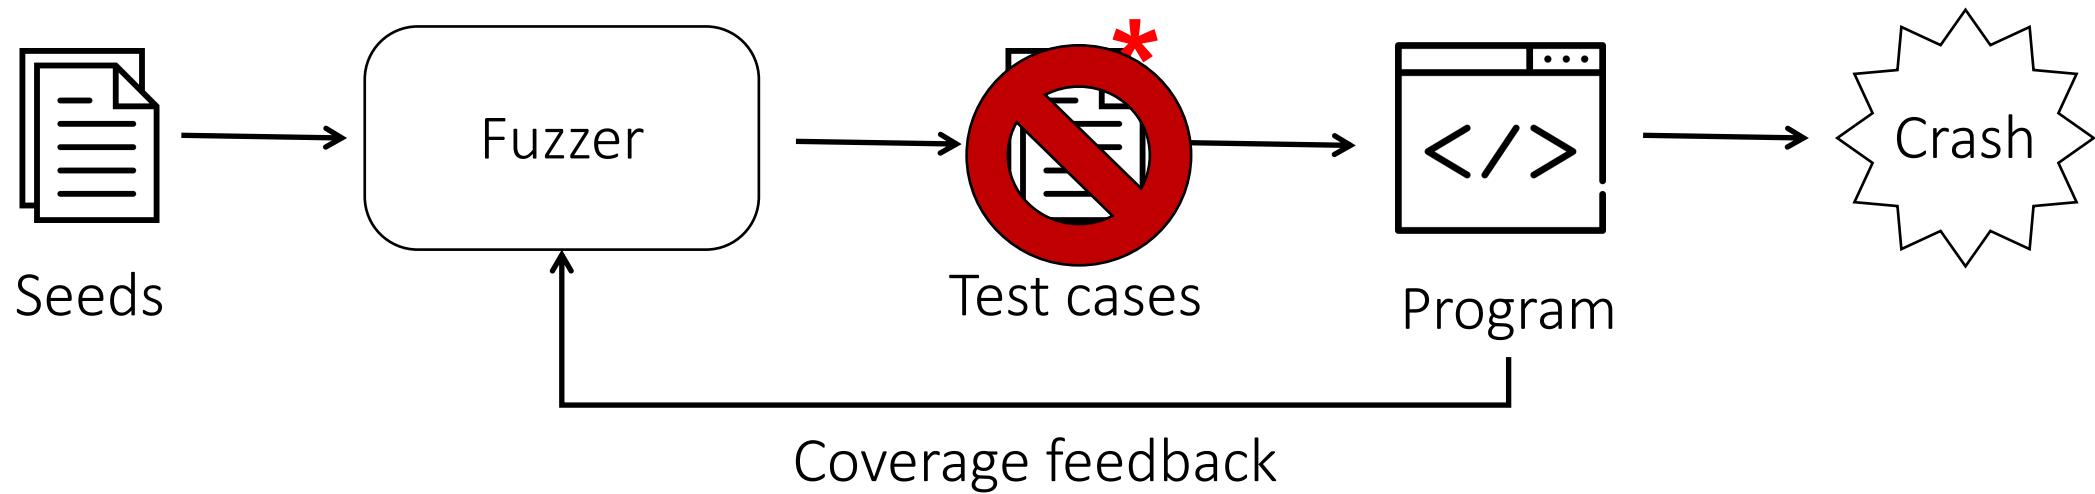
If the test cases  
are *incorrect*...?

get test cases from  
execution

In hybrid fuzzing, generating incorrect inputs is fine because of coverage-guided fuzzing



In hybrid fuzzing, generating incorrect inputs is fine because of coverage-guided fuzzing



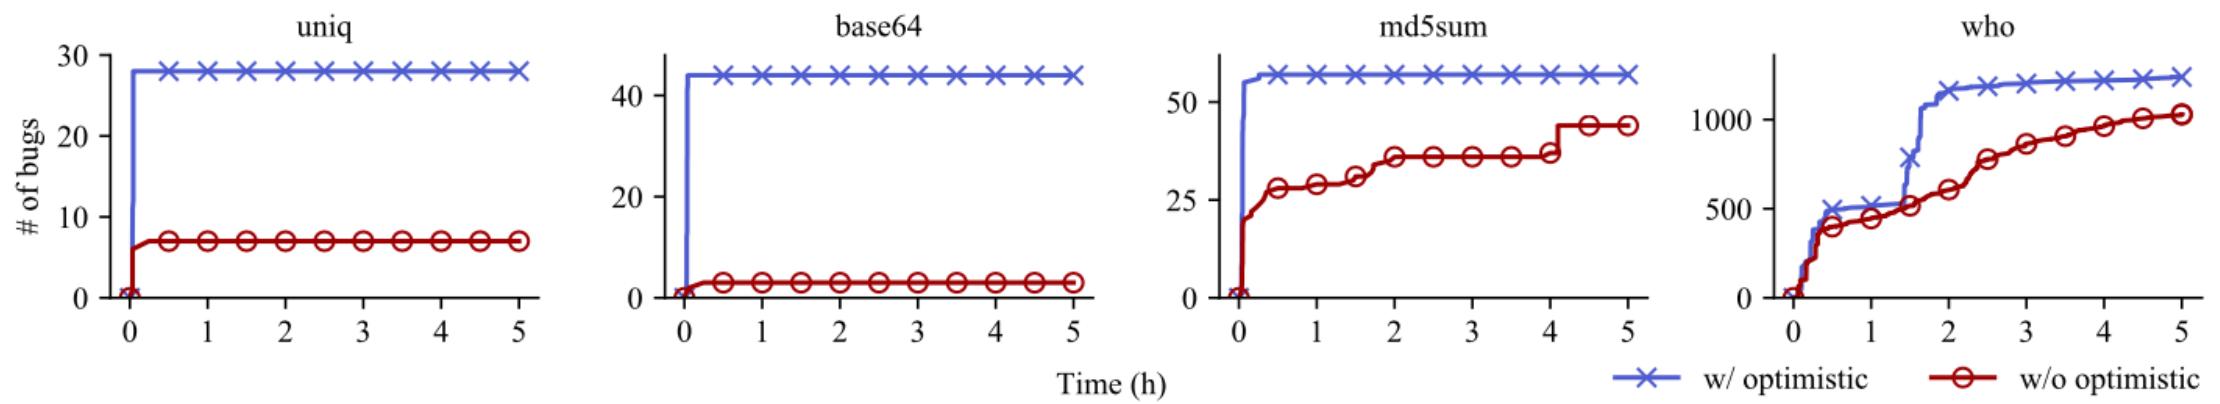
Coverage feedback  
will *filter out*  
incorrect test cases!

+

Concolic execution

# Optimistic solving helps to find more bugs

- LAVA-M dataset
  - Inject hard-to-reach bugs in real-world applications



Remind: Completeness of concolic execution often blocks its further exploration

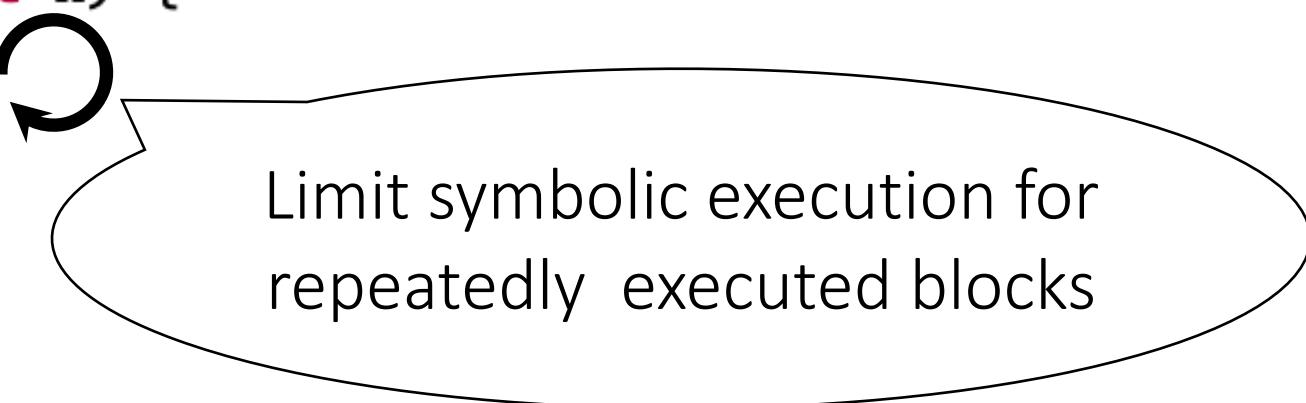
```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```

Remind: Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf); 
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```

Remind: Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```



Limit symbolic execution for repeatedly executed blocks

# Remind: Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```



Limit symbolic execution for repeatedly executed blocks

Can further explore even with such a complicated routine

# Remind: Completeness of concolic execution often blocks its further exploration

```
1 // 'buf' and 'x' are symbolic
2 int completeness(char* buf, int x) {
3     very_complicated_logic(buf);
4
5     if (x * x == 1234 * 1234)
6         crash();
7 }
```



Limit symbolic execution for repeatedly executed blocks

Can further explore even with such a complicated routine

Incomplete constraints

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()  
  
// x != 0 is missed because  
of basic block pruning
```

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()  
  
// x != 0 is missed because  
of basic block pruning
```

```
if y == Oxdeadbeef :
```

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()
```

// x != 0 is missed because  
of basic block pruning

```
if y == Oxdeadbeef :
```



Independent constraints: Use x != 0 in the input

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()
```

// x != 0 is missed because  
of basic block pruning

```
if y == Oxdeadbeef :
```



```
if x == Oxdeadbeef :
```

Independent constraints: Use x != 0 in the input

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()  
  
// x != 0 is missed because  
of basic block pruning
```

```
if y == Oxdeadbeef :
```



Independent constraints: Use x != 0 in the input

```
if x == Oxdeadbeef :
```



Subsumed constraints

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()  
  
// x != 0 is missed because  
of basic block pruning
```

```
if y == Oxdeadbeef :
```



Independent constraints: Use x != 0 in the input

```
if x == Oxdeadbeef :
```



Subsumed constraints

```
if x != Oxdeadbeef :
```

# Incomplete constraints are not significant in practice for hybrid fuzzing

```
x = input()  
y = input()  
  
// x != 0 is missed because  
of basic block pruning
```

```
if y == Oxdeadbeef :
```



Independent constraints: Use x != 0 in the input

```
if x == Oxdeadbeef :
```



Subsumed constraints

```
if x != Oxdeadbeef :
```



Failing...

# Evaluating QSYM

- Scaling to real-world software?
- How good is QSYM compared to
  - The state-of-art hybrid fuzzing (Driller)

# QSYM scales to real-world software

- 13 bugs in real-world software (already tested by fuzzing)

<b>Program</b>	<b>CVE</b>	<b>Bug Type</b>	<b>Fuzzer</b>
lepton	CVE-2017-8891	Out-of-bounds read	AFL
openjpeg	CVE-2017-12878	Heap overflow	OSS-Fuzz
	Fixed by other patch	NULL dereference	
tcpdump	CVE-2017-11543*	Heap overflow	AFL
file	CVE-2017-1000249*	Stack overflow	OSS-Fuzz
libarchive	Wait for patch	NULL dereference	OSS-Fuzz
audiofile	CVE-2017-6836	Heap overflow	AFL
	Wait for patch	Heap overflow × 3	
	Wait for patch	Memory leak	
ffmpeg	CVE-2017-17081	Out-of-bounds read	OSS-Fuzz
objdump	CVE-2017-17080	Out-of-bounds read	AFL

# QSYM scales to real-world software

- 13 bugs in real-world software (already tested by fuzzing)

Real-world software

Program	CVE	Bug Type	Fuzzer
lepton	CVE-2017-8891	Out-of-bounds read	AFL
openjpeg	CVE-2017-12878	Heap overflow	OSS-Fuzz
	Fixed by other patch	NULL dereference	
tcpdump	CVE-2017-11543*	Heap overflow	AFL
file	CVE-2017-1000249*	Stack overflow	OSS-Fuzz
libarchive	Wait for patch	NULL dereference	OSS-Fuzz
audiofile	CVE-2017-6836	Heap overflow	AFL
	Wait for patch	Heap overflow × 3	
	Wait for patch	Memory leak	
ffmpeg	CVE-2017-17081	Out-of-bounds read	OSS-Fuzz
objdump	CVE-2017-17080	Out-of-bounds read	AFL

# QSYM scales to real-world software

- 13 bugs in real-world software (already tested by fuzzing)

Program	CVE	Bug Type	Fuzzer
lepton	CVE-2017-8891	Out-of-bounds read	AFL
openjpeg	CVE-2017-12878	Heap overflow	OSS-Fuzz
	Fixed by other patch	NULL dereference	
tcpdump	CVE-2017-11543*	Heap overflow	AFL
file	CVE-2017-1000249*	Stack overflow	OSS-Fuzz
libarchive	Wait for patch	NULL dereference	OSS-Fuzz
audiofile	CVE-2017-6836	Heap overflow	AFL
	Wait for patch	Heap overflow × 3	
	Wait for patch	Memory leak	
ffmpeg	CVE-2017-17081	Out-of-bounds read	OSS-Fuzz
objdump	CVE-2017-17080	Out-of-bounds read	AFL

Already  
heavily  
fuzzed

# QSYM can generate test cases that fuzzing is hard to find

- e.g.) ffmpeg: Not reachable by fuzzing

```
if( !((ox^(ox+dxw))
       | (ox^(ox+dxh))
       | (ox^(ox+dxw+ dxh))
       | (oy^(oy+dyw))
       | (oy^(oy+dyh))
       | (oy^(oy+dyw+ dyh))) >> (16 + shift)
       && !(dxx | dxy | dyx | dyy) & 15
       && !(need_emu&&(h>MAX_H || stride > MAX_STRIDE)))
{ // the bug is here ; }
```

# QSYM can generate test cases that fuzzing is hard to find

- e.g.) ffmpeg: Not reachable by fuzzing

```
if( !(c
```

Cannot be found by cloud fuzzing  
(Note: 4 trillion test cases per week)

```
    | (oy^(oy+dyh))  
    | (oy^(oy+dyw+ dyh))) >> (16 + shift)  
    && !(dxx | dxy | dyx | dyy) & 15  
    && !(need_emu&&(h>MAX_H || stride > MAX_STRIDE))  
{ // the bug is here ; }
```

# QSYM can generate test cases that fuzzing is hard to find

- e.g.) ffmpeg: Not reachable by fuzzing

```
if( !(c
```

Cannot be found by cloud fuzzing  
(Note: 4 trillion test cases per week)

```
| (oy^(oy+dyh))
```



Found by a *single workstation*  
using QSYM

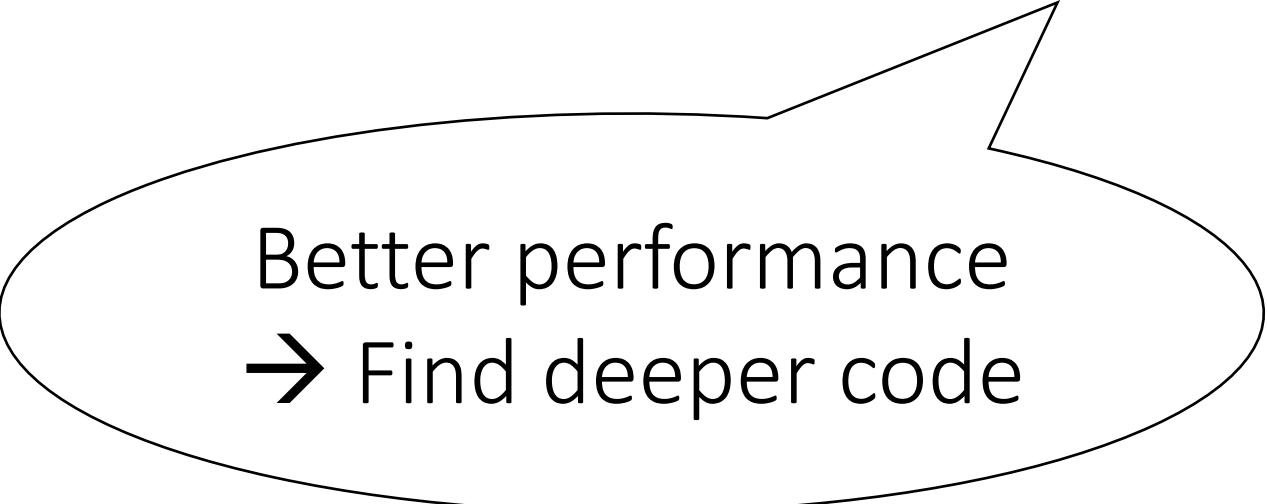
```
{ // the bug is here , ,
```

```
RIDE)))
```

QSYM outperforms Driller, the state-of-the-art hybrid fuzzer

# QSYM outperforms Driller, the state-of-the-art hybrid fuzzer

- Dataset: 126 CGC binaries
- Compare code coverage achieved by a single run of concolic execution
- QSYM achieved more code coverage in ***104 (82%)*** binaries



Better performance  
→ Find deeper code

# QSYM is also practically impactful

- e.g., Rode0day: A monthly competition for automatic bug finding tool

<b>Table 2. The overall rankings for top Rode0day competitors after 10 competitions.</b>		
<b>Place</b>	<b>Elo score</b>	<b>Team name</b>
1	1,087	afl-lazy
2	1,069	itszn
3	1,027	H3ku
4	1,017	REDQUEEN
5	1,062	NU-AFL-QSYM

Fasano, Andrew, et al. "The Rode0day to Less-Buggy Programs." *IEEE Security & Privacy* (2019)

# QSYM is also practically impactful

- e.g., Rode0day: A monthly competition for automatic bug finding tool

<b>Table 2. The overall rankings for top Rode0day competitors after 10 competitions.</b>		
<b>Place</b>	<b>Elo score</b>	<b>Team name</b>
1	1,087	afl-lazy
2	1,069	itszn
3	1,027	H3ku
4	1,017	REDQUEEN
5	1,062	NU-AFL-QSYM

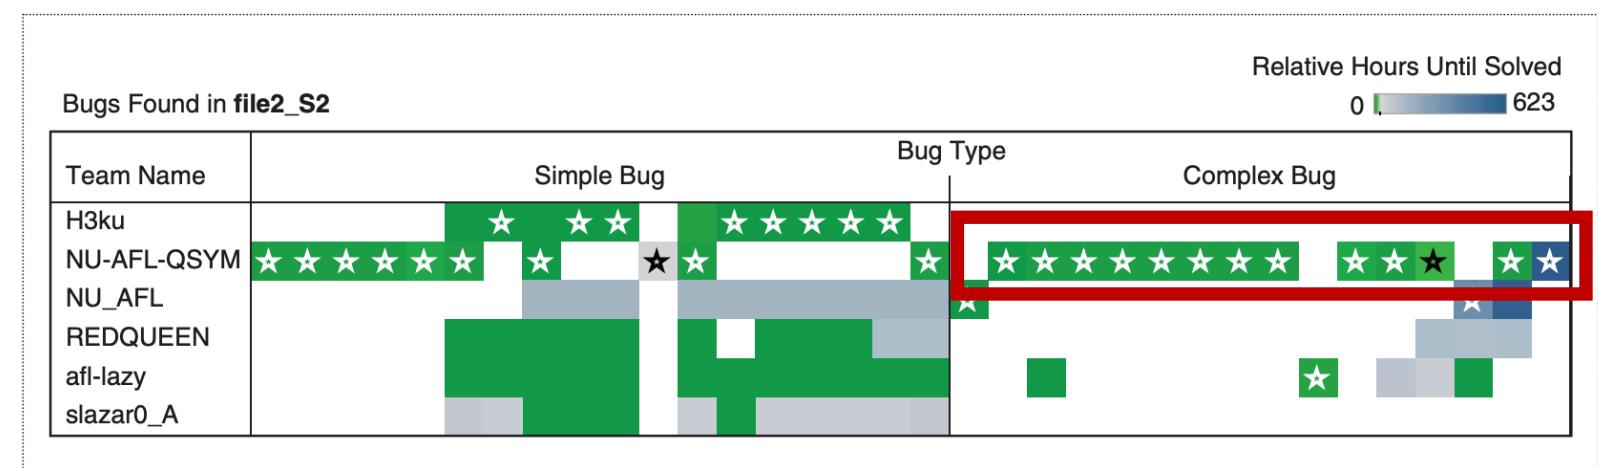


Figure 2. A visualization of when teams found bugs in file2\_S2 during the November 2018 Rode0day. Green indicates the bugs found within 24 h of a team's first score. The stars denote the first team to find a bug.

Fasano, Andrew, et al. "The Rode0day to Less-Buggy Programs." *IEEE Security & Privacy* (2019)

# Today's talk

QSYM: A Binary-level  
Concolic Execution Engine  
for Hybrid fuzzing

- Binary
- User applications

Hybridra: A Hybrid Fuzzer  
for Kernel File Systems

- Source code
- File systems

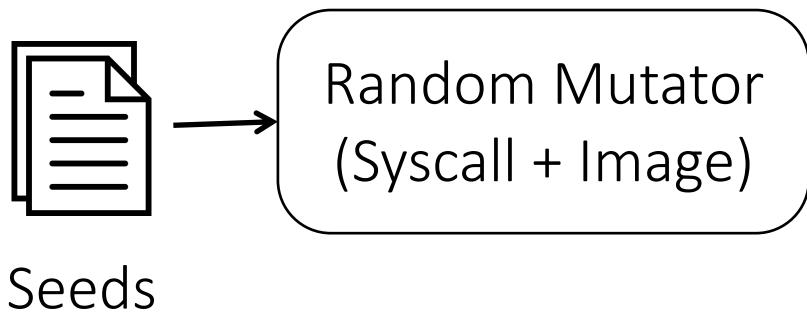
# Hybridra improves Hydra by supporting concolic image mutation



Seeds

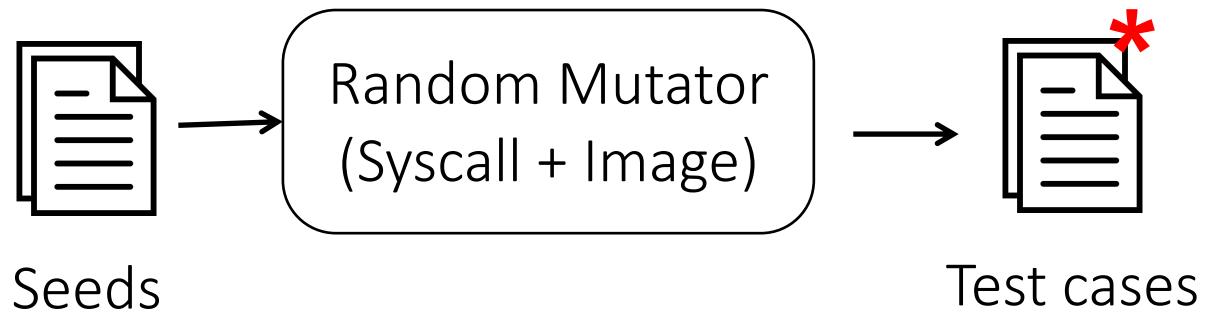
Hydra

# Hybridra improves Hydra by supporting concolic image mutation



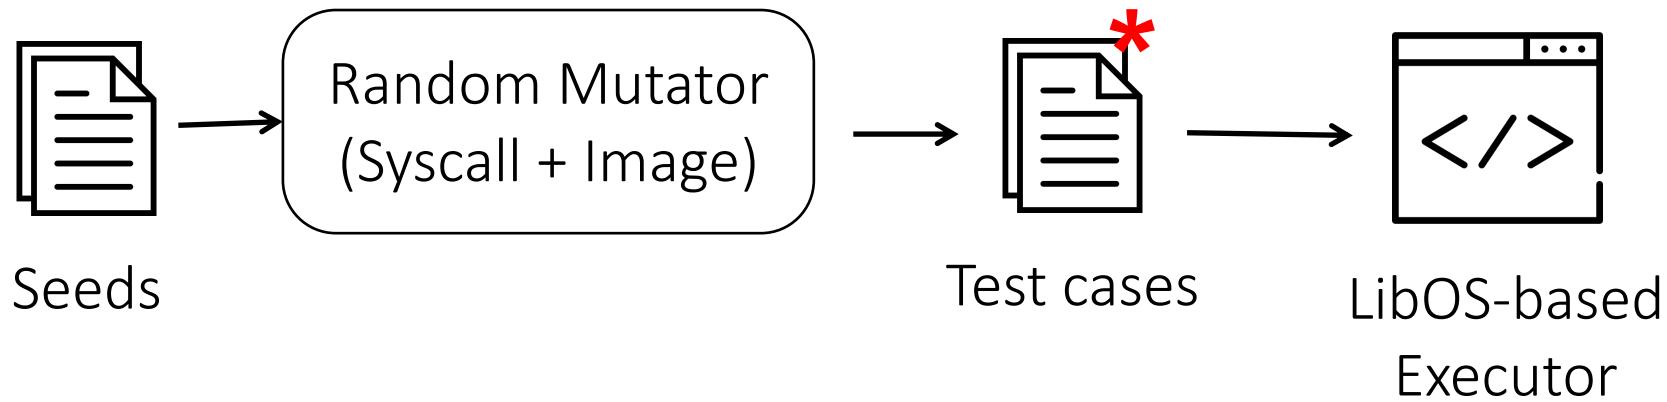
Hydra

# Hybridra improves Hydra by supporting concolic image mutation



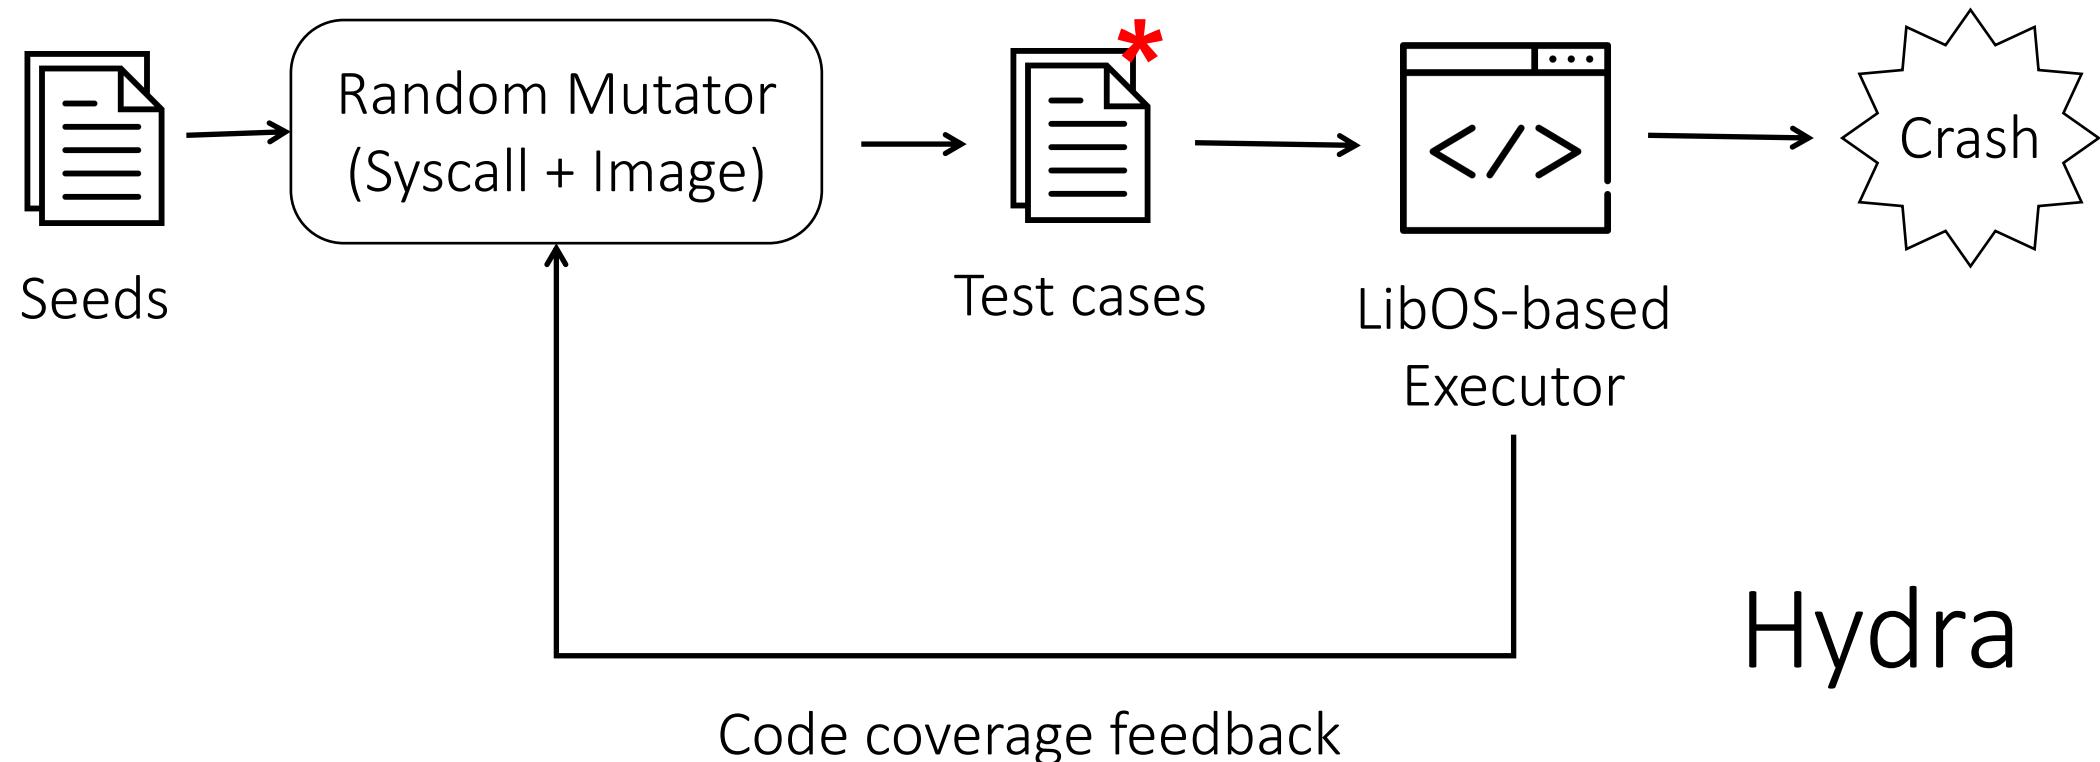
Hydra

# Hybridra improves Hydra by supporting concolic image mutation

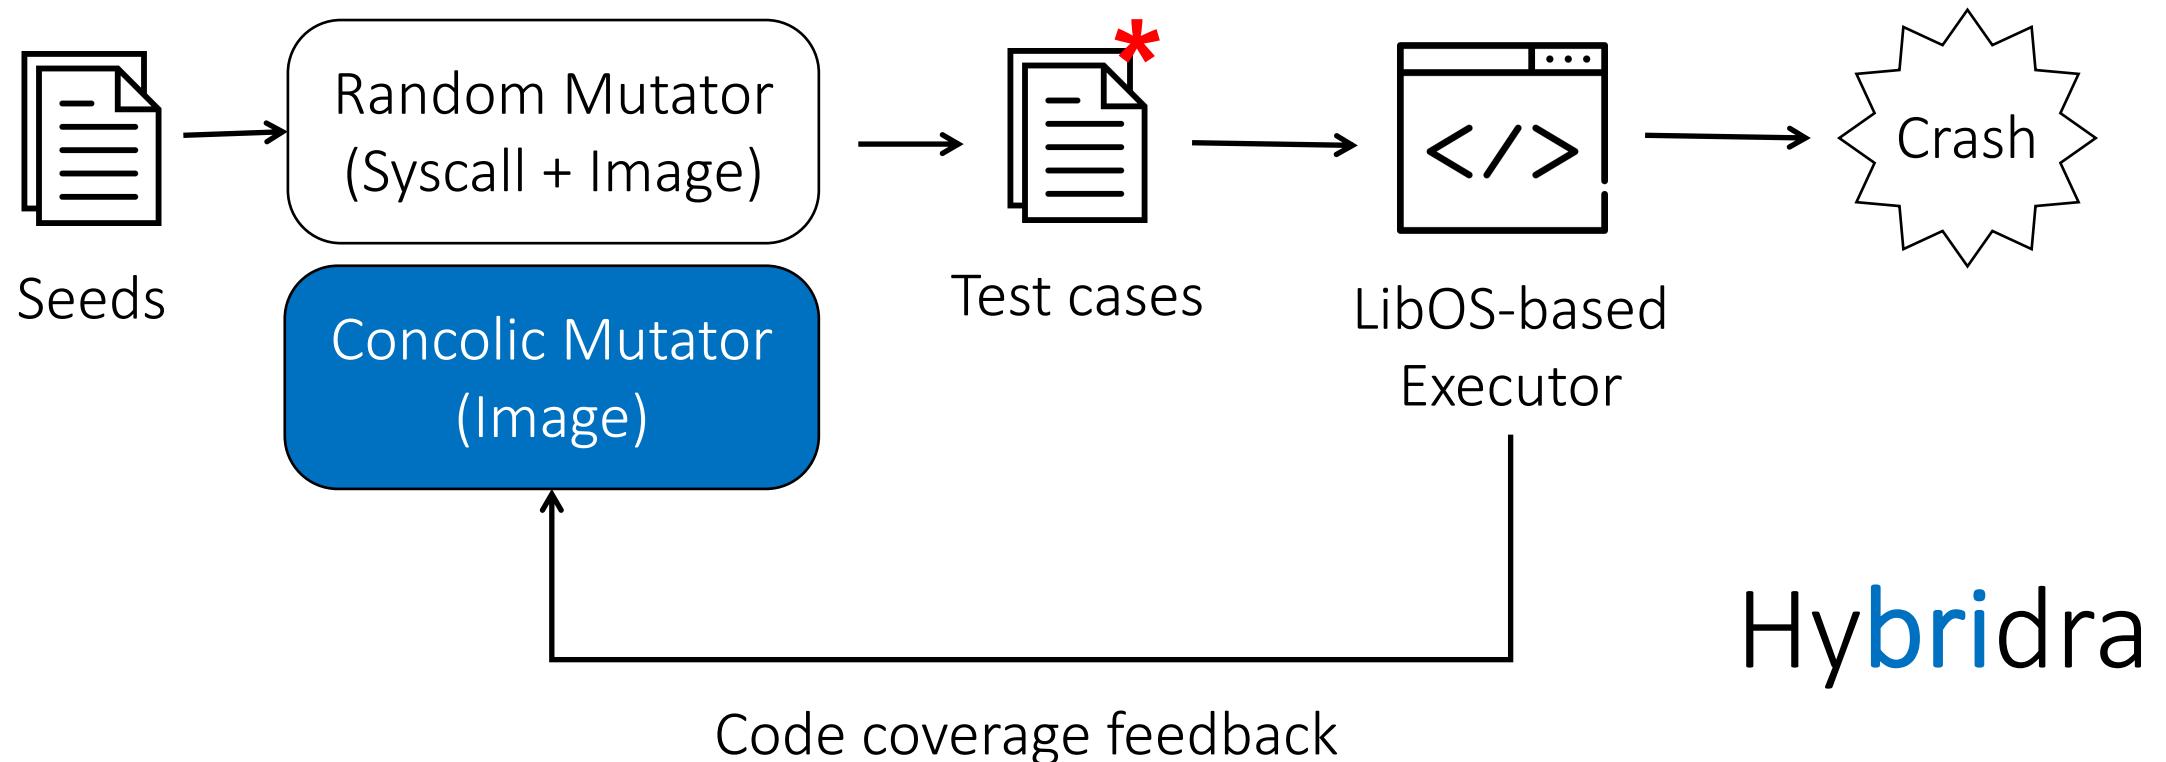


Hydra

# Hybridra improves Hydra by supporting concolic image mutation

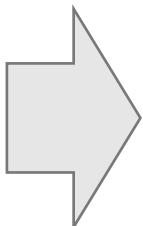


# Hybridra improves Hydra by supporting concolic image mutation



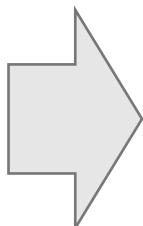
# Hybridra: Key ideas

Generating constraints is too slow



Compilation-based  
concolic execution  
(For source code)

Not effective in  
generating test cases



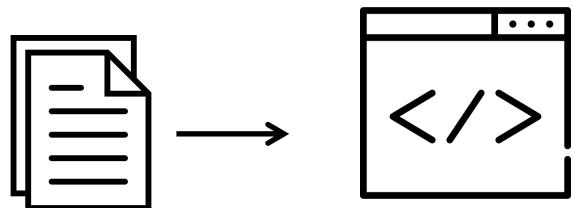
Staged reduction  
+ Heuristics from QSYM

# Design: Concolic Image Mutator



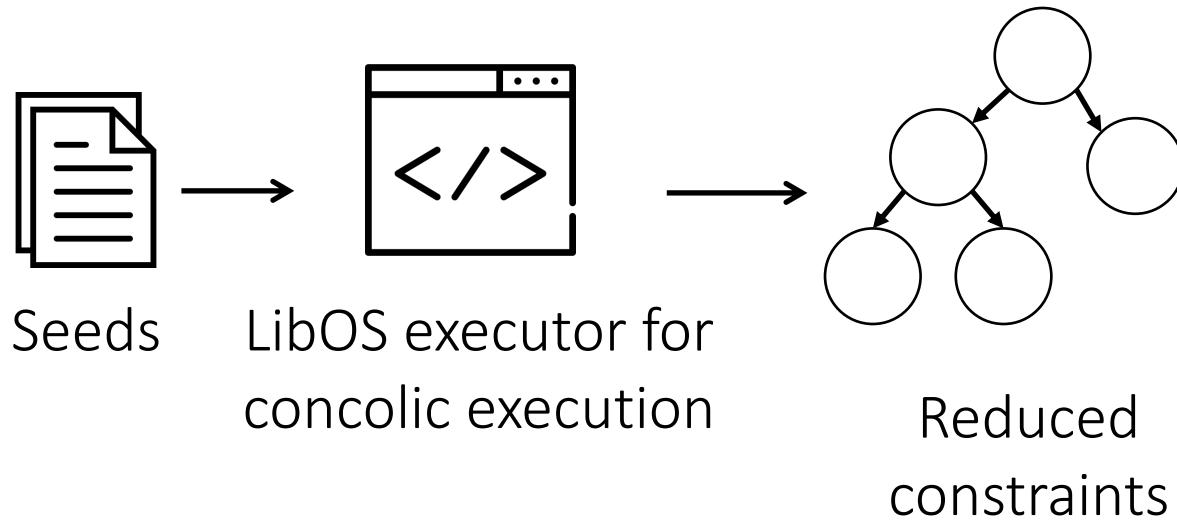
Seeds

# Design: Concolic Image Mutator

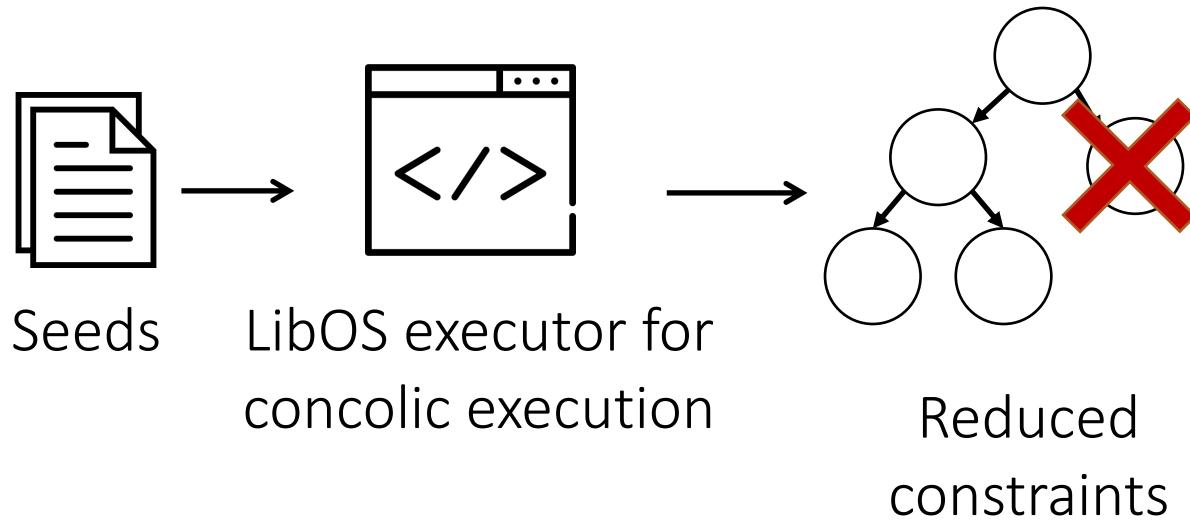


Seeds    LibOS executor for  
concolic execution

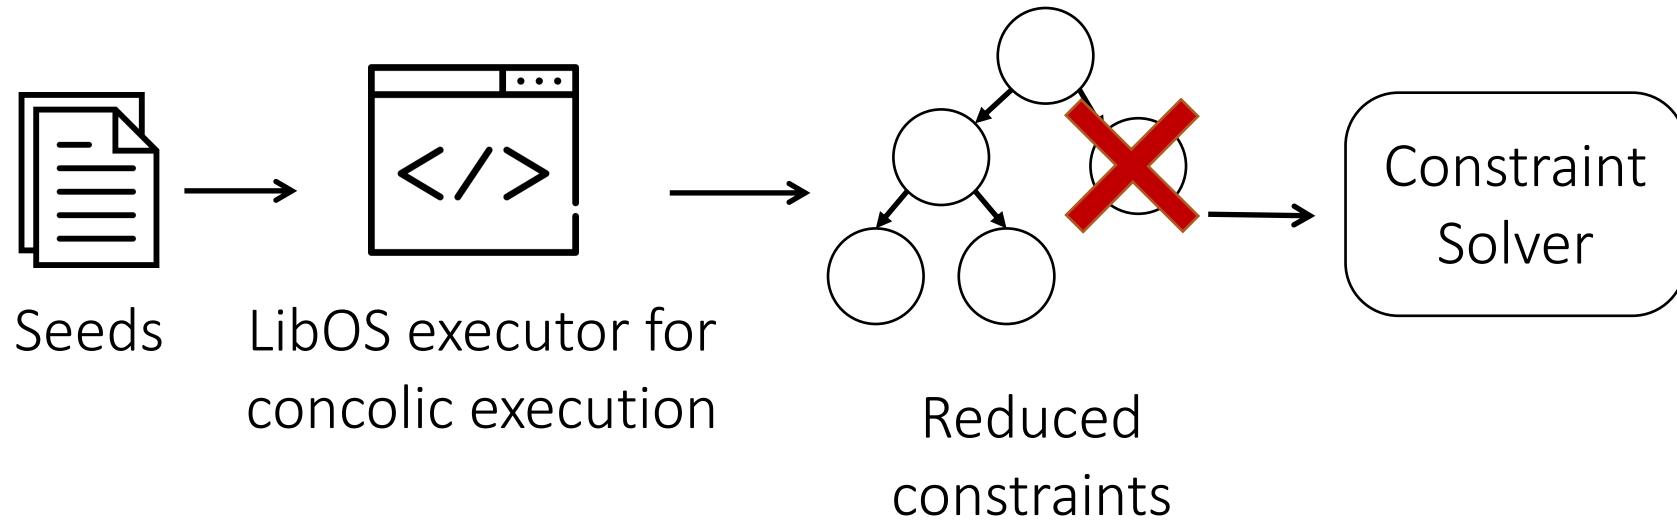
# Design: Concolic Image Mutator



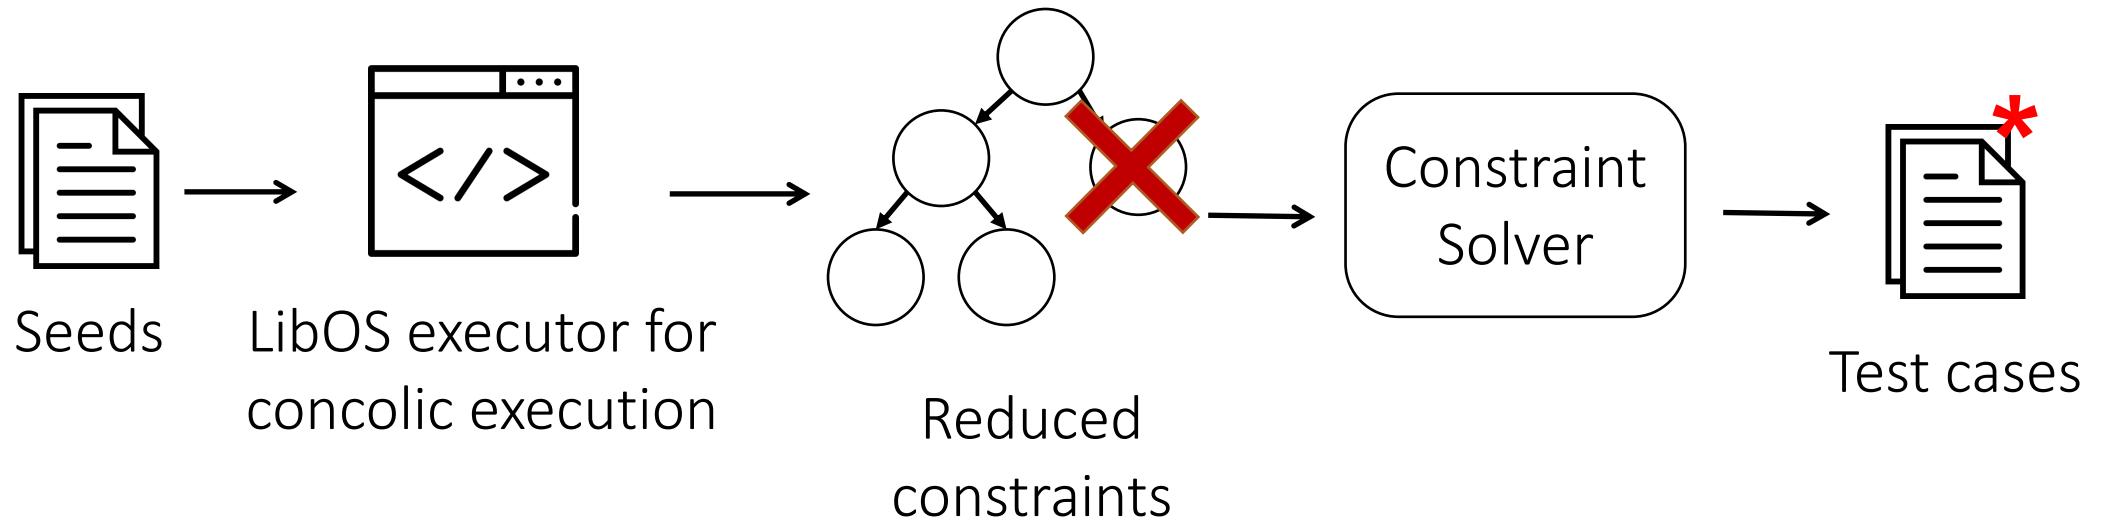
# Design: Concolic Image Mutator



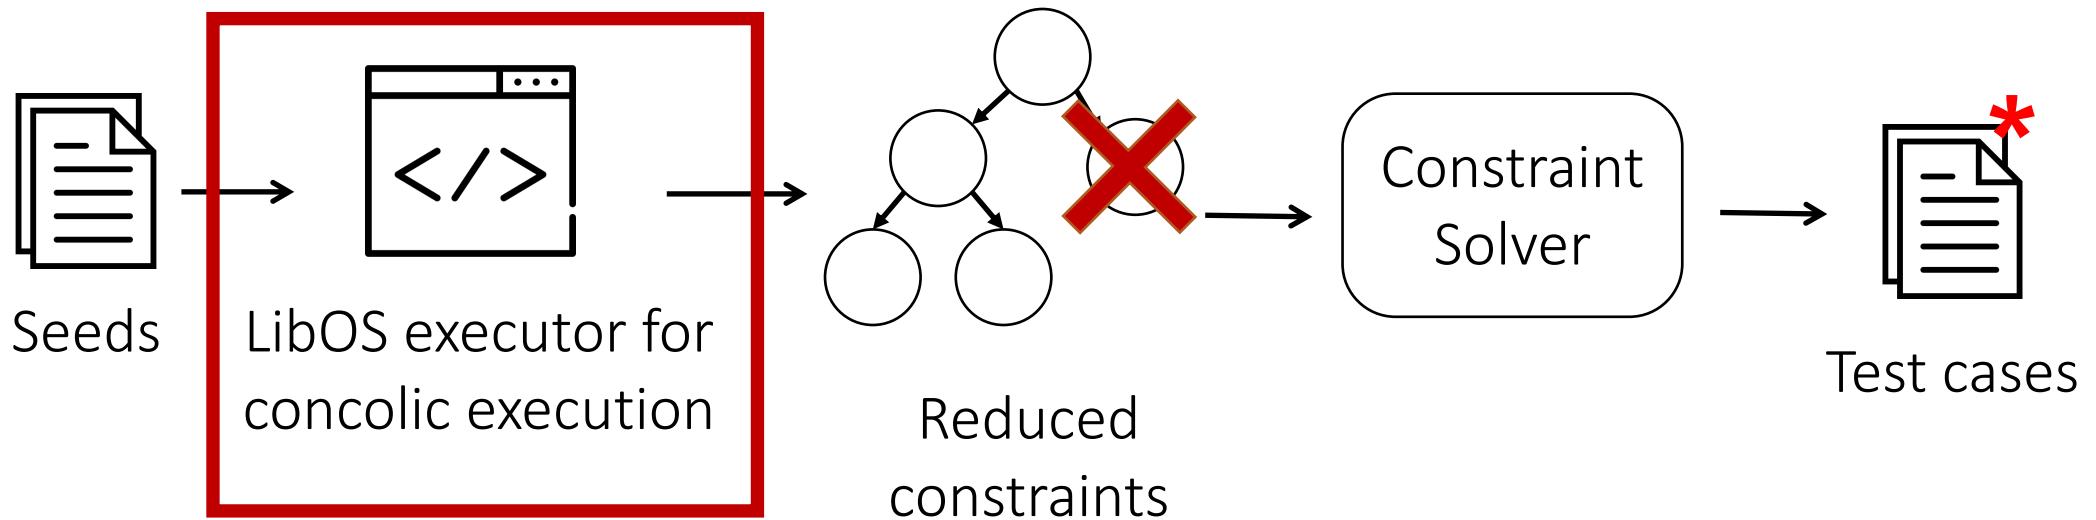
# Design: Concolic Image Mutator



# Design: Concolic Image Mutator



# Design: Concolic Image Mutator



Hybridra utilizes compilation-based concolic execution

# Hybridra utilizes compilation-based concolic execution

```
+ Symbol* symA = getSymbol(a);
+ Symbol* symB = getSymbol(b);
+ Symbol* symC = addSymbol(symA, symB);

int c = a + b;

+ Symbol* symD = getSymbol(d);
+ // Make test cases
+ checkEqual(symC, symD);

if (c == d) {
    ...
}
```

# Hybridra utilizes compilation-based concolic execution

```
+ Symbol* symA = getSymbol(a);
+ Symbol* symB = getSymbol(b);
+ Symbol* symC = addSymbol(symA, symB);

int c = a + b;

+ Symbol* symD = getSymbol(d);
+ // Make test cases
+ checkEqual(symC, symD);

if (c == d) {
    ...
}
```

# Hybridra utilizes compilation-based concolic execution

```
+ Symbol* symA = getSymbol(a);
+ Symbol* symB = getSymbol(b);
+ Symbol* symC = addSymbol(symA, symB);

int c = a + b;

+ Symbol* symD = getSymbol(d);
+ // Make test cases
+ checkEqual(symC, symD);
```

200x performance improvement  
compared to QSYM  
(NOTE: code is required)

Hybridra's compilation-based concolic execution, Kirenenko, is useful by supporting multi-threading

	CUTE	SymCC	Kirenenko
<b>Language</b>	C	LLVM IR	LLVM IR
<b>Memory modeling</b>	Page table	Page table	Shadow memory
<b>Multi-threading</b>			✓

Hybridra's compilation-based concolic execution, Kirenenko, is useful by supporting multi-threading

	CUTE	SymCC	Kirenenko
<b>Language</b>	C	LLVM IR	LLVM IR
<b>Memory modeling</b>	Page table	Page table	Shadow memory
<b>Multi-threading</b>			✓

Hybridra's compilation-based concolic execution, Kirenenko, is useful by supporting multi-threading

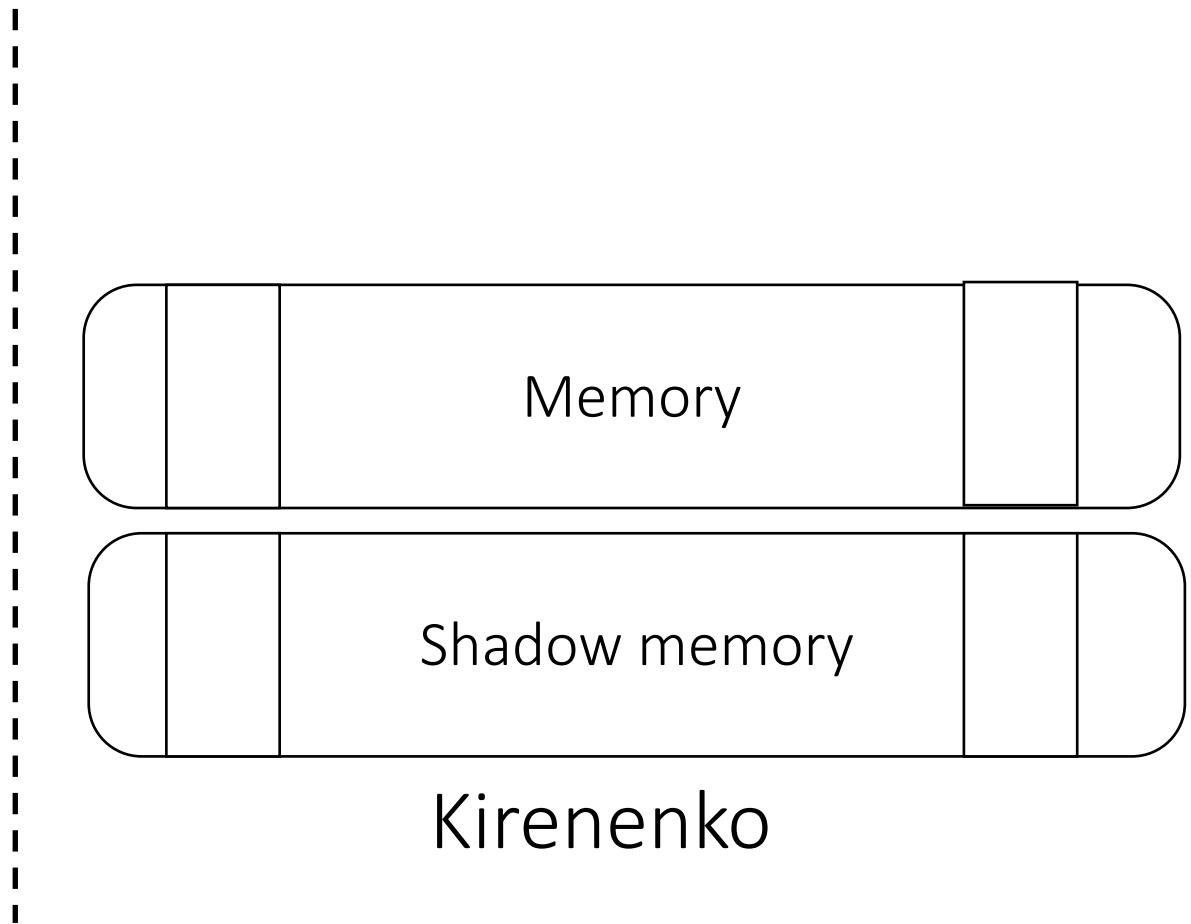
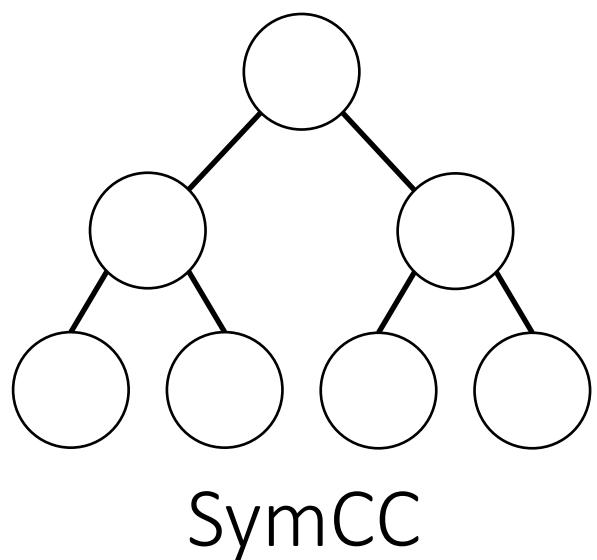
	CUTE	SymCC	Kirenenko
Language	C	LLVM IR	LLVM IR
Memory modeling	Page table	Page table	Shadow memory
Multi-threading			✓

File system	
btrfs	z3 exception
ext4	Deadlock
f2fs	Deadlock
xfs	Deadlock

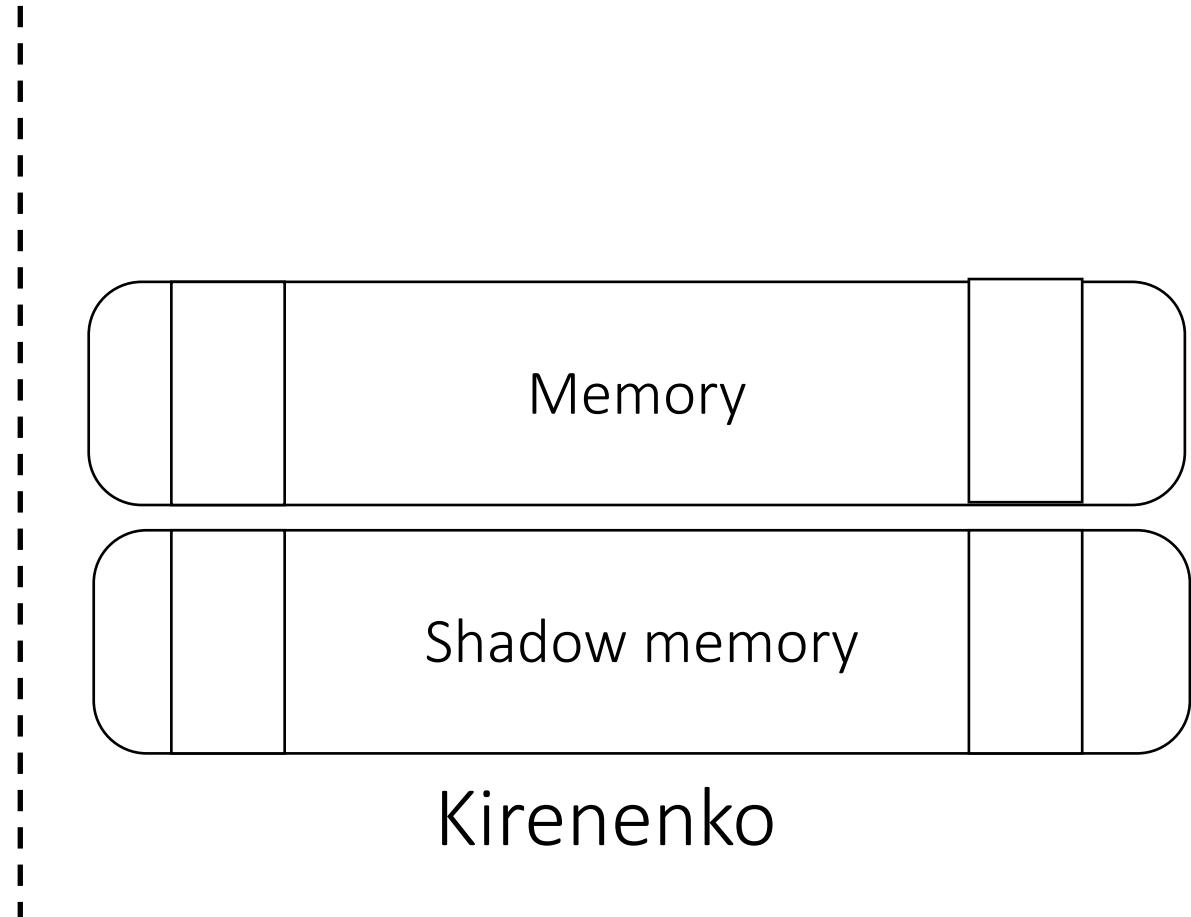
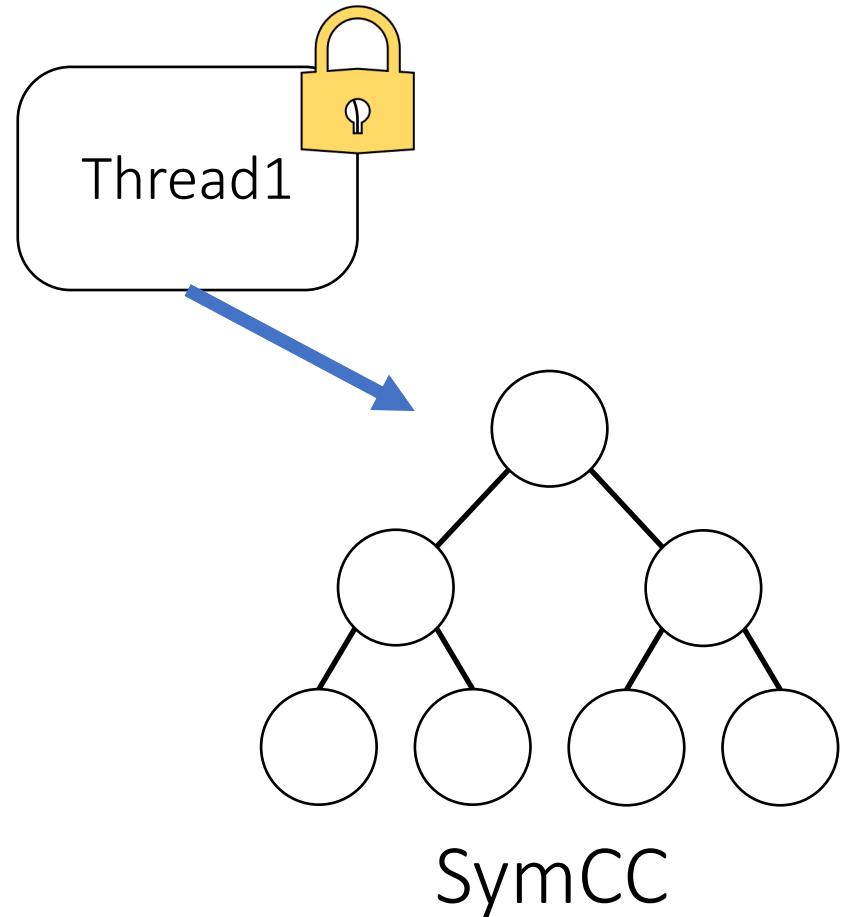
Hybridra's compilation-based concolic execution, Kirenenko, is useful by supporting multi-threading

	CUTE	SymCC	Kirenenko
<b>Language</b>	C	LLVM IR	LLVM IR
<b>Memory modeling</b>	Page table	Page table	Shadow memory
<b>Multi-threading</b>			✓

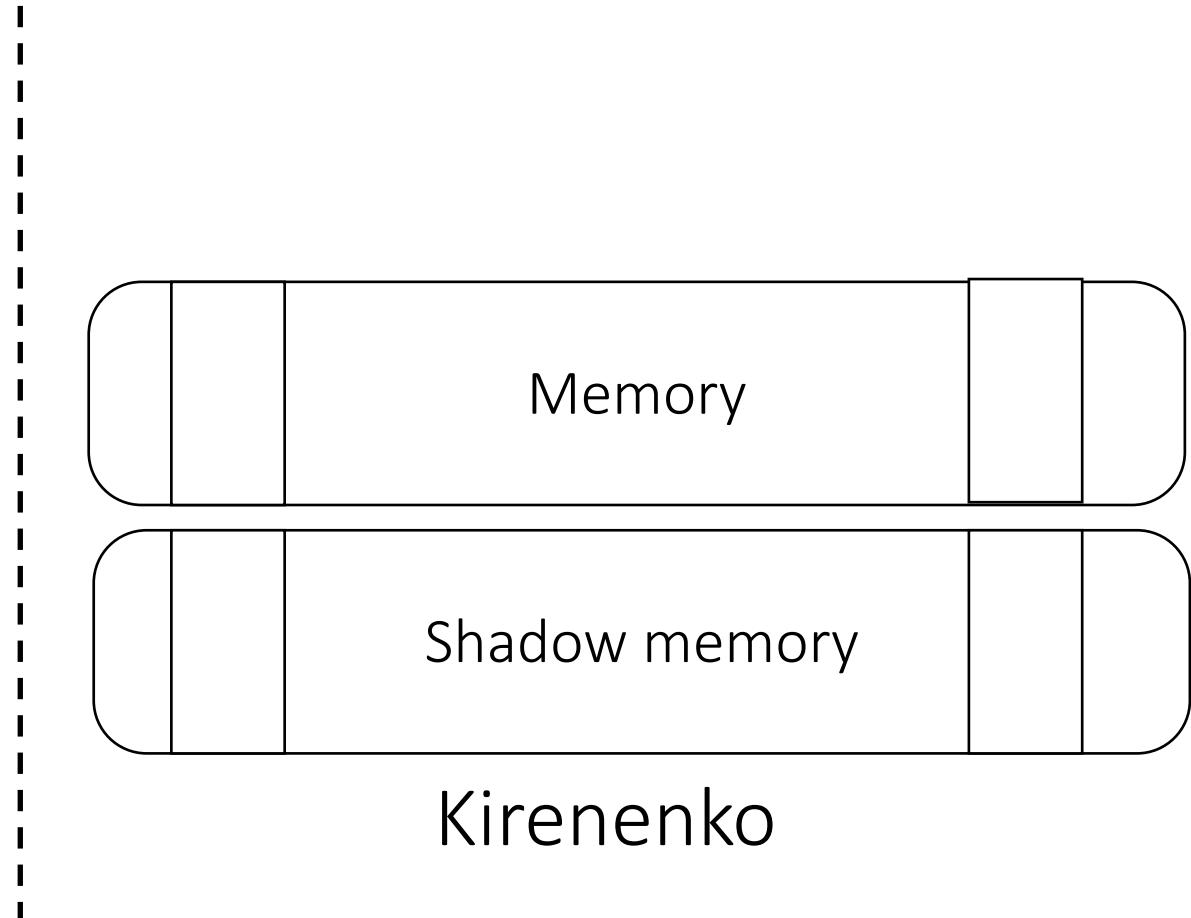
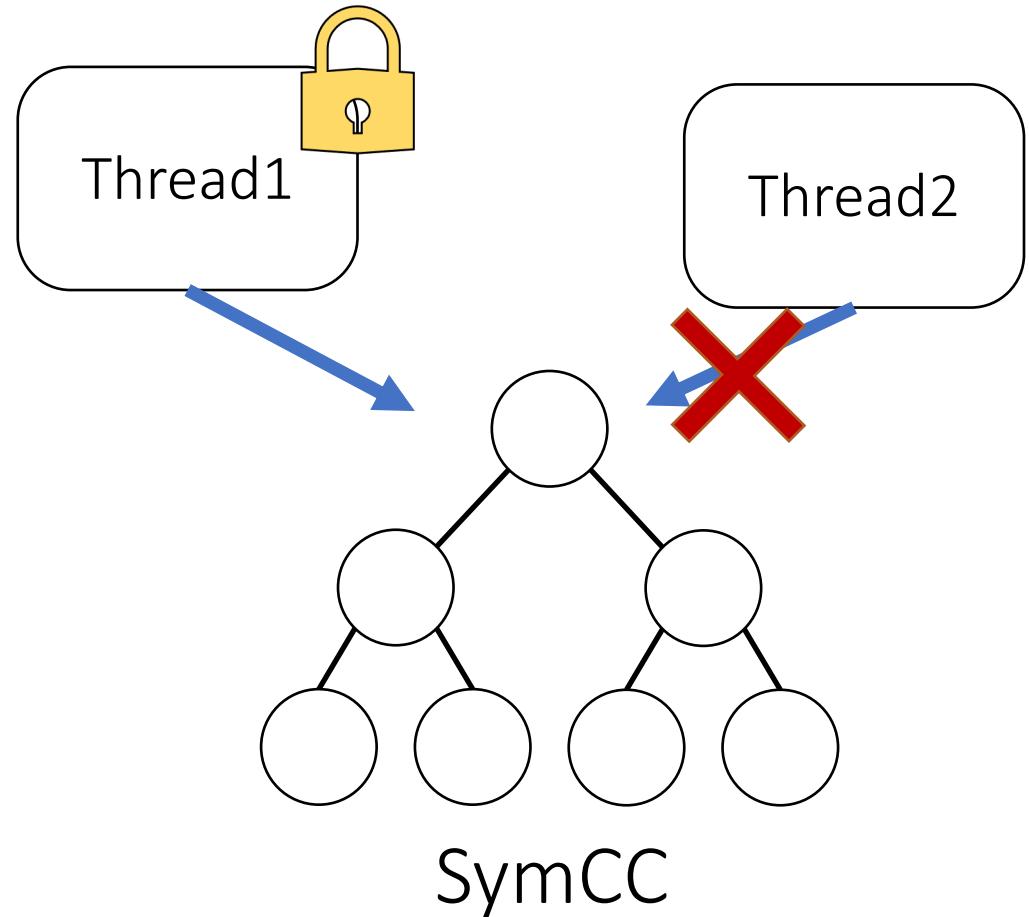
# Comparison: Memory modeling



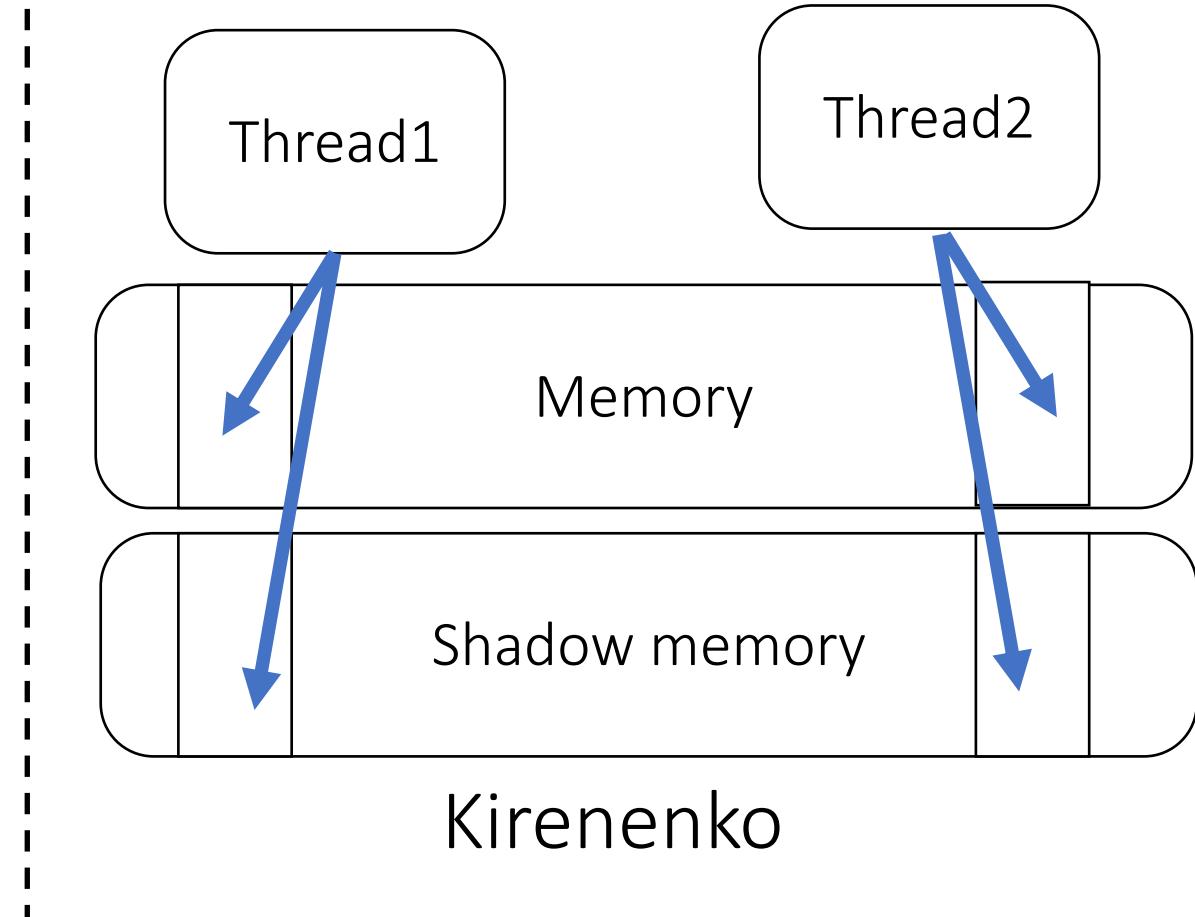
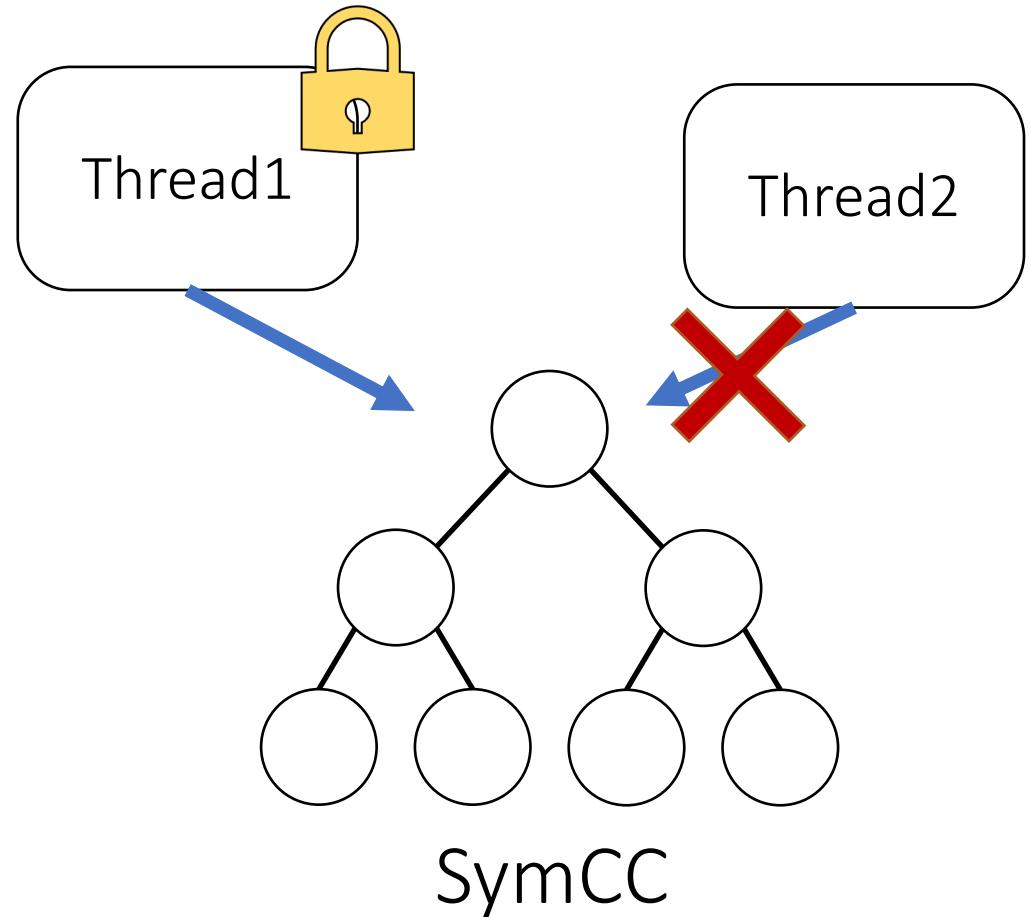
# Comparison: Memory modeling



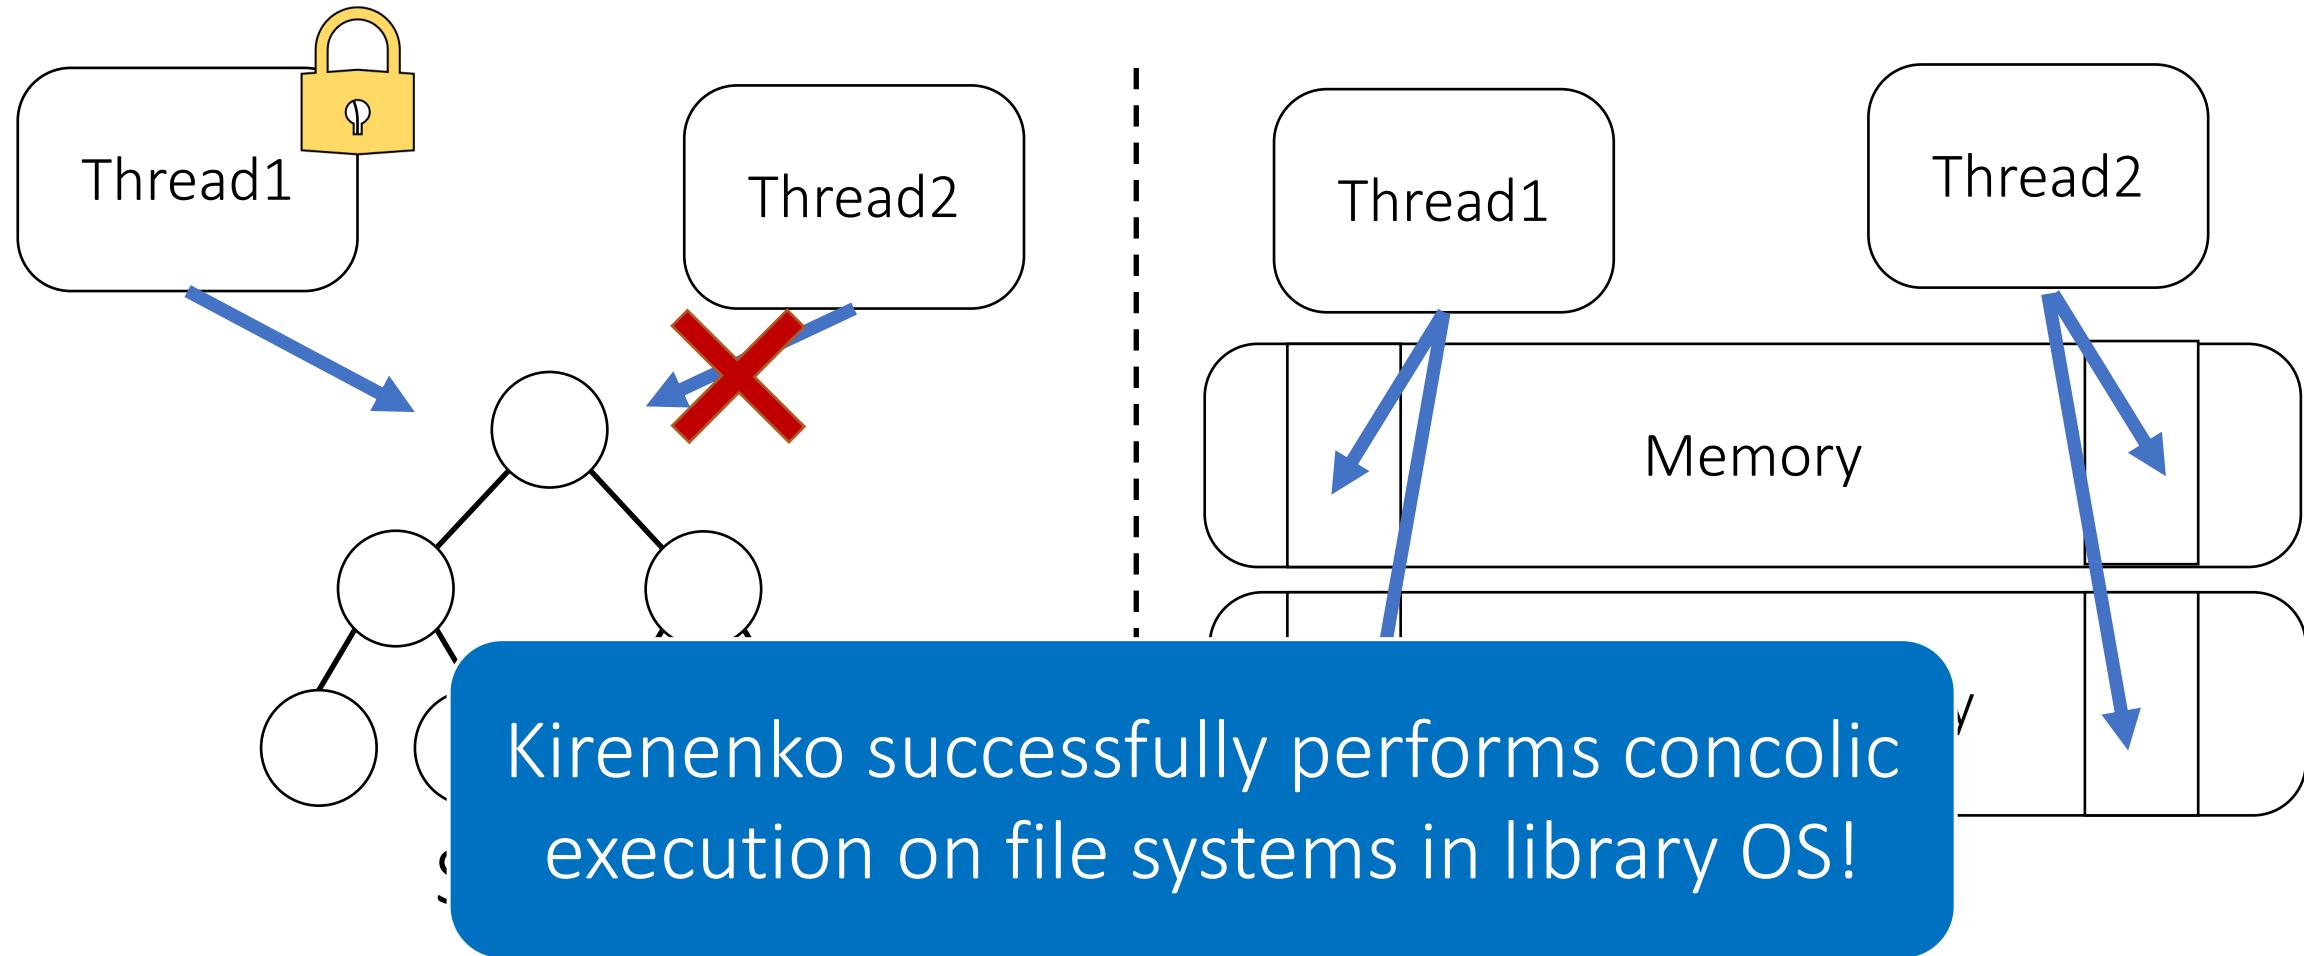
# Comparison: Memory modeling



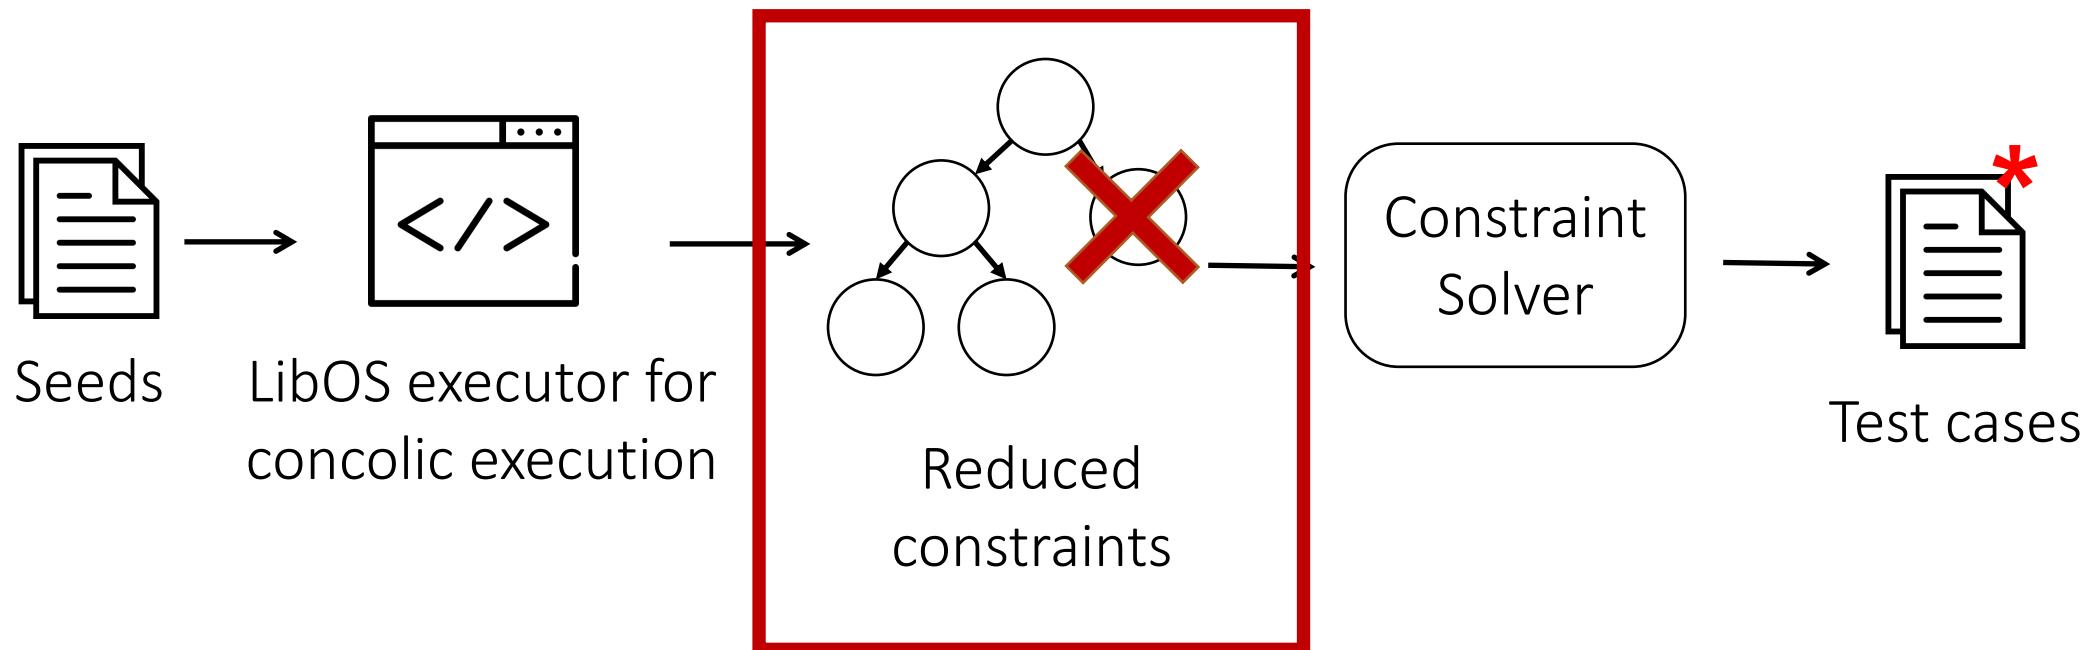
# Comparison: Memory modeling



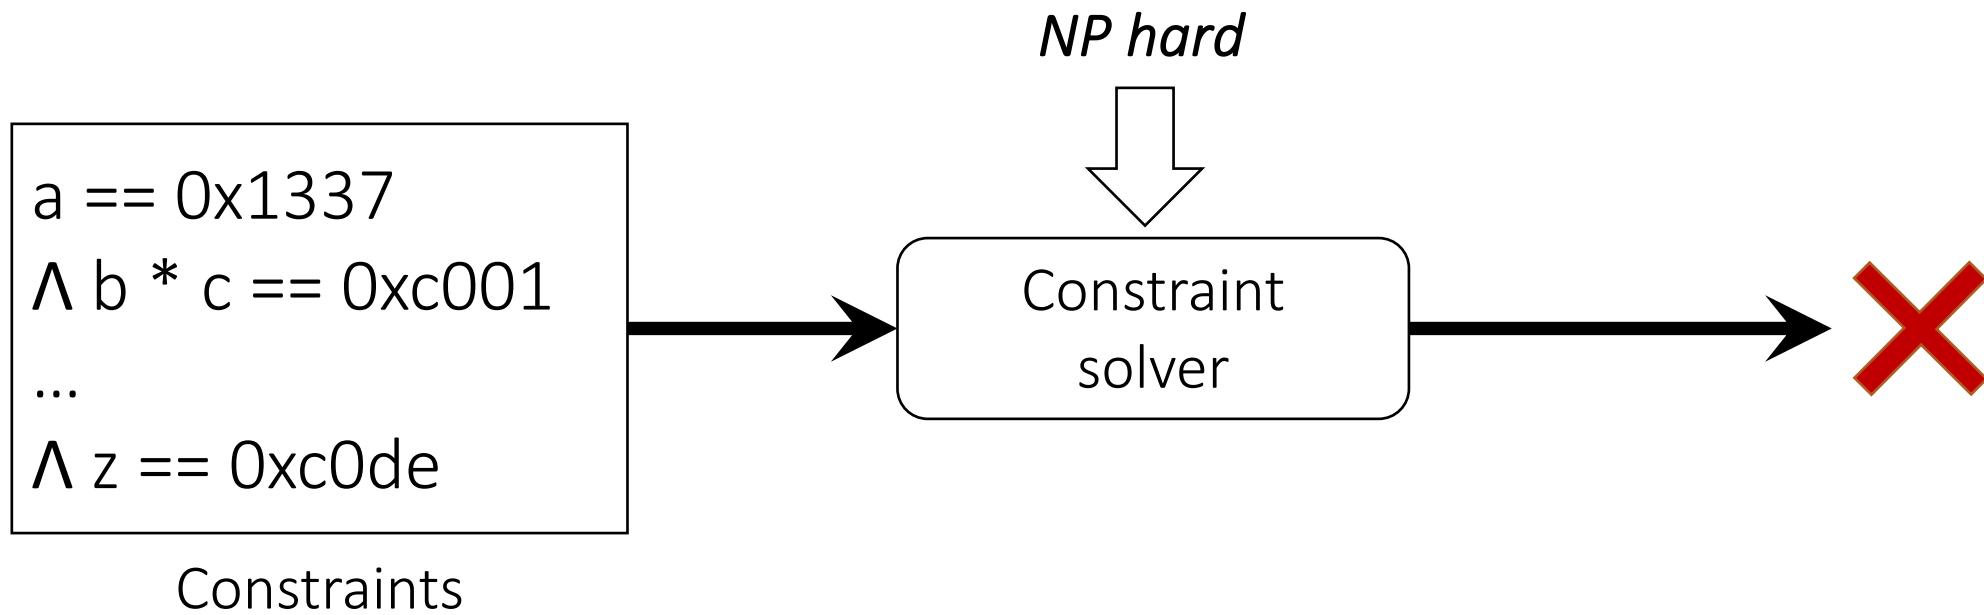
# Comparison: Memory modeling



# Design: Concolic Image Mutator



# Remind: Constraints solving is hard!



# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

```
a == 0x1337  
Λ 13 * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: Complexity  
(e.g., Linear reduction)

# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

```
a == 0x1337  
Λ 13 * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: Complexity  
(e.g., Linear reduction)

Fast algorithm

# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

```
a == 0x1337  
Λ 13 * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: Complexity  
(e.g., Linear reduction)

Fast algorithm

Limited expressiveness

# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

```
a == 0x1337  
Λ 13 * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: Complexity  
(e.g., Linear reduction)

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: # of constraints  
(e.g., Basic block pruning)

Fast algorithm

Limited expressiveness

# Two types of constraints reduction exist

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Constraints

```
a == 0x1337  
Λ 13 * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: Complexity  
(e.g., Linear reduction)

```
a == 0x1337  
Λ b * c == 0xc001  
...  
Λ z == 0xc0de
```

Reduction: # of constraints  
(e.g., Basic block pruning)

Fast algorithm

Limited expressiveness

Expressive

# Two types of constraints reduction exist

a == 0x1337  
 $\wedge$  b \* c == 0xc001  
...  
 $\wedge$  z == 0xc0de

Constraints

a == 0x1337  
 $\wedge$  13 \* c == 0xc001  
...  
 $\wedge$  z == 0xc0de

Reduction: Complexity  
(e.g., Linear reduction)

~~a == 0x1337~~  
 $\wedge$  b \* c == 0xc001  
...  
 $\wedge$  z == 0xc0de

Reduction: # of constraints  
(e.g., Basic block pruning)

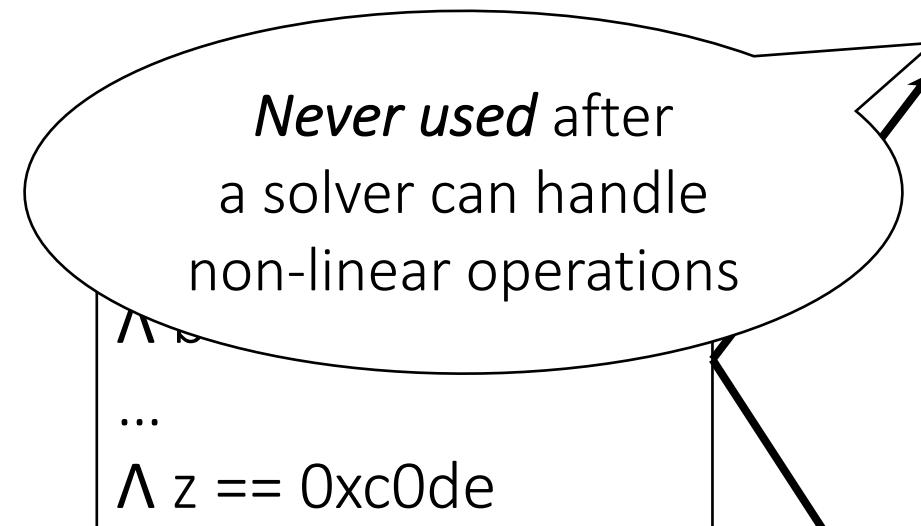
Fast algorithm

Limited expressiveness

Expressive

No algorithmic improvement

# Two types of constraints reduction exist



$a == 0x1337$   
 $\wedge 13 * c == 0xc001$   
...  
 $\wedge z == 0xc0de$

Reduction: Complexity  
(e.g., Linear reduction)

~~$a == 0x1337$~~   
 $\wedge b * c == 0xc001$   
...  
 $\wedge z == 0xc0de$

Reduction: # of constraints  
(e.g., Basic block pruning)

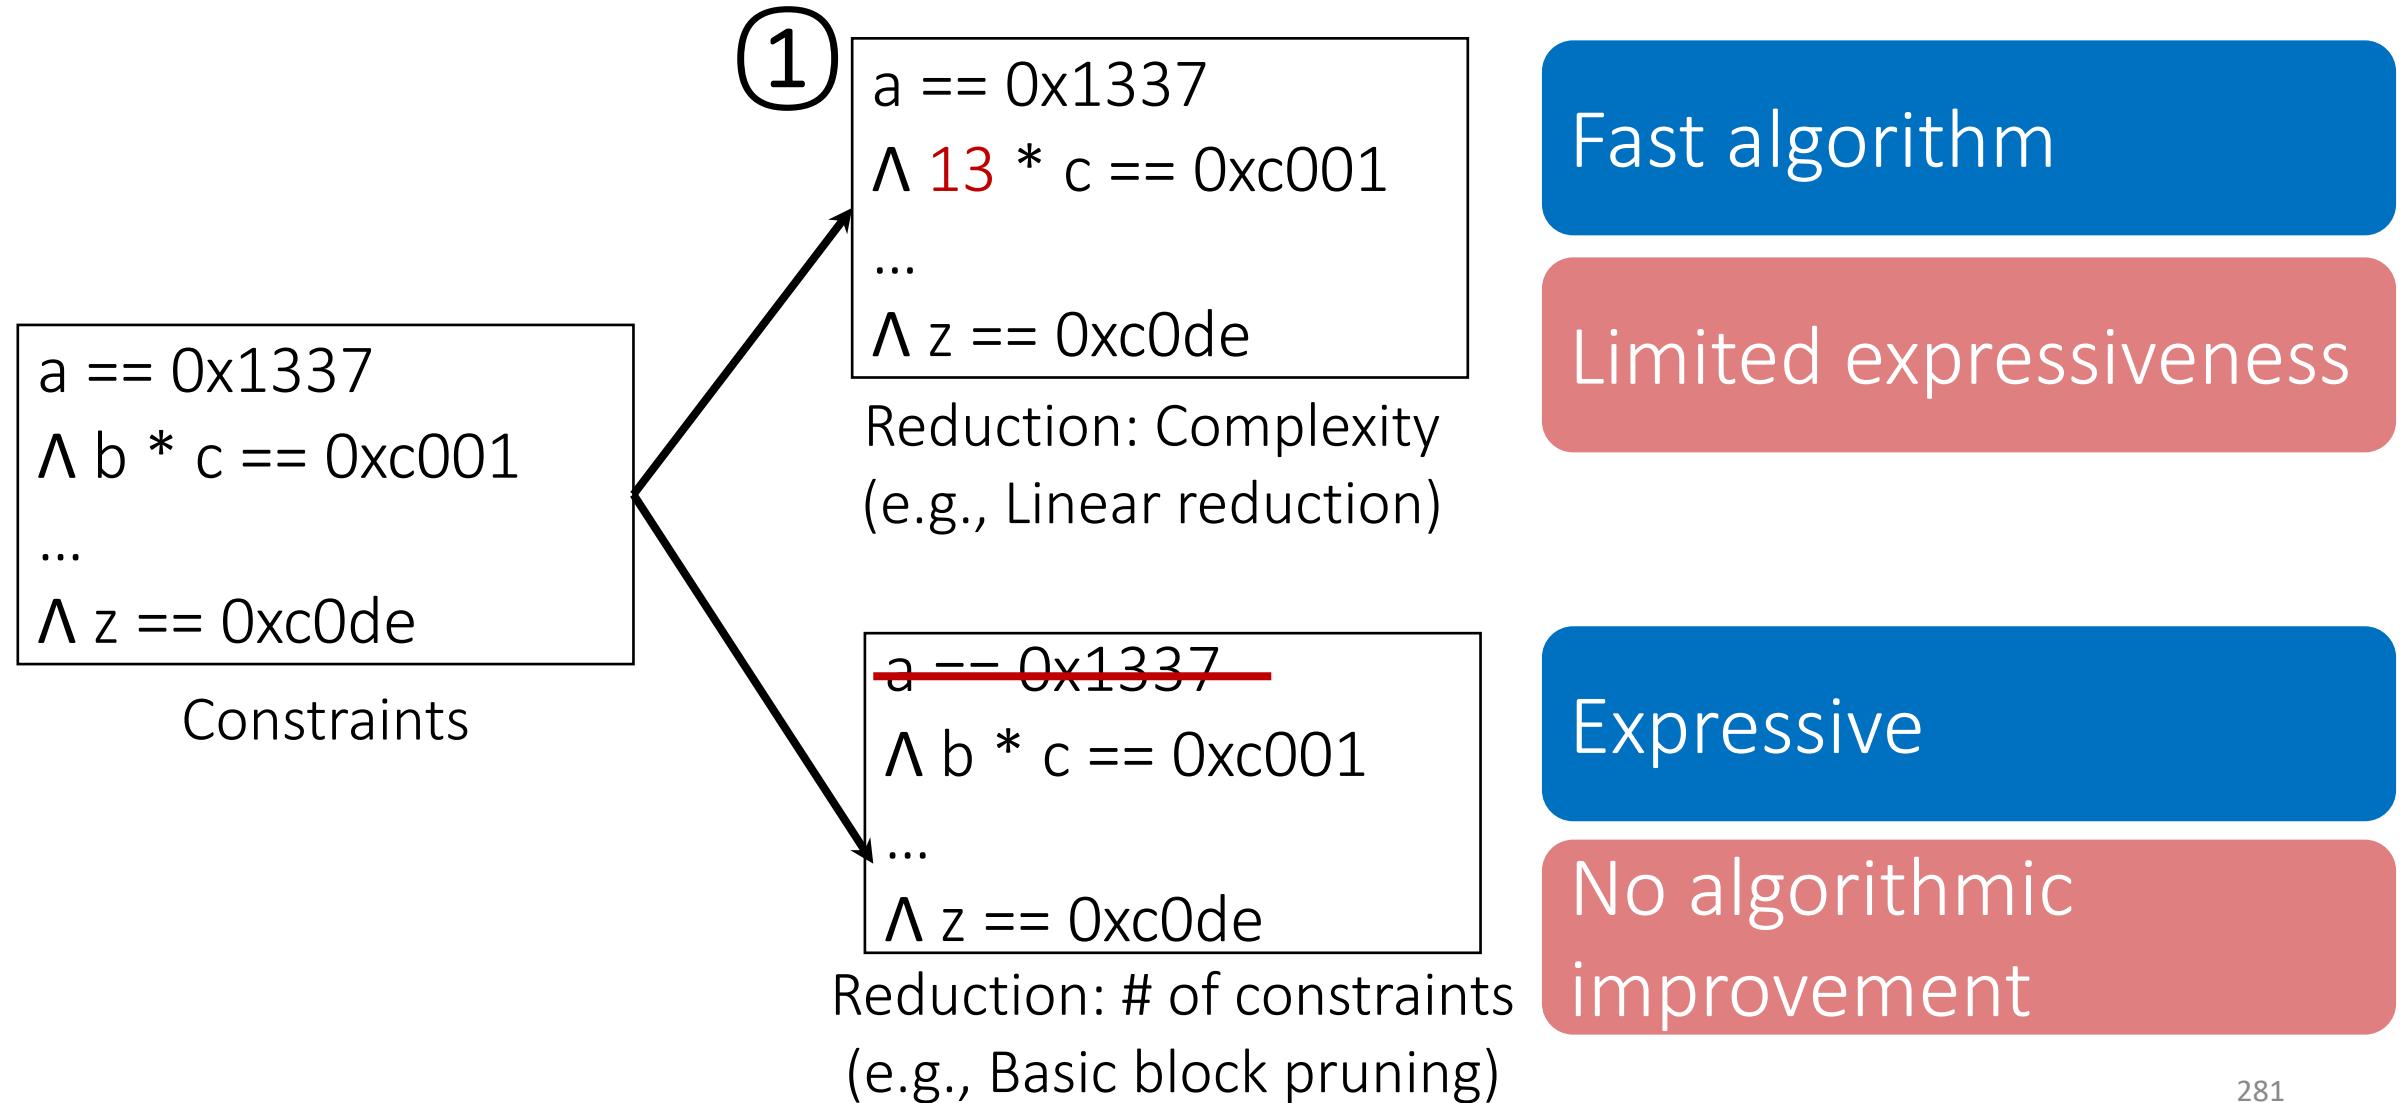
Fast algorithm

Limited expressiveness

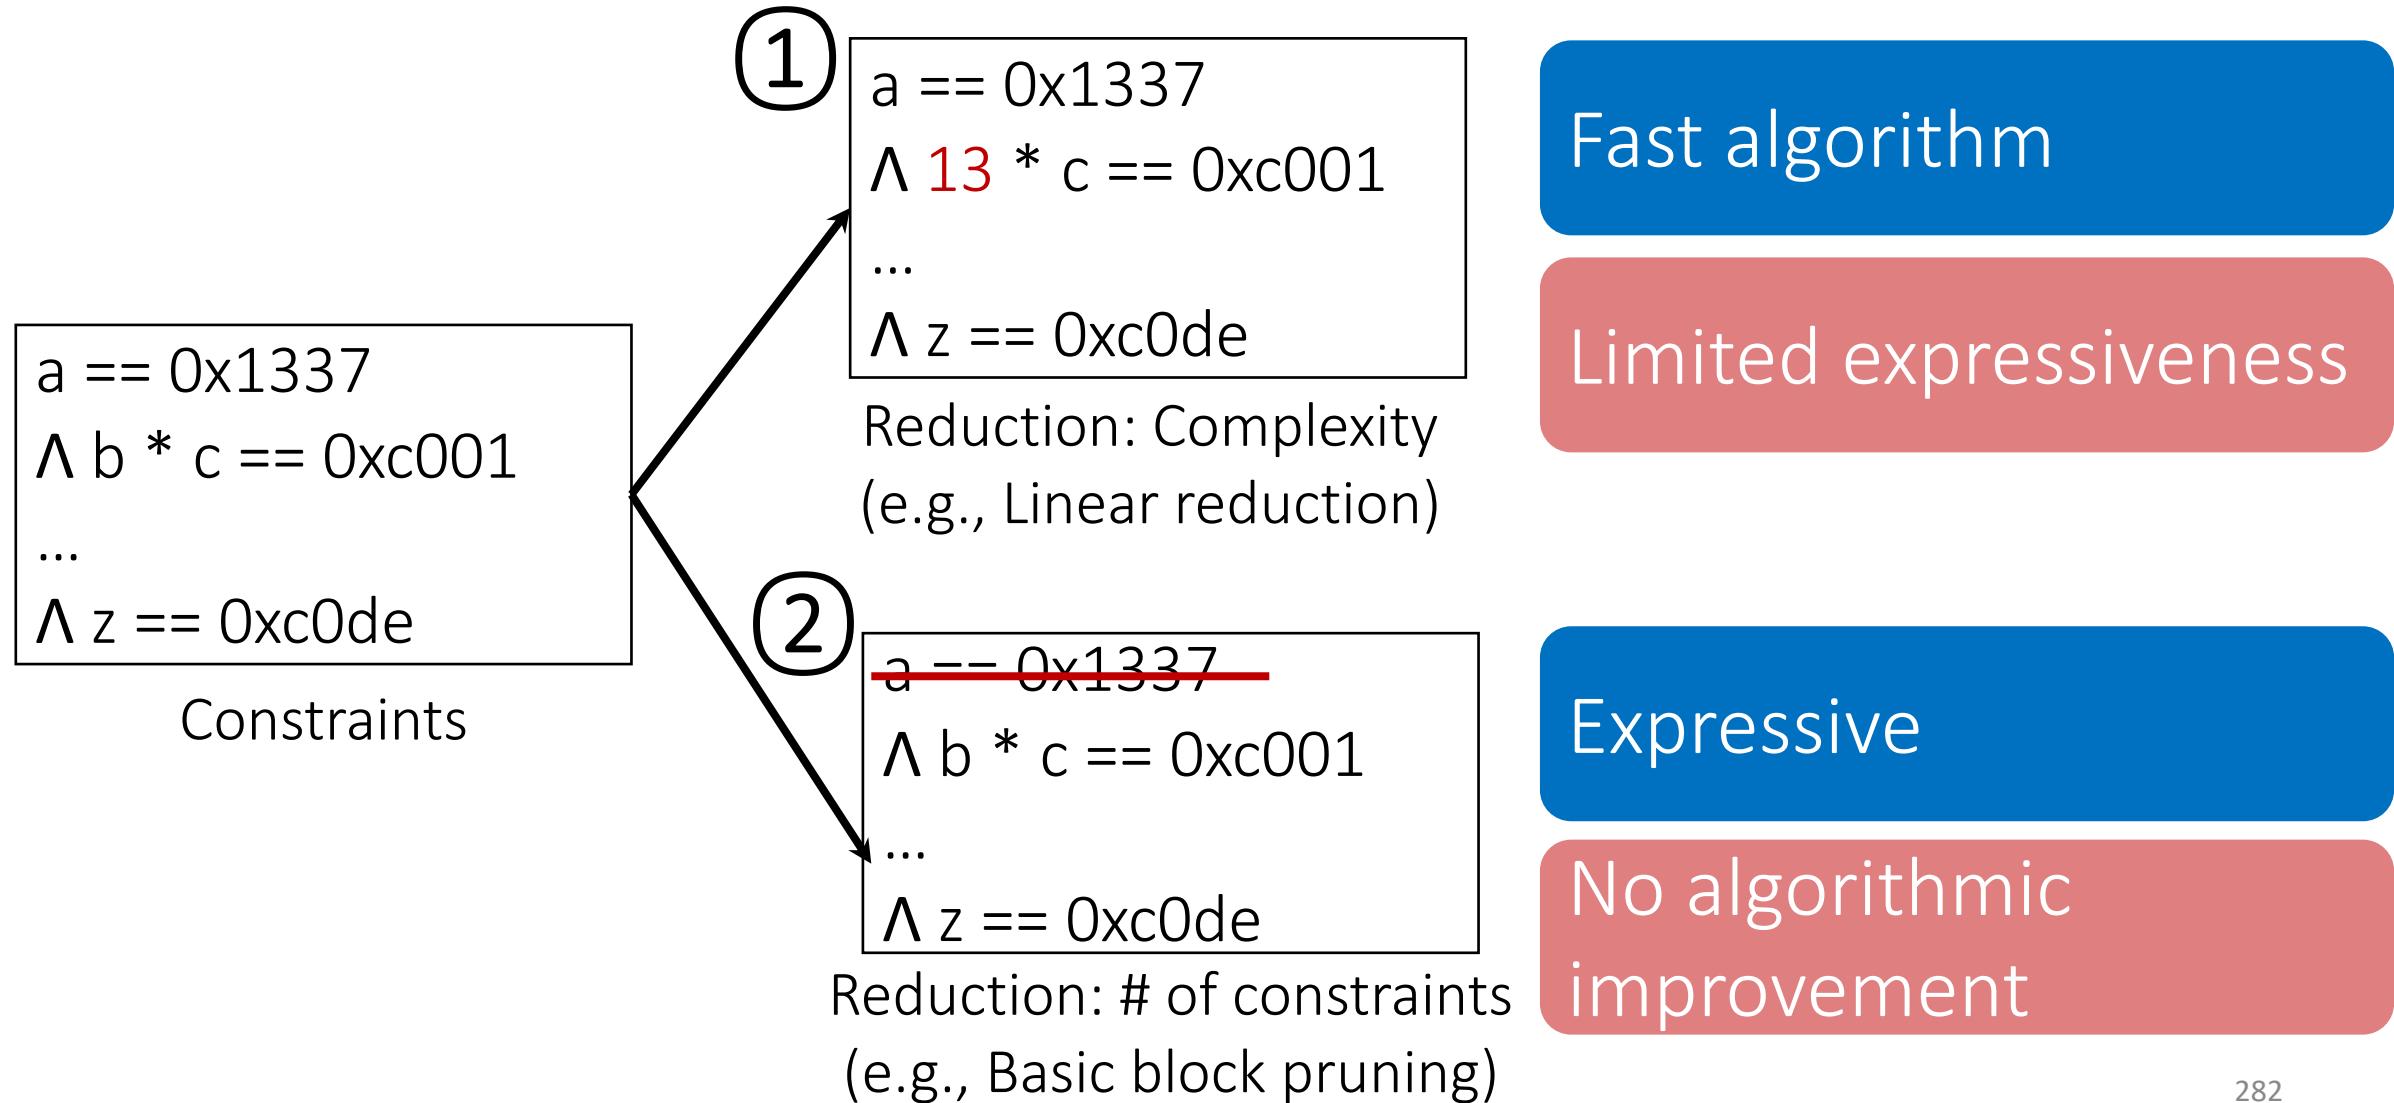
Expressive

No algorithmic improvement

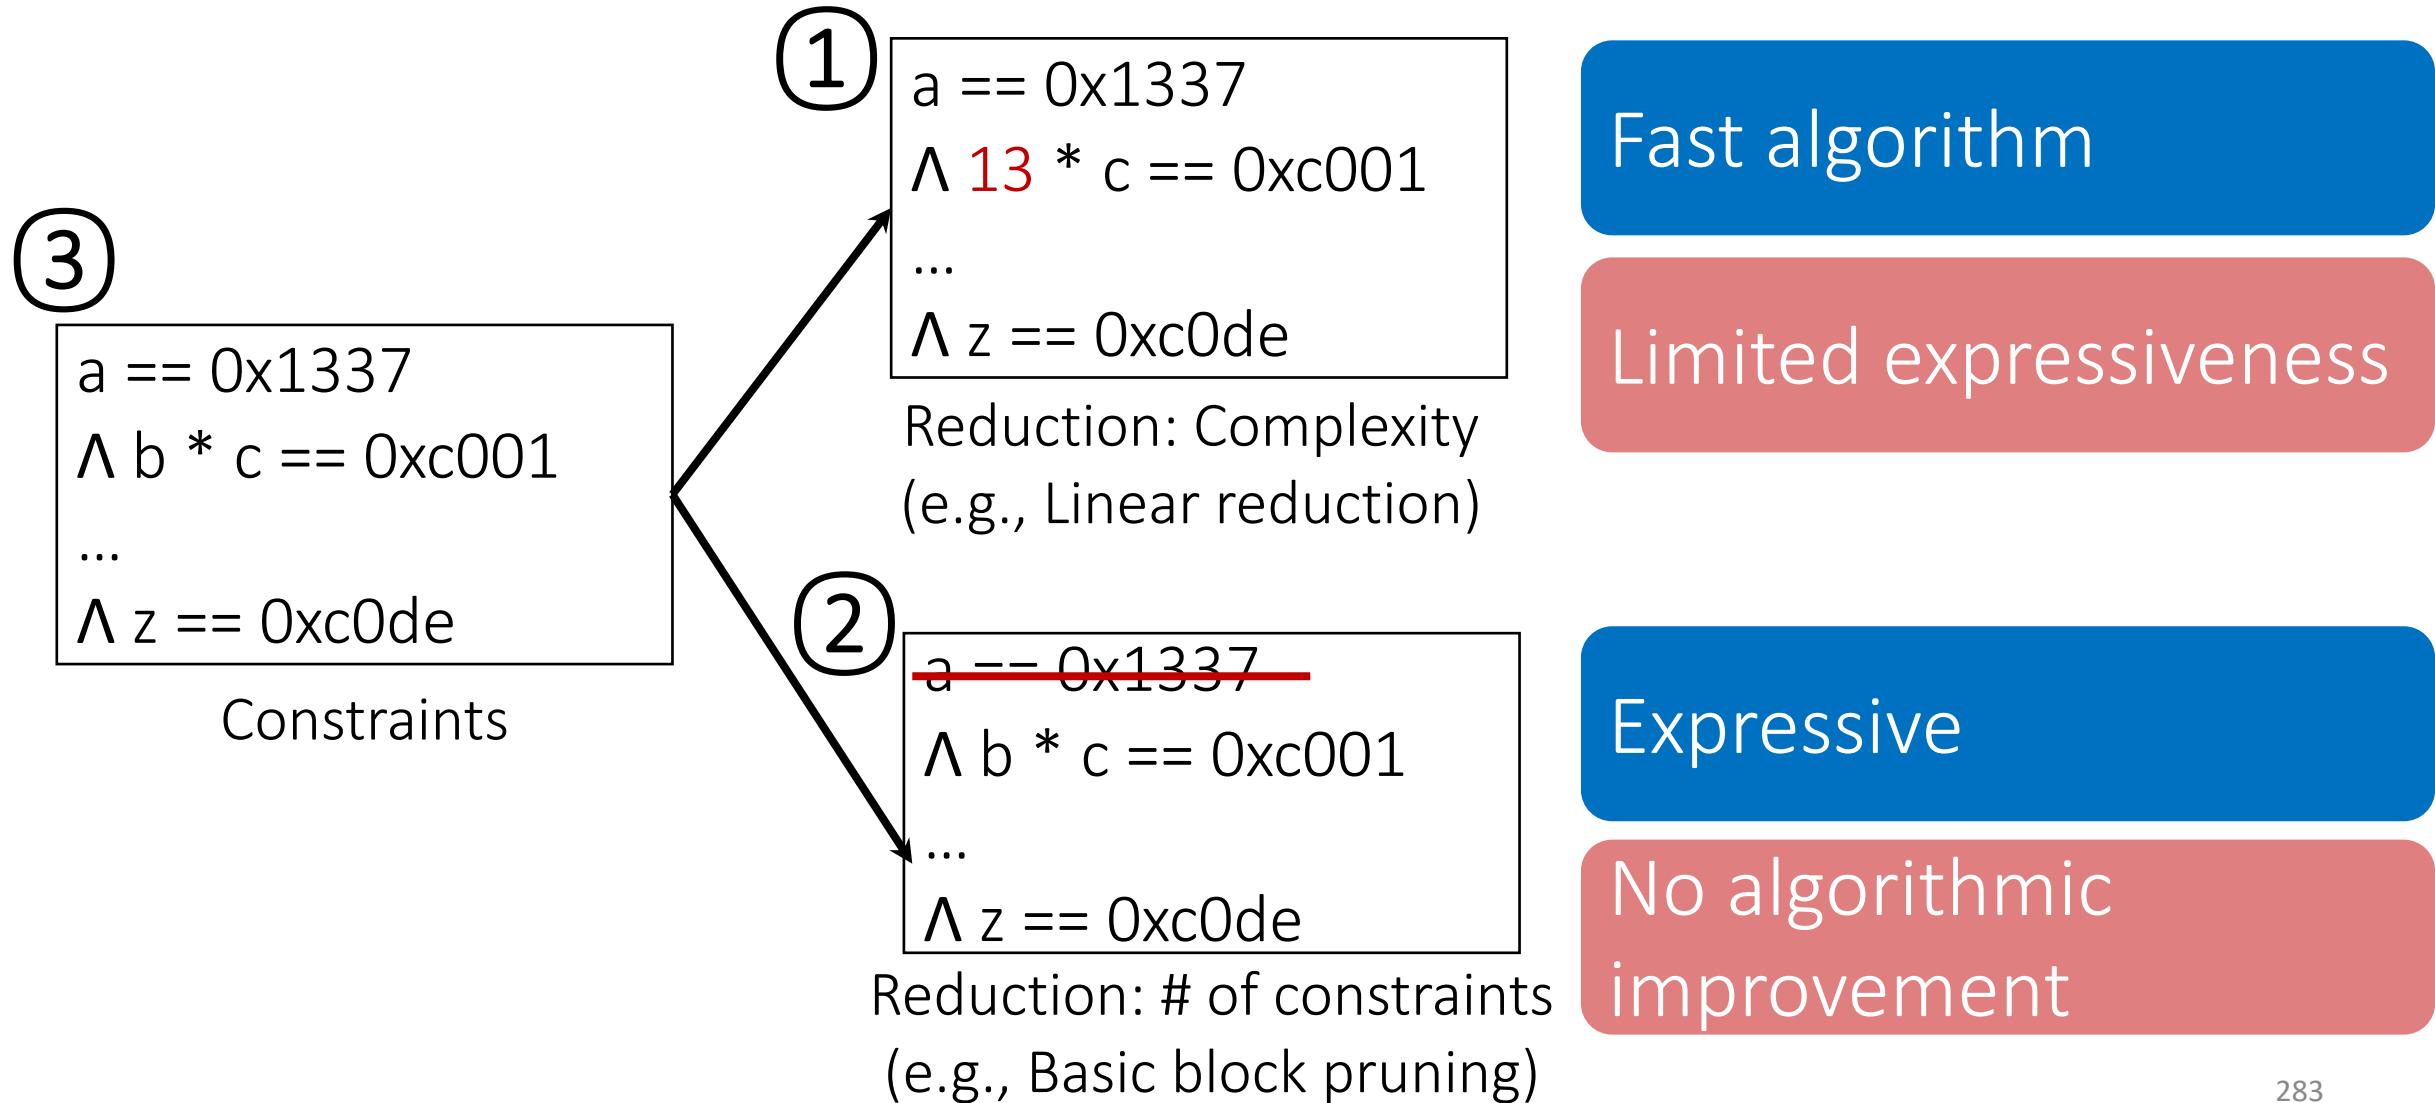
# Staged reduction: combine both reduction mechanisms



# Staged reduction: combine both reduction mechanisms

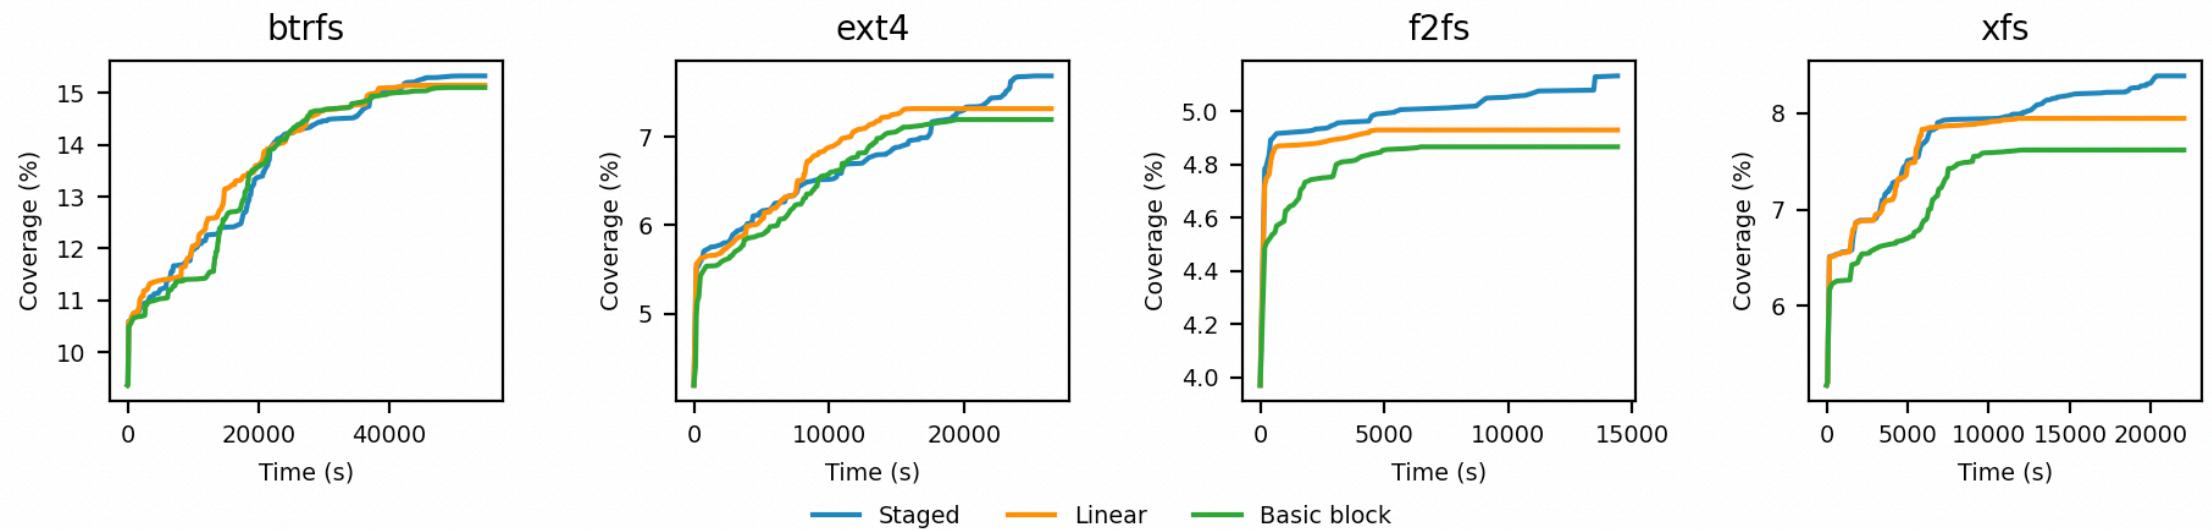


# Staged reduction: combine both reduction mechanisms

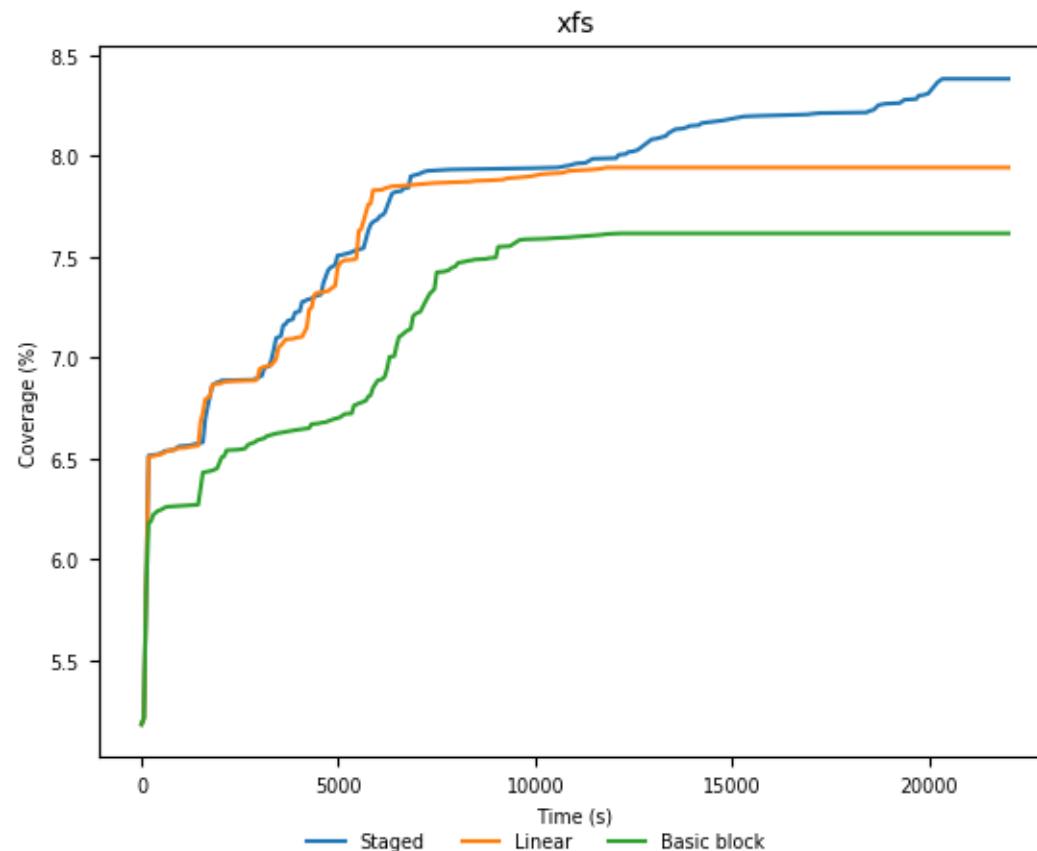


# Staged reduction outperforms each reduction mechanism

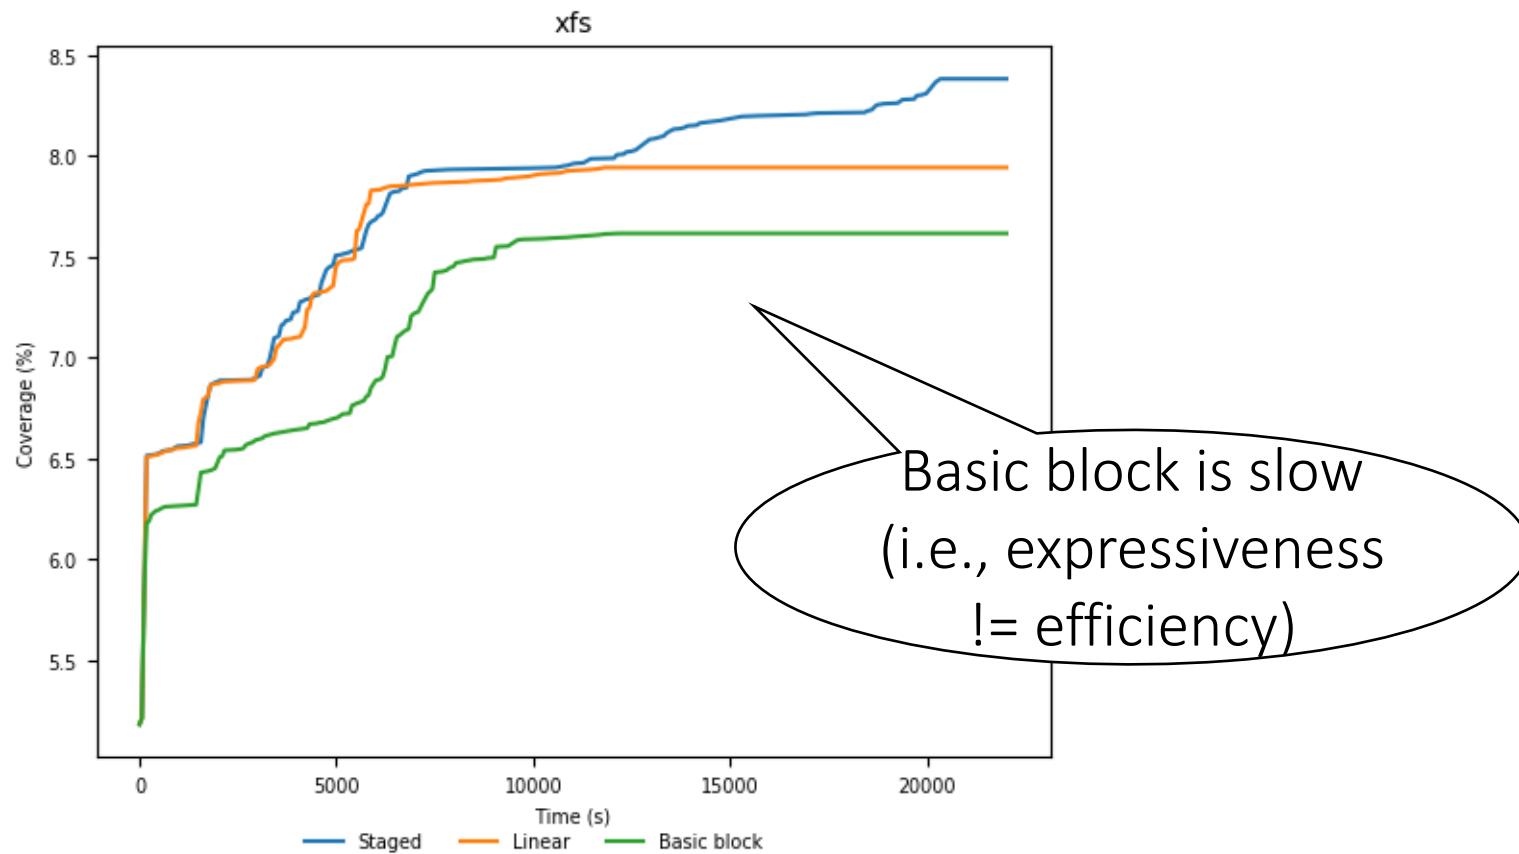
- Setting: Concolic image only, fixed timeout (9 min, 24 hours)



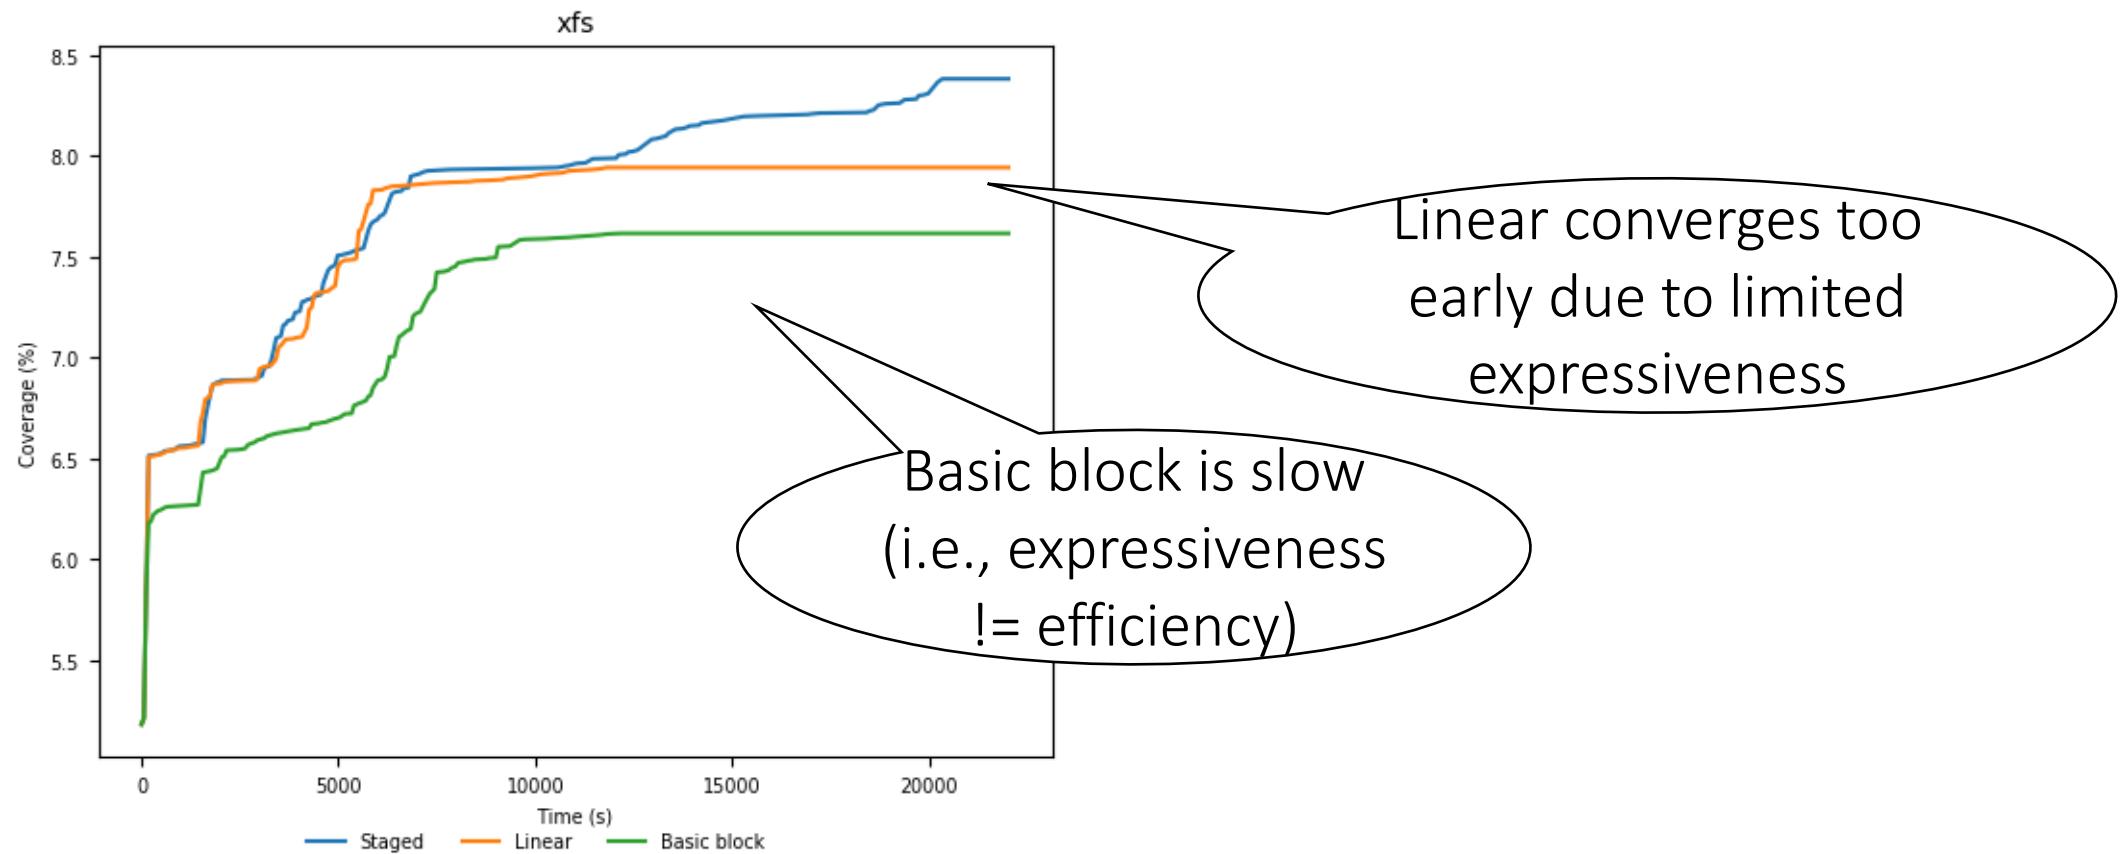
Staged reduction outperforms each reduction mechanism



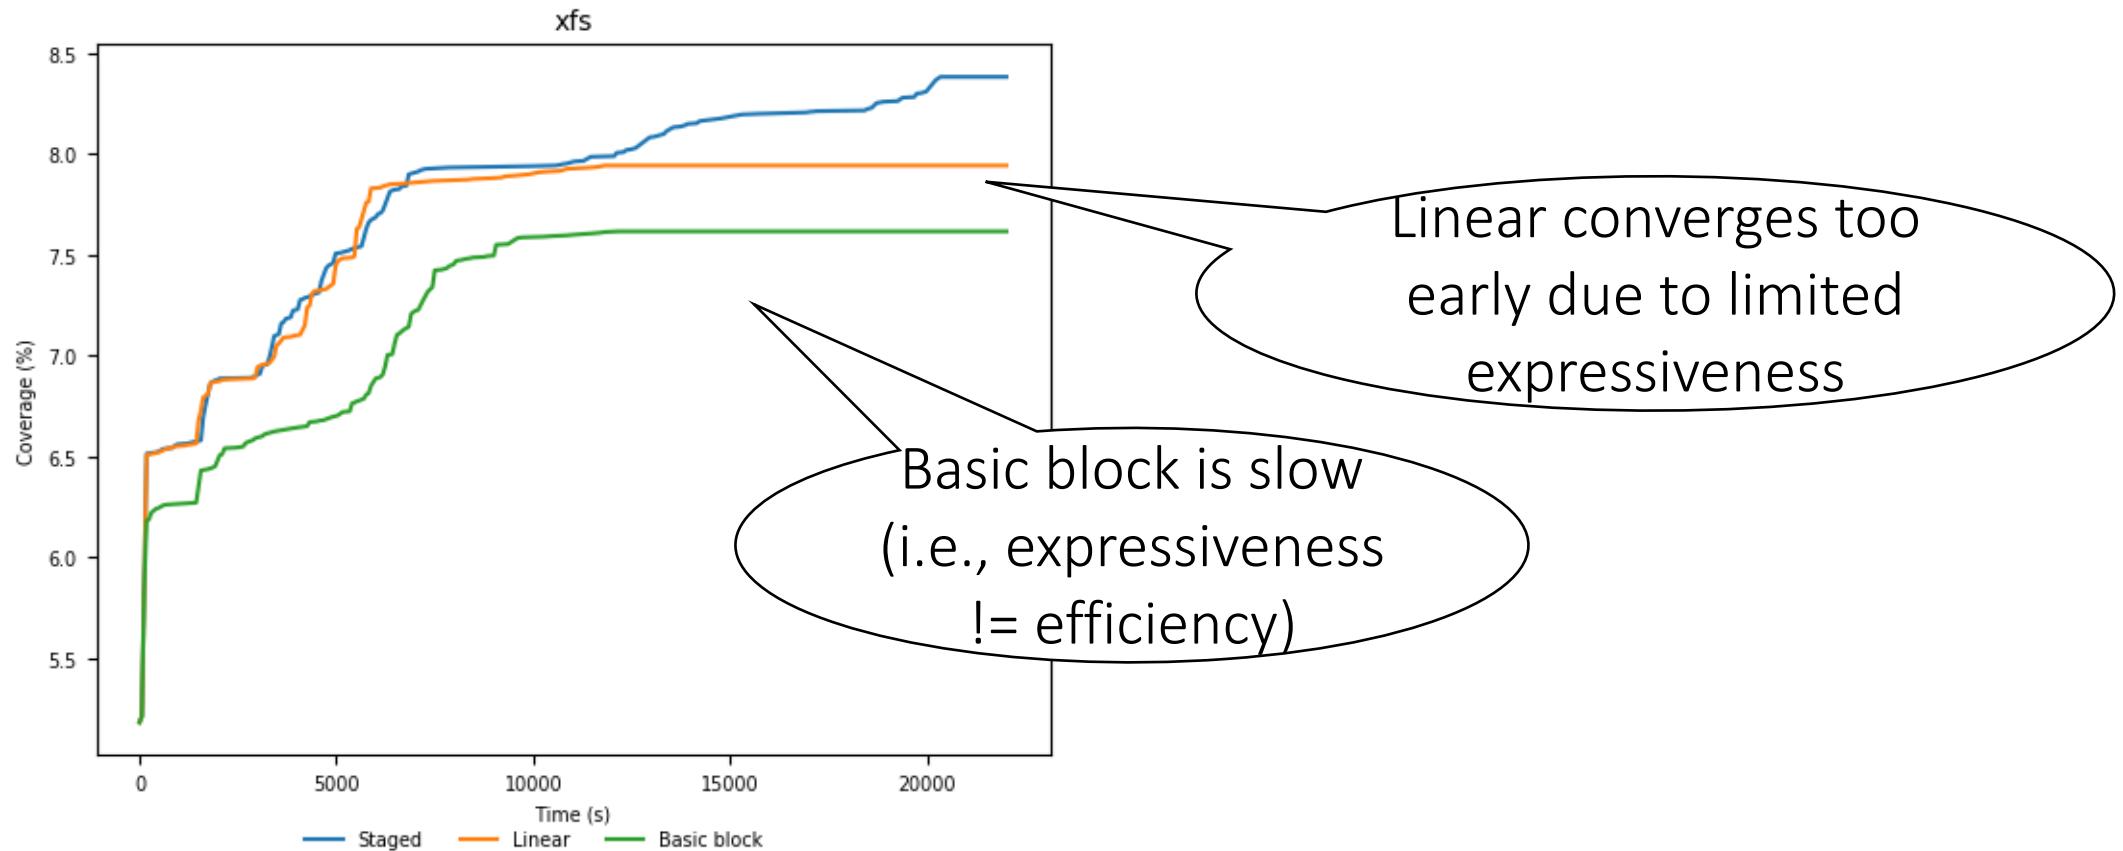
Staged reduction outperforms each reduction mechanism



Staged reduction outperforms each reduction mechanism



Staged reduction outperforms each reduction mechanism



Combining both techniques is useful  
to achieve higher code coverage!

# Evaluation

- Effective to discover new bugs in file systems?
- Outperforms the fuzzing-only solution, Hydra?

# Hybridra is effective in finding bugs in file systems

- We fuzz for 2 weeks
  - Each fuzzing takes 24 hours
- Target: Linux v5.3 (LKL), but the latest Linux is v5.8

<b>File system</b>	<b>File</b>	<b>Function</b>	<b>Type</b>	<b>Concolic</b>	<b>New</b>
btrfs	fs/btrfs/extent_io.c	extent_io_tree_panic	Null pointer dereference	✓	✓
	fs/btrfs/free-space-cache.c	tree_insert_offset	BUG()	✓	
	fs/btrfs/extent-tree.c	btrfs_drop_snapshot	BUG()		
	fs/btrfs/extent-tree.c	walk_down_proc	BUG()	✓	
	fs/btrfs/relocation.c	merge_reloc_root	BUG()		
	fs/btrfs/root-tree.c	btrfs_find_root	BUG()	✓	✓
	fs/btrfs/ctree.c	setup_items_for_insert	BUG()		✓
	fs/btrfs/volumes.c	calc_stripe_length	Divide by zero		✓
ext4	fs/ext4/super.c	ext4_clear_journal_err	BUG()		
f2fs	fs/f2fs/segment.c	f2fs_build_segment_manager	Out-of-bounds read		

# Hybridra is effective in finding bugs in file systems

- We fuzz for 2 weeks
  - Each fuzzing takes 24 hours
- Target: Linux v5.3 (LKL), but the latest Linux is v5.8

Four new bugs

File system	File	Function	Type	Concolic	New
btrfs	fs/btrfs/extent_io.c	extent_io_tree_panic	Null pointer dereference	✓	✓
	fs/btrfs/free-space-cache.c	tree_insert_offset	BUG()		✓
	fs/btrfs/extent-tree.c	btrfs_drop_snapshot	BUG()		
	fs/btrfs/extent-tree.c	walk_down_proc	BUG()	✓	
	fs/btrfs/relocation.c	merge_reloc_root	BUG()		
	fs/btrfs/root-tree.c	btrfs_find_root	BUG()	✓	✓
	fs/btrfs/ctree.c	setup_items_for_insert	BUG()		✓
	fs/btrfs/volumes.c	calc_stripe_length	Divide by zero		✓
ext4	fs/ext4/super.c	ext4_clear_journal_err	BUG()		
f2fs	fs/f2fs/segment.c	f2fs_build_segment_manager	Out-of-bounds read		

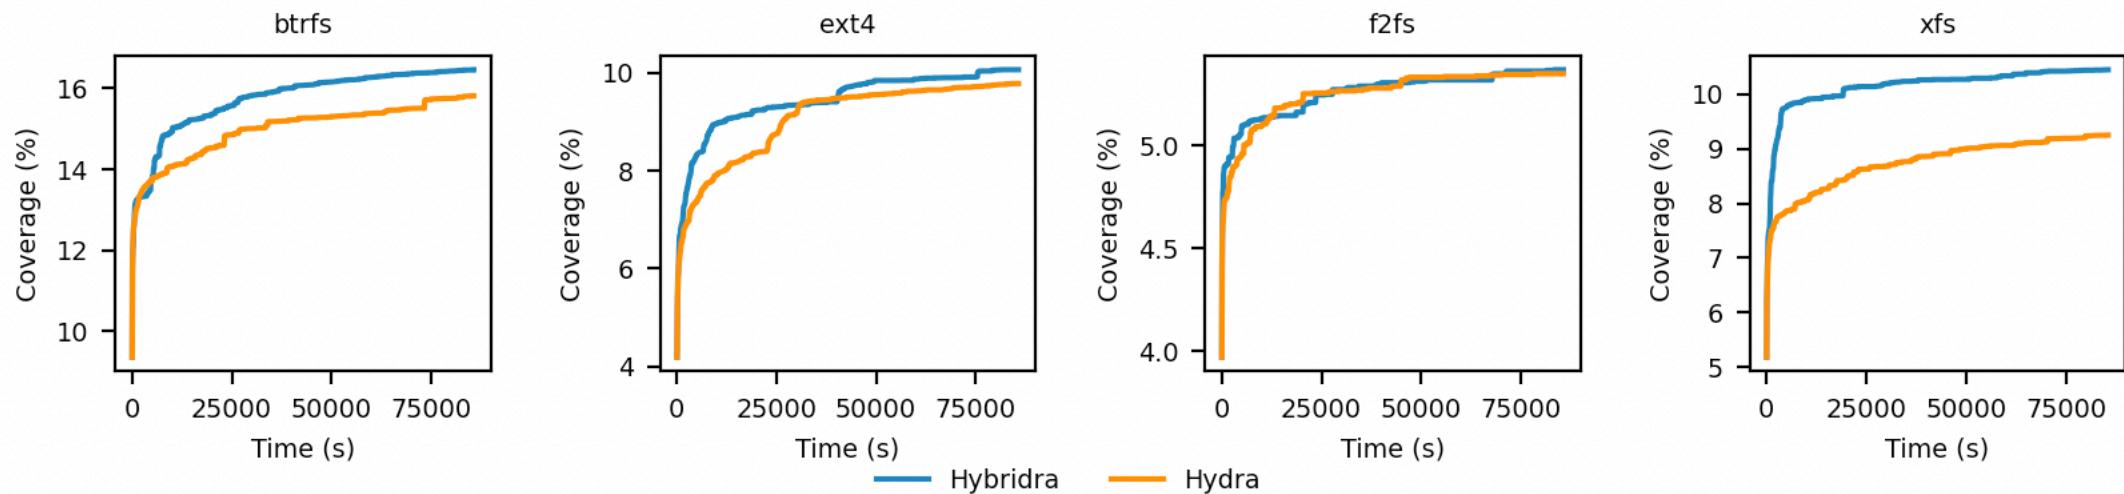
# Hybridra is effective in finding bugs in file systems

- We fuzz for 2 weeks
  - Each fuzzing takes 24 hours
- Target: Linux v5.3 (LKL), but the latest Linux is v5.8

File system	File	Function	Type	Concolic	New
btrfs	fs/btrfs/extent_io.c	extent_io_tree_panic	Null pointer dereference	✓	✓
	fs/btrfs/free-space-cache.c	tree_insert_offset	BUG()	✓	
	fs/btrfs/extent-tree.c	btrfs_drop_snapshot	BUG()		
	fs/btrfs/extent-tree.c	walk_down_proc	BUG()	✓	
	fs/btrfs/relocation.c	merge_reloc_root	BUG()		
	fs/btrfs/root-tree.c	btrfs_find_root	BUG()	✓	✓
	fs/btrfs/ctree.c	setup_items_for_insert	BUG()		✓
	fs/btrfs/volumes.c	calc_stripe_le	BUG()		✓
ext4	fs/ext4/super.c	ext	Many bugs directly from concolic execution e.g., BUG(x != 0);		
f2fs	fs/f2fs/segment.c	f2f			

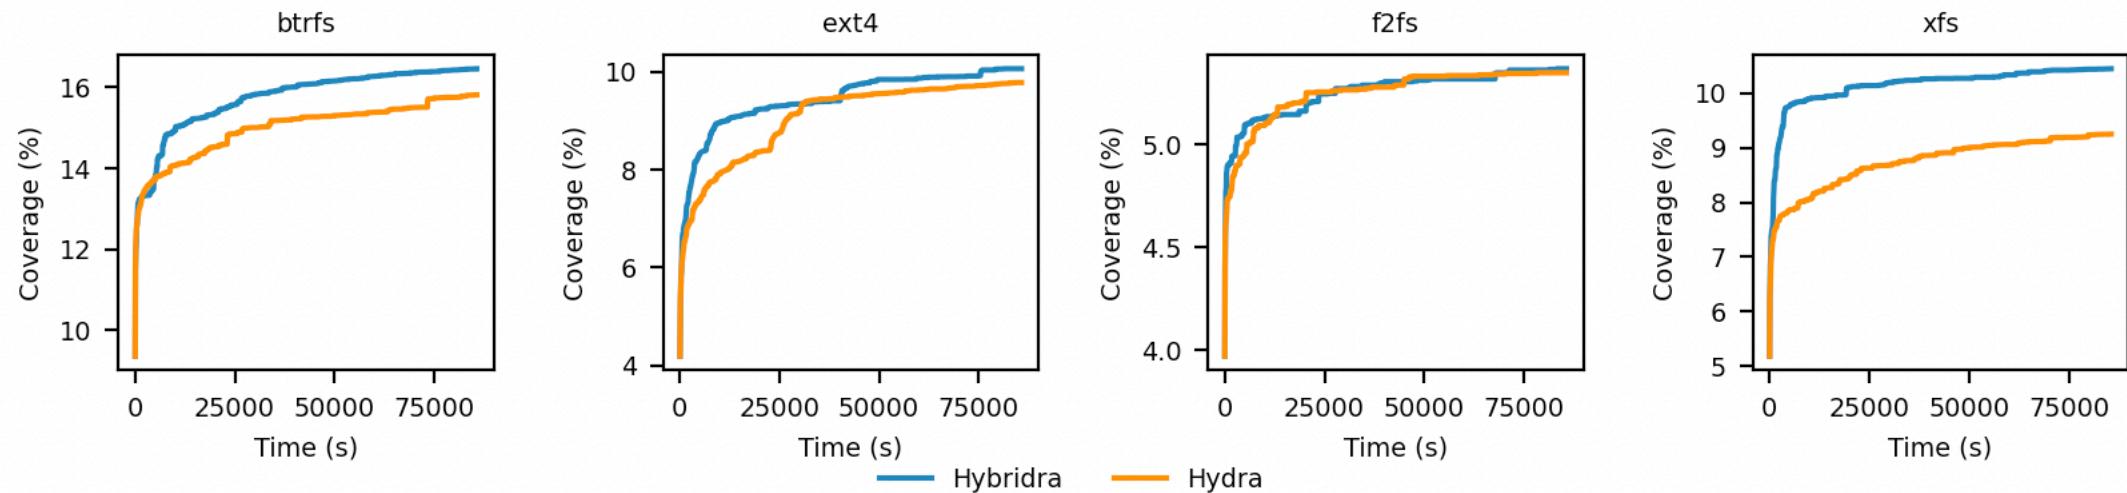
# Hybridra outperforms the fuzzing-only approach, Hydra

- Setting: Image only (+ Random), fixed timeout (24 hours)



# Hybridra outperforms the fuzzing-only approach, Hydra

- Setting: Image only (+ Random), fixed timeout (24 hours)



Concolic execution can help fuzzing in file systems  
by discovering interesting test cases!

# Discussion & Limitation

# Discussion & Limitation

- Apply to other applications
  - Our library OS (LKL) also supports network simulation.
  - Thus, it is possible to extend it to network stacks
  - We can apply other user-mode kernel (e.g., Kunit) to test other features
- Limitations
  - Currently, Hybridra does not support floating point and vector operation
  - The limited number of symbols ( $2^{30}$ ) because of shadow memory
    - In our evaluation, this is fine for testing file systems

# Conclusion

- Designing a concolic executor tailored for hybrid fuzzing is important for scaling hybrid fuzzing to real-world software
  - Systematic approaches for fast symbolic simulation
  - New heuristics for test case generation
- This dissertation demonstrates this idea with
  - QSYM: Hybrid fuzzing for binary-only applications
  - Hybridra: Hybrid fuzzing for file systems

# Acknowledgements

- Georgia Tech
  - Taesoo Kim
  - Meng Xu
  - Jinho Jung
  - Wen Xu
  - Soyeon Park
  - Daehee Jang
- Oregon State University
  - Yeongjin Jang
- Virginia Tech
  - Changwoo Min
- Seoul National University
  - Byoungyoung lee
  - Yunheung Paek
- Microsoft Research
  - Sangho Lee
  - Weidong Cui
  - Xinyang Ge
  - Ben Niu
- University of Michigan
  - Baris Kasikci
  - Upamanyu Sharma
- Arizona State University
  - Ryuou Wang
- Google
  - Chanil Cheon
- Facebook
  - Dhaval Kapil
- University of Pennsylvania
  - Xujie Si
  - Mayur Naik
- UNIST
  - Hyungon Moon
- NSRI
  - Su yong Kim
- KAIST
  - Yongdae Kim
  - Kyoungsoo Park
  - Yung Yi

Thank you!